PS5 Solutions

Jingle Fu

Problem 1

Solution (a).

Yes, there are missing values in the data set. The initial number of observations is 99457, The dataset's reported missing values indicate that only *age* has missing observations (119 missing values). After dropping these, we end up with 99338 observations.

```
rm(list = ls())
library(tidyverse)
library(ggplot2)
library(dplyr)
library(broom)
library(stats)
library(stargazer)
library(car)

set.seed(2024)

dat <- read.csv("dat_SalesCustomers.csv")
variables_to_check <- c("category", "price", "gender", "age", "payment_
    method")
missing_counts <- sapply(dat[variables_to_check], function(x) sum(is.na(
    x)))
print("Number of missing values in each variable:")
print(missing_counts)

dat_clean <- dat[complete.cases(dat[variables_to_check]), ]

num_observations <- nrow(dat_clean)
print(paste("Number of observations after removing missing values:", num
    _observations))

```

Solution (b).

Define

$$\text{paid\_in\_cash}_i = \mathbf{1}\{\text{payment\_method}_i = \text{Cash}\}$$
$$\text{male}_i = \mathbf{1}\{\text{gender}_i = \text{Male}\}$$

The fraction of transactions carried out in cash is

$$\frac{1}{n}\sum_{i=1}^{n}\text{paid\_in\_cash}_i.$$

Empirically, this is about 44.69%.
The fraction of overall sales carried out in cash is

$$\frac{\sum_{i=1}^{n}\text{paid\_in\_cash}_i \cdot price_i}{\sum_{i=1}^{n} price_i}.$$

Empirically, this fraction is about 44.79%.
These results indicate that cash payments represent nearly half of all transactions and sales value.

```
dat_clean$paid_in_cash <- ifelse(dat_clean$payment_method == "Cash", 1,
    0)
dat_clean$male <- ifelse(dat_clean$gender == "Male", 1, 0)
fraction_cash_transactions <- mean(dat_clean$paid_in_cash)
print(paste("Fraction of transactions carried out in cash:",
            round(fraction_cash_transactions * 100, 2), "%"))

total_sales <- sum(dat_clean$price)
cash_sales <- sum(dat_clean$price[dat_clean$paid_in_cash == 1])
fraction_cash_sales <- cash_sales / total_sales
print(paste("Fraction of overall sales carried out in cash:",
            round(fraction_cash_sales * 100, 2), "%"))

```

Solution (c).
We now consider only the first $n = 1000$ observations. Let the categories be divided into five mutually exclusive groups: Clothes and Shoes (C), Cosmetics (Cos), Food (F), Technology (T), and Other (O). Define indicator variables:

$$\begin{aligned}
d_{C,i} &= \mathbf{1}\{\text{category}_i = \text{Clothes and Shoes}\}\\
d_{Cos,i} &= \mathbf{1}\{\text{category}_i = \text{Cosmetics}\}\\
d_{F,i} &= \mathbf{1}\{\text{category}_i = \text{Food}\}\\
d_{T,i} &= \mathbf{1}\{\text{category}_i = \text{Technology}\}\\
d_{O,i} &= 1 - (d_{C,i} + d_{Cos,i} + d_{F,i} + d_{T,i}).
\end{aligned}$$

The fraction of transactions in category j is

$$\frac{1}{1000}\sum_{i=1}^{1000} d_{j,i}.$$

The fraction of sales in category j is

$$\frac{\sum_{i=1}^{1000} d_{j,i} \cdot price_i}{\sum_{i=1}^{1000} price_i}.$$

Empirically:

- Transactions fraction: Clothes/Shoes: 43.8%, Cosmetics: 14.8%, Food: 14.0%, Technology: 5.0%, Other: 22.4%.
- Sales fraction: Clothes/Shoes: 70.58%, Cosmetics: 2.72%, Food: 0.32%, Technology: 23.9%, Other: 2.49%.

The result shows that most transactions and sales are in the Clothes/Shoes category. Technology, though having the lowest transaction fraction, has the second-highest sales fraction, meaning that it has the highest average price.

```r
dat_1000 <- dat_clean[1:1000, ]

dat_1000$clothes_shoes <- ifelse(dat_1000$category %in% c("Clothing", "Shoes"), 1, 0)
dat_1000$cosmetics <- ifelse(dat_1000$category == "Cosmetics", 1, 0)
dat_1000$food <- ifelse(dat_1000$category %in% c("Food", "Food & Beverage"), 1, 0)
dat_1000$technology <- ifelse(dat_1000$category == "Technology", 1, 0)

dat_1000$other_category <- ifelse(dat_1000$clothes_shoes + dat_1000$cosmetics + dat_1000$food + dat_1000$technology == 0, 1, 0)

all(dat_1000$clothes_shoes + dat_1000$cosmetics + dat_1000$food + dat_1000$technology + dat_1000$other_category == 1)

fraction_transactions <- c(
  "Clothes and Shoes" = mean(dat_1000$clothes_shoes),
  "Cosmetics" = mean(dat_1000$cosmetics),
  "Food" = mean(dat_1000$food),
  "Technology" = mean(dat_1000$technology),
  "Other" = mean(dat_1000$other_category)
)
print("Fraction of transactions in each category:")
print(round(fraction_transactions * 100, 2))

total_sales_1000 <- sum(dat_1000$price)
sales_clothes_shoes <- sum(dat_1000$price[dat_1000$clothes_shoes == 1])
sales_cosmetics <- sum(dat_1000$price[dat_1000$cosmetics == 1])
sales_food <- sum(dat_1000$price[dat_1000$food == 1])
sales_technology <- sum(dat_1000$price[dat_1000$technology == 1])
sales_other <- sum(dat_1000$price[dat_1000$other_category == 1])

```

```r
fraction_sales <- c(
  "Clothes and Shoes" = sales_clothes_shoes / total_sales_1000,
  "Cosmetics" = sales_cosmetics / total_sales_1000,
  "Food" = sales_food / total_sales_1000,
  "Technology" = sales_technology / total_sales_1000,
  "Other" = sales_other / total_sales_1000
)

print("Fraction of sales in each category:")
print(round(fraction_sales * 100, 2))

```

Solution (d).

To find the Maximum Likelihood Estimator (MLE) $\hat{\beta}$, we differentiate the log-likelihood with respect to β. Let $\phi(\cdot)$ denote the standard normal PDF. We use:

$$\frac{d}{dt}\log(\Phi(t)) = \frac{\phi(t)}{\Phi(t)}, \quad \text{and} \quad \frac{d}{dt}\log(1-\Phi(t)) = -\frac{\phi(t)}{1-\Phi(t)}.$$

For each element β_j of β, the derivative of the log-likelihood is:

$$\frac{\partial \ell(\beta; Z_n)}{\partial \beta_j} = \sum_{i=1}^{n}\left[y_i \frac{\phi(x_i'\beta)}{\Phi(x_i'\beta)} - (1-y_i)\frac{\phi(x_i'\beta)}{1-\Phi(x_i'\beta)}\right] x_{ij}.$$

Stacking all partial derivatives together, the gradient (score vector) is:

$$\nabla_\beta \ell(\beta; Z_n) = \sum_{i=1}^{n}\left[\frac{y_i - \Phi(x_i'\beta)}{\Phi(x_i'\beta)(1-\Phi(x_i'\beta))}\phi(x_i'\beta)\right] x_i.$$

Often written more simply as:

$$\nabla_\beta \ell(\beta; Z_n) = \sum_{i=1}^{n}\left[y_i \frac{\phi(x_i'\beta)}{\Phi(x_i'\beta)} - (1-y_i)\frac{\phi(x_i'\beta)}{1-\Phi(x_i'\beta)}\right] x_i.$$

Step 1: Characterizing the MLE $\hat{\beta}$

The MLE $\hat{\beta}$ sets the gradient to zero:

$$\nabla_\beta \ell(\hat{\beta}; Z_n) = 0.$$

Substituting back:

$$\sum_{i=1}^{n}\left[y_i \frac{\phi(x_i'\hat{\beta})}{\Phi(x_i'\hat{\beta})} - (1-y_i)\frac{\phi(x_i'\hat{\beta})}{1-\Phi(x_i'\hat{\beta})}\right] x_i = 0.$$

This is a system of k nonlinear equations in the k unknowns $\hat{\beta} = (\hat{\beta}_1, \ldots, \hat{\beta}_k)'$.

Step 2: No Closed-Form Solution

Unlike in linear regression or the logit model (even the logit doesn't have a closed form), the Probit model does not admit a closed-form solution for $\hat{\beta}$. The equation above must be solved using numerical optimization techniques such as the Newton-Raphson algorithm or other iterative methods.

Step 3: Numerical Optimization

A common iterative procedure is:

Algorithm 1: Newton-Raphson Method

Input: Initialize β_0, tolerence level $\varepsilon > 0$

for $m = 1$ **to** M **do**
  Given β^m, compute $\nabla_\beta \ell(\beta^{(m)}; Z_n)$ and $[H(\beta^{(m)}; Z_n)]$;
  Set $\beta^{(m+1)} = \beta^{(m)} - [H(\beta^{(m)}; Z_n)]^{-1} \nabla_\beta \ell(\beta^{(m)}; Z_n),$;
  if $\|\beta^{m+1} - \beta^m\| < \varepsilon$ **then**
    $\hat{\beta} = \beta^{m+1}$;
  else
    Proceed to the next iteration;
  end
end

where $H(\beta; Z_n)$ is the Hessian matrix of second derivatives evaluated at β. Convergence is achieved when changes in β or the norm of the gradient are below a given tolerance.
The regression result is as follows:

$$\hat{\beta} = \begin{bmatrix} \hat{\beta}_0 \\ \hat{\beta}_{price} \\ \hat{\beta}_{male} \\ \hat{\beta}_{age} \\ \hat{\beta}_{cosmetics} \\ \hat{\beta}_{food} \\ \hat{\beta}_{technology} \end{bmatrix} = \begin{bmatrix} 0.0682 \\ 0.000112 \\ -0.0502 \\ -0.00183 \\ -0.2879 \\ -0.1195 \\ 0.0640 \\ -0.4195 \end{bmatrix}.$$

Interpretation:

- The coefficient on price is positive but very small, suggesting a tiny positive association of price with the probability of cash payment (not statistically significant).
- male is negative, but not significant, suggesting no strong gender effect on the probability of cash usage.
- age coefficient is negative and small, not statistically significant either.
- Some category dummies (like Clothes/Shoes) are significantly different from zero, indicating that the reference category (likely "Other") differs in payment method probability.

Table 1: Optimization model

	Dependent variable:
	paid_in_cash
price	0.0001
	(0.0001)
male	−0.050
	(0.081)
age	−0.002
	(0.003)
clothes_shoes	−0.288**
	(0.130)
cosmetics	−0.120
	(0.133)
food	0.064
	(0.135)
technology	−0.420
	(0.314)
Constant	0.068
	(0.148)
Observations	1,000
Log Likelihood	−685.217
Akaike Inf. Crit.	1,386.434
Note:	*p<0.1; **p<0.05; ***p<0.01

```r
X <- as.matrix(cbind(1, dat_1000[, c("price", "male", "age", "clothes_shoes", "cosmetics", "food", "technology")]))
y <- dat_1000$paid_in_cash

neg_log_likelihood <- function(beta, X, y) {
  X_beta <- X %*% beta
  log_phi_Xb <- pnorm(X_beta, log.p = TRUE)
  log_phi_minus_Xb <- pnorm(-X_beta, log.p = TRUE)
  ll <- sum(y * log_phi_Xb + (1 - y) * log_phi_minus_Xb)
  return(-ll)
}

neg_log_likelihood_grad <- function(beta, X, y) {
  X_beta <- X %*% beta
  phi_Xb <- dnorm(X_beta)
  Phi_Xb <- pnorm(X_beta)
  Phi_minus_Xb <- pnorm(-X_beta)
  epsilon <- 1e-16
  Phi_Xb <- pmax(Phi_Xb, epsilon)
  Phi_minus_Xb <- pmax(Phi_minus_Xb, epsilon)
  gradient <- -t(X) %*% ((y * phi_Xb / Phi_Xb) - ((1 - y) * phi_Xb / Phi_minus_Xb))
  return(as.vector(gradient))
}

initial_beta <- rep(0, ncol(X))

result <- optim(par = initial_beta, fn = neg_log_likelihood, gr = neg_log_likelihood_grad, X = X, y = y, method = "BFGS")

if (result$convergence == 0) {
  cat("Optimization converged.\n")
} else {
  cat("Optimization did not converge.\n")
}

beta_hat <- result$par
print("Estimated coefficients (beta_hat):")
print(beta_hat)

### Programming Method
model <- glm(paid_in_cash ~ price + male + age + clothes_shoes + cosmetics + food + technology,
    data = dat_1000, family = binomial(link = "probit"))
stargazer(model, type = "latex", title = "Optimization model", out = "d.tex")

beta_hat2 <- coef(model)
```

```r
print("Estimated coefficients (beta_hat):")
print(beta_hat2)

```

Solution (e).

We define

$$\gamma_1(\beta) = \Phi(x_2'\beta) - \Phi(x_1'\beta),$$

where x_1' is a vector for a 30-year-old male buying Clothes/Shoes for 500 TRY, and x_2' is the same vector but with age increased to 60 years old. Only the age element of x_i changes from 30 to 60.

For our estimated $\hat{\beta}$,

$$\gamma_1(\hat{\beta}) \approx -0.02096.$$

This suggests that increasing age from 30 to 60 reduces the probability of cash payment by about 2.1 percentage points for this specific profile.

For $\gamma_2(\beta)$, we do not condition on category. We take a weighted average of the partial effects across the five categories, with weights given by their share in total sales:

$$\gamma_2(\hat{\beta}) = \sum_j w_j \left[\Phi(x_{2,j}'\beta) - \Phi(x_{1,j}'\beta)\right],$$

where w_j is the sales fraction for category j.

Empirically,

$$\gamma_2(\hat{\beta}) \approx -0.02077,$$

very close to $\gamma_1(\hat{\beta})$, indicating a similar overall effect once categories are averaged by their sales importance.

```r
x_age_30 <- c(1, 500, 1, 30, 1, 0, 0, 0)

x_age_60 <- x_age_30
x_age_60[4] <- 60  # Update age to 35

prob_age_30 <- pnorm(sum(x_age_30 * beta_hat))
prob_age_60 <- pnorm(sum(x_age_60 * beta_hat))

gamma_1 <- prob_age_60 - prob_age_30
print(paste("Gamma_1 (effect of age increasing by 5 years):", gamma_1))

gamma_c <- numeric(length(fraction_sales))
names(gamma_c) <- names(fraction_sales)

for (cat in names(fraction_sales)) {
  clothes_shoes <- ifelse(cat == "Clothes and Shoes", 1, 0)
  cosmetics <- ifelse(cat == "Cosmetics", 1, 0)
  food <- ifelse(cat == "Food", 1, 0)
```

```r
  technology <- ifelse(cat == "Technology", 1, 0)

  x_age_30_2 <- c(1, 500, 1, 30, clothes_shoes, cosmetics, food,
    technology)
  x_age_60_2 <- x_age_30_2
  x_age_60_2[4] <- 60  # Update age to 35

  prob_age_30_2 <- pnorm(sum(x_age_30_2 * beta_hat))
  prob_age_60_2 <- pnorm(sum(x_age_60_2 * beta_hat))

  gamma_c[cat] <- prob_age_60_2 - prob_age_30_2
}

gamma_2 <- sum(fraction_sales * gamma_c)
print(paste("Gamma_2 (weighted effect over categories):", gamma_2))

```

Solution (f).

Consider the linear model

$$y_i = x_i'\beta + u_i$$

where $u_i \mid x_i \sim N(0,1)$.

Step 1: Define the Objective Function

Define $\mathcal{B} = \{\beta \in \mathbb{R} : \|\beta\| \leq c\}$ for some very large c.

The objective function is given by:

$$\hat{\beta} = \arg\min_{\beta\in\mathcal{B}} Q_n(\beta) = \arg\min_{\beta\in\mathcal{B}} \frac{1}{n}\sum_{i=1}^{n}(y_i - x_i'\beta)^2$$

Step 2: Express the Limiting Behavior of the Objective Function

The expected form of the objective function $Q_n(\beta)$ is:

$$\begin{aligned}
Q(\beta) &= E[(y_i - x_i'\beta)^2] \\
&= E[(x_i'\beta_0 + u_i - x_i'\beta)^2] \\
&= E[(x_i'(\beta_0 - \beta) + u_i)^2] \\
&= E[(x_i'(\beta_0 - \beta))^2] + 2E[x_i'(\beta_0 - \beta)u_i] + E[u_i^2] \\
&= E[(x_i'(\beta_0 - \beta))^2] + \sigma^2 \\
&= (\beta_0 - \beta)'E[x_i x_i'](\beta_0 - \beta) + 1
\end{aligned}$$

since $E[u_i \mid x_i] = 0$ and $\sigma^2 = 1$.

Step 3: Show Minimization of $Q(\beta)$ at β_0

As $E[x_i x_i']$ is positive definite, the function $Q(\beta)$ is minimized at β_0 because:

$$(\beta_0 - \beta)'E[x_i x_i'](\beta_0 - \beta) \geq 0$$

which is zero if and only if $\beta = \beta_0$.

Step 4: Prove Uniform Convergence of $Q_n(\beta)$ to $Q(\beta)$

It's obvious that $m(x_i, y_i, \beta) = (y_i - x_i'\beta)^2$ satisfies the first three conditions of the Uniform Law of Large Numbers(ULLN).

We then prove the fourth one:

$$E\left[\sup_{\beta\in\mathcal{B}} \|m(x_i, y_i, \beta)\|\right] \leq E\left[|y_i|^2\right] + \sup_{\beta\in\mathcal{B}} 2E\left[|y_i|\,|x_i|\,\|\beta\|\right] + \sup_{\beta\in\mathcal{B}} E\left[|x_i|^2\left\|\beta^2\right\|\right]$$
$$< \infty$$

By ULLN:

$$Q_n(\beta) \xrightarrow{p} Q(\beta), n \to \infty$$

Step 5: Demonstrate the Consistency of $\hat{\beta}_n$

According to extremum estimator theory, if $Q_n(\beta)$ converges uniformly to $Q(\beta)$ and $Q(\beta)$ has a unique global minimum at β_0, then:

$$\hat{\beta}_n \xrightarrow{p} \beta_0$$

Consistency of Probit Estimator:

First, we define $f(w_i;\theta) = \Phi(x_i'\beta)^{y_t}\Phi(-x_i'\beta)^{1-y_t}$, then $\log f(w_i;\theta) = y_i \log\Phi(x_i'\beta) + (1 - y_i)\log\Phi(-x_i'\beta)$.

Step 1: Theorem of Consistency

Theorem 1. If $Q_n(w_i;\theta)$ is a function of w_i and θ such that:

(A) Parameter space $\Theta \in \mathbb{R}^k$ is compact, $\theta_0 \in \Theta$;

(B) $Q_n(w_i;\theta)$ is continuous in $\theta \in \Theta$ for all w_i.

(C) $Q_n(\theta)$ converges in probability to $Q(\theta)$ uniformly in $\theta \in \Theta$, and $Q(\theta)$ has a unique global minimum at θ_0.

Define $Q_n(\hat{\theta}_n) = \max_{\theta\in\Theta} Q_n(\theta)$.

Then, $\hat{\theta}_n \xrightarrow{p} \theta_0$.

Proof for Theorem 1.

Let N be a neignbourhood in $\mathbb{R}^k$ containing θ_0. Then $\overline{N} \cap \Theta$ is compact $\Rightarrow \max_{\theta\in\overline{N}\cap\Theta} Q(\theta)$ exists.

Denote $\varepsilon = Q(\theta_0) - \max_{\theta\in\overline{N}\cap\Theta} Q(\theta)$.

Define incident A_n as:

$$A_n : \left|\frac{1}{n}Q_n(\theta) - Q(\theta)\right| < \frac{\varepsilon}{2}.$$

This implies that:

$$\begin{cases} Q(\hat{\theta}_n) > \frac{1}{n} Q_n(\hat{\theta}_n) - \frac{\varepsilon}{2} \\ \frac{1}{n} Q_n(\hat{\theta}_n) > Q(\theta_0) - \frac{\varepsilon}{2} \end{cases}$$

But, as $Q_N(\hat{\theta}_n) \geq Q_n(\theta_0)$,

$$Q(\hat{\theta}_n) > Q(\theta_0) - \varepsilon \Rightarrow \hat{\theta}_n \in N.$$

Thus, $\mathbb{P}[A_n] \leq \mathbb{P}[\hat{\theta}_n \in N]$. Since we have $\lim_{n\to\infty} \mathbb{P}[A_n] = 1$ by (C), we have $\mathbb{P}[\hat{\theta}_n \in N] \to 1$. Hence $\hat{\theta}_n \xrightarrow{p} \theta_0$. □

In our cese, we take the parameter space $\mathcal{B} = \left\{\beta \in \mathbb{R}^k : \|\beta\| < c\right\}$ for some large c, then, we have our compact parameter space Θ, (A) is satisfied.
As we take $Q_n(\theta) = \frac{1}{n}\ell(\beta; Z_n) = \frac{1}{n}\sum_{i=1}^n \log f(w_i; \beta)$, it's continuous in θ for all w_i, (B) is satisfied.
So, we need to prove two conditions for the theorem to hold:

1. $Q_n(\theta)$ converges in probability to $Q(\theta)$ uniformly in $\theta \in \Theta$.
2. (Identification) $Q(\theta)$ has a unique global minimum at θ_0.

Step 2: Identification of Probit Model

Definition 1 (Identification)**.** The information matrix $I(\theta)$ is defined as:

$$I(\theta) = E\left[\frac{\partial \ell^2(w_i;\theta)}{\partial\theta\partial\theta'}\right].$$

If $I(\theta)$ is positive definite, then θ is identified.
If θ is identified, it means that if $\theta \neq \theta_0$, then $f(w_i;\theta) \neq f(w_i;\theta_0)$.

Lemma 1 (Information Inequality)**.**
If θ is identified, and $\mathbb{E}\left[\log f(w_i;\theta) < \infty\right]$ for all θ, then $Q(\theta) = \mathbb{E}\left[|\log f(w_i;\theta)|\right]$ has a unique maximum at θ_0.

Proof for Lemma 1. For a random variable Y, by Jensen's inequality, we have:

$$-\log \mathbb{E}[Y] < \mathbb{E}\left[-\log Y\right].$$

In our case, we define a new random variable for $\theta \neq \theta_0$:

$$Y = \frac{f(w_i;\theta)}{\mathbb{E}[f(w_i;\theta_0)]}.$$

Then, we have:

$$\begin{aligned} Q(\theta_0) - Q(\theta) &= \mathbb{E}\left[-\log Y\right] > -\log \mathbb{E}[Y] \\ &= -\log \int f(w_i;\theta) d\theta = 0 \end{aligned}$$

□

Step 3: Prove uniform convergence of Probit Model

To prove the uniform convergence of the Probit model, we give the second theorem:

Theorem 2 (Uniform Law of Large Numbers)**.**
If x_i(or say, data) are i.i.d, and $\log f(w_i;\theta)$ is a function of x_i, y_i, θ such that:

(a) Parameter space $\Theta \in \mathbb{R}^k$ is compact, $\theta_0 \in \Theta$;

(b) $m(w_i;\theta)$ is continuous at each $\theta \in \Theta$ with probability to 1;

(c) There exist a dominant function $d(w_i)$ such that $\|\log f(w_i;\theta)\| \leq \|d(w_i)\|\ \forall \theta \in \Theta$;

(d) $E\left[d(w_i)\right] < \infty$.

then, $\mathbb{E}\left[\log f(w_i;\theta)\right]$ is continuous and

$$\sup_{\theta\in\Theta} \left\| \frac{1}{n}\sum_{i=1}^{n} \log f(w_i;\theta) - \mathbb{E}\left[\log f(w_i;\theta)\right] \right\| \xrightarrow{p} 0.$$

Proof for Theorem 2.
For $\forall \theta_0 \in \Theta$, we define $\mathcal{B}(\theta_0, \delta) = \{\theta \in \Theta : \|\theta - \theta_0\| < \delta\}$ and

$$\Delta_\delta(w_i;\theta) = \sup_{\theta\in\mathcal{B}(\theta_0,\delta)} \left(\log f(w_i;\theta) - \mathbb{E}\left[\log f(w_i;\theta)\right]\right).$$

For $\theta_0 \in \Theta$, we have:

$$\mathbb{E}\left[\Delta_\delta(w_i;\theta) \to 0\right] \text{ as } \delta \to 0.$$

because

1. $\Delta_\delta(w_i;\theta_0) \to \log f(w_i;\theta) - \mathbb{E}\left[\log f(w_i;\theta)\right]$ almost surely as $\delta \to 0$, because:

$$\mathbb{P}\left[\lim_{\delta\to 0} \sup_{\theta\in\mathcal{B}(\theta_0,\delta)} \log f(w_i;\theta) = \log f(w_i;\theta_0)\right] = 1$$

   by condition (b) and that $\mathbb{E}\left[\log f(w_i;\theta)\right]$ is continuous at θ_0.

2. By condition (c) and (d), we have:

$$\Delta_\delta(w_i;\theta_0) \leq 2 \sup_{\theta\in\mathcal{B}(\theta_0,\delta)} |\log f(w_i;\theta)| \leq 2d(w_i)$$

So, for all $\theta \in \Theta$, $\varepsilon > 0$, $\exists \delta_\varepsilon(\theta)$, such that

$$\mathbb{E}\left[\Delta_{\delta_\varepsilon(\theta)}(w_i;\theta)\right] < \varepsilon.$$

Obviously, we can cover the entire parameter space with a finite number of $\mathcal{B}(\theta, \delta_\varepsilon(\theta)) : \theta \in \Theta$, which is:

$$\mathcal{B}\left(\theta_k, \delta_\varepsilon(\theta_k) : k = 1, 2, \cdots, K\right) \text{ s.t. } \Theta = \bigcup_{k=1}^{K} \mathcal{B}\left(\theta_k, \delta_\varepsilon(\theta_k)\right).$$

Note that:

$$\begin{aligned}
&\sup_{\theta\in\Theta}\left[\frac{1}{n}\sum_{i=1}^{n}\log f(w_i;\theta) - \mathbb{E}\left[\log f(w_i;\theta)\right]\right] \\
&= \max_k \sup_{\theta\in\mathcal{B}(\theta_k,\delta_\varepsilon(\theta_k))}\left[\frac{1}{n}\sum_{i=1}^{n}\log f(w_i;\theta) - \mathbb{E}\left[\log f(w_i;\theta)\right]\right] \\
&\leq \max_k \frac{1}{n}\left[\sum_{i=1}^{n}\sup_{\theta\in\mathcal{B}(\theta_k,\delta_\varepsilon(\theta_k))}\log f(w_i;\theta) - \mathbb{E}\left[\log f(w_i;\theta)\right]\right] \\
&\leq \mathcal{O}_p(1) + \max_k \mathbb{E}^* \Delta_{\delta_\varepsilon(\theta_k)}(w_i;\theta_k) \\
&= \mathcal{O}_p(1) + \varepsilon.
\end{aligned}$$

where the second inequality holds by the Weak Law of Large Numbers (WLLN) because:

$$\left|\sup_{\theta\in\mathcal{B}(\theta_k,\delta_\varepsilon(\theta_k))}\log f(w_i;\theta)\right| \leq d(w_i); \mathbb{E}[d(w_i)] < \infty.$$

and the third inequality holds by the definition of $\delta_\varepsilon(\theta_k)$. By analogous argument, we can prove that:

$$\inf_{\theta\in\Theta}\left[\frac{1}{n}\sum_{i=1}^{n}\log f(w_i;\theta) - \mathbb{E}\left[\log f(w_i;\theta)\right]\right] \geq O_p(1) - \varepsilon.$$

Combing the two results, we have:

$$\left|\frac{1}{n}\sum_{i=1}^{n}\log f(w_i;\theta) - \mathbb{E}[\log f(w_i;\theta)]\right| \to \mathcal{O}_p(1) = 0.$$

□

Finishing the proof of Theorem 2, we can find that the Probit model still have to satisfy conditions (c) and (d) to hold the theorem.

Step 4: Proof of Conditions (c) and (d) for ULLN

In this part, we show that identification and the uniform convergence of the Probit model are combined by the existence of $\mathbb{E}[x_i x_i']$ and its nonsingularity.

Proof for ULLN conditions (c) and (d).
For this proof, we take two steps:
Step 1: $\mathbb{E}\left[|\log f(w_i;\theta)|\right]$ is finite.
Let $\theta \neq \theta_0$, then

$$\begin{aligned}\mathbb{E}\left[(x_i'(\theta-\theta_0))^2\right] &= (\theta-\theta_0)'\mathbb{E}[x_i x_i'](\theta-\theta_0) > 0 \\ &\Rightarrow x_i'(\theta-\theta_0) \neq 0 \\ &\Rightarrow x_i'\theta \neq x_i'\theta_0\end{aligned}$$

Since Φ is strictly monotone, this gives us $\Phi(x_i'\theta) \neq \Phi(-x_i'\theta)$. So that $f(w_i;\theta) = \Phi(x_i'\beta)^{y_t}\Phi(-x_i'\beta)^{1-y_t} \neq f(w_i;\theta_0)$.
We know that $\frac{d\log\Phi(v)}{dx} = \frac{\phi(v)}{\Phi(v)}$ is convex and asymptotic to 0 as $v \to \infty$ and to $-v$ as $v \to -\infty$.
We take the mean-value expansion around $\theta = 0$:

$$\begin{aligned}|\log\Phi(x_i'\theta)| &= \left|\log\Phi(0) + \lambda(x_i'\tilde{\theta})x_i'\theta\right| \\ &\leq |\log\Phi(0)| + \left|\lambda(x_i'\tilde{\theta})x_i'\theta\right| \\ &\leq |\log\Phi(0)| + C\left(1 + \left|x_i'\tilde{\theta}\right|\right)|x_i'\theta| \\ &\leq |\log\Phi(0)| + C\left(1 + \|x_i\|\,\|\theta\|\right)\|x_i\|\,\|\theta\|\end{aligned}$$

where λ is the reverse Mills ratio.
Since $1-\Phi(v) = \Phi(-v)$ and y are bounded, we have:

$$|\log f(w_i;\theta)| \leq |\log\Phi(0)| + C\left(1 + \|x_i\|\,\|\theta\|\right)\|x_i\|\,\|\theta\|$$

where C is a constant.
Thus, we could say that $\mathbb{E}\left[|\log f(w_i;\theta)|\right]$ is finite.
Step 2: $\mathbb{E}[d(w_i)]$ exist, and is finite.
Based on Step 1, we could directly take

$$d(w_i) = C\left(1 + \|x_i\|^2\right).$$

It's obvious that $\mathbb{E}[d(w_i)]$ is finite. □

Combining Lemma 1, Lemma 2, Theorem 1, and Theorem 2, we could say that the Probit model estimator is consistent.

Solution (g).

```
M <- 100
n <- nrow(dat_1000)
beta_age_bootstrap <- numeric(M)
```

```r
gamma_1_bootstrap <- numeric(M)
gamma_2_bootstrap <- numeric(M)
set.seed(2024)

for (m in 1:M) {
  indices <- sample(1:n, size = n, replace = TRUE)
  dat_bootstrap <- dat_1000[indices, ]

  model_boot <- glm(paid_in_cash ~ price + male + age + clothes_shoes + cosmetics + food + technology, data = dat_bootstrap, family = binomial(link = "probit"))
  beta_hat_boot <- coef(model_boot)

  beta_age_bootstrap[m] <- beta_hat_boot["age"]

  x_age_30 <- c(1, 500, 1, 30, 1, 0, 0, 0)
  x_age_60 <- x_age_30
  x_age_60[4] <- 60  # Update age to 35
  prob_age_30 <- pnorm(sum(x_age_30 * beta_hat_boot))
  prob_age_60 <- pnorm(sum(x_age_60 * beta_hat_boot))
  gamma_1_bootstrap[m] <- prob_age_60 - prob_age_30

  gamma_c <- numeric(length(fraction_sales))
  names(gamma_c) <- names(fraction_sales)

  for (cat in names(fraction_sales)) {
    clothes_shoes <- ifelse(cat == "Clothes and Shoes", 1, 0)
    cosmetics <- ifelse(cat == "Cosmetics", 1, 0)
    food <- ifelse(cat == "Food", 1, 0)
    technology <- ifelse(cat == "Technology", 1, 0)

    x_age_30_2 <- c(1, 500, 1, 30, clothes_shoes, cosmetics, food, technology)
    x_age_60_2 <- x_age_30
    x_age_60_2[4] <- 60  # Update age to 35

    prob_age_30_2 <- pnorm(sum(x_age_30_2 * beta_hat_boot))
    prob_age_60_2 <- pnorm(sum(x_age_60_2 * beta_hat_boot))
    gamma_c[cat] <- prob_age_60_2 - prob_age_30_2
  }

  gamma_2_bootstrap[m] <- sum(fraction_sales * gamma_c)
}

hist(beta_age_bootstrap, main = "Bootstrap Distribution of Coefficient on Age", xlab = "Coefficient on Age", breaks = 20)
hist(gamma_1_bootstrap, main = "Bootstrap Distribution of Gamma_1", xlab = "Gamma_1", breaks = 20)
```

```
47 hist(gamma_2_bootstrap, main = "Bootstrap Distribution of Gamma_2", xlab
       = "Gamma_2", breaks = 20)
48
```

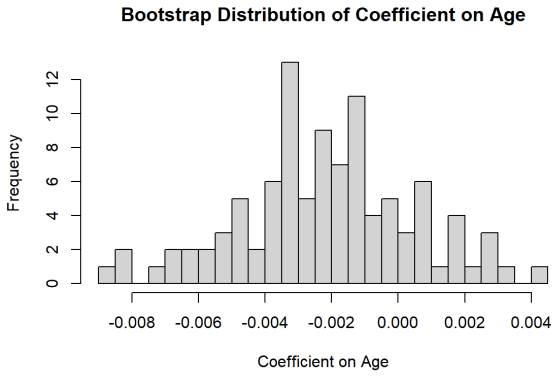

Figure 1: Bootstrap Distribution of Coefficient on Age

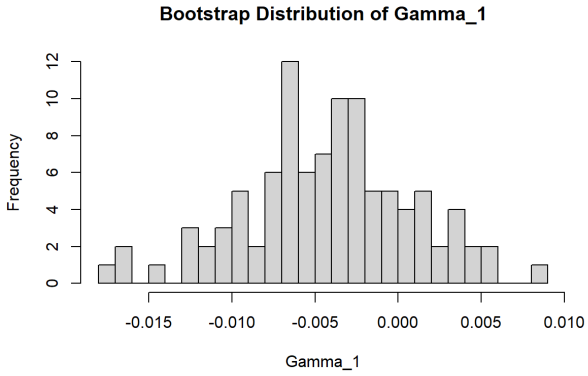

Figure 2: Bootstrap Distribution of Gamma_1

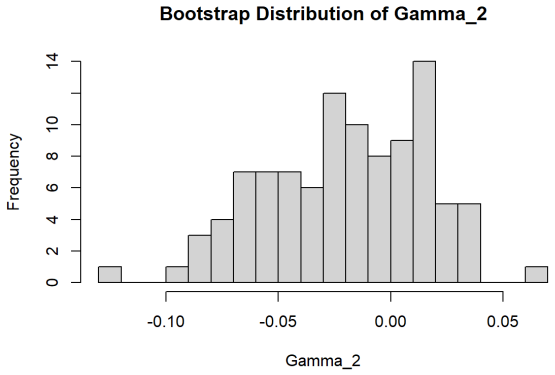

Figure 3: Bootstrap Distribution of Gamma_2

Solution (h).
For a sample $\{(y_i, x_i)\}_{i=1}^n$, the log-likelihood function is:

$$\ell(\beta; Z_n) = \sum_{i=1}^{n} \left[y_i \log(\Phi(x_i'\beta)) + (1 - y_i) \log(1 - \Phi(x_i'\beta)) \right].$$

To derive the Hessian, focus on a single observation i and then we will take expectations:

$$\ell_i(\beta) = y_i \log(\Phi(t_i)) + (1 - y_i) \log(1 - \Phi(t_i)), \quad t_i = x_i'\beta.$$

First Derivative (Score):
First, take the derivative with respect to t_i:

$$\frac{\partial \ell_i(\beta)}{\partial t_i} = y_i \frac{\phi(t_i)}{\Phi(t_i)} - (1 - y_i) \frac{\phi(t_i)}{1 - \Phi(t_i)},$$

where $\phi(\cdot)$ is the standard normal PDF.
Combine the fractions:

$$\frac{\partial \ell_i(\beta)}{\partial t_i} = \phi(t_i) \left[\frac{y_i}{\Phi(t_i)} - \frac{1 - y_i}{1 - \Phi(t_i)} \right].$$

Find a common denominator $\Phi(t_i)(1 - \Phi(t_i))$:

$$\frac{\partial \ell_i(\beta)}{\partial t_i} = \frac{\phi(t_i)}{\Phi(t_i)(1 - \Phi(t_i))} (y_i - \Phi(t_i)).$$

Since $y_i - \Phi(t_i) = y_i - p_i$, we have:

$$\frac{\partial \ell_i(\beta)}{\partial t_i} = \frac{\phi(t_i)}{p_i(1 - p_i)} (y_i - p_i).$$

To get the gradient w.r.t. β, use chain rule:

$$\frac{\partial \ell_i(\beta)}{\partial \beta} = \frac{\partial \ell_i(\beta)}{\partial t_i} x_i = \frac{\phi(t_i)}{p_i(1 - p_i)} (y_i - p_i) x_i.$$

This is the score vector for a single observation.
Second Derivative (Hessian):
Now differentiate again with respect to β:

$$\frac{\partial^2 \ell_i(\beta)}{\partial \beta \partial \beta'} = \frac{\partial}{\partial \beta} \left(\frac{\phi(t_i)}{p_i(1 - p_i)} (y_i - p_i) x_i \right).$$

Since $t_i = x_i'\beta$, $\frac{\partial t_i}{\partial \beta} = x_i$. Thus, second derivatives w.r.t. β come through differentiating

w.r.t. t_i, then applying chain rule again:

$$\frac{\partial^2 \ell_i(\beta)}{\partial \beta \partial \beta'} = \left(\frac{\partial^2 \ell_i(\beta)}{\partial t_i^2} \right) x_i x_i'.$$

So the main task is to find:

$$\frac{\partial^2 \ell_i(\beta)}{\partial t_i^2}.$$

We have:

$$\frac{\partial \ell_i(\beta)}{\partial t_i} = \frac{\phi(t_i)}{p_i(1-p_i)}(y_i - p_i).$$

Take the derivative w.r.t. t_i:

$$\frac{\partial^2 \ell_i(\beta)}{\partial t_i^2} = \frac{\partial}{\partial t_i} \left[\frac{\phi(t_i)}{p_i(1-p_i)}(y_i - p_i) \right].$$

This involves the product rule and quotient rule. However, *the key simplification occurs when we take expectations at the true parameter* β_0. Under the true model, $E[y_i] = p_i$, so $E[y_i - p_i] = 0$. Terms involving $(y_i - p_i)$ vanish when taking expectation.
At the true parameter, the Fisher information (which is $-E[\partial^2 \ell_i(\beta_0)/\partial\beta\partial\beta']$) simplifies dramatically. Instead of going through the full complex algebra of the second derivative in the y_i form, we use the known result from standard Probit derivations:
Under correct specification, the expected Hessian w.r.t. t_i at β_0 is known to be:

$$E\left[\frac{\partial^2 \ell_i(\beta_0)}{\partial t_i^2} \right] = -\frac{\phi(t_i)^2}{p_i(1-p_i)}.$$

Thus:

$$E\left[\frac{\partial^2 \ell_i(\beta_0)}{\partial \beta \partial \beta'} \right] = E\left[\frac{\partial^2 \ell_i(\beta_0)}{\partial t_i^2} x_i x_i' \right] = E\left[-\frac{\phi(t_i)^2}{p_i(1-p_i)} x_i x_i' \right].$$

Multiplying by -1, the Fisher Information matrix (which is H in the problem) is:

$$H = E\left[\frac{\phi(x_i'\beta_0)^2}{\Phi(x_i'\beta_0)[1 - \Phi(x_i'\beta_0)]} x_i x_i' \right].$$

Since $1 - \Phi(t_i) = \Phi(-t_i)$:

$$H = E\left[\frac{\phi(x_i'\beta_0)^2}{\Phi(x_i'\beta_0)\Phi(-x_i'\beta_0)} x_i x_i' \right].$$

In the dataset result, the given histograms for the bootstrap distributions and the asymptotic distributions show approximately symmetric, bell-shaped distributions. This suggests that the normal approximation may be reasonable.

```
X <- as.matrix(cbind(1, dat_1000[, c("price", "male", "age", "clothes_
    shoes", "cosmetics", "food", "technology")]))
X_beta_hat <- X %*% beta_hat
```

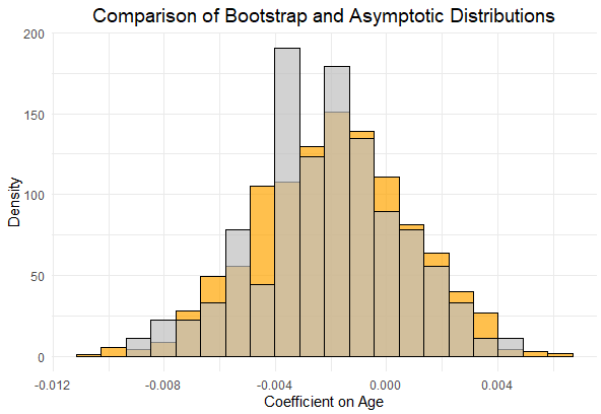

Figure 4: Comparison of Bootstrap and Asymptotic Distributions

```
3
4  log_phi_Xb <- dnorm(X_beta_hat, log = TRUE)
5  log_Phi_Xb <- pnorm(X_beta_hat, log.p = TRUE)
6
7  log_Phi_minus_Xb <- pnorm(-X_beta_hat, log.p = TRUE)
8  log_factor <- 2 * log_phi_Xb - log_Phi_Xb - log_Phi_minus_Xb
9  factor <- exp(log_factor)
10 factor[!is.finite(factor)] <- 0
11 factor <- as.vector(factor)
12 X_weighted <- sweep(X, 1, factor, FUN = "*")
13
14 H_hat <- t(X) %*% X_weighted / nrow(X)
15 V_hat <- solve(H_hat)
16
17 variance_beta_age <- V_hat[4, 4]
18 beta_age_sd <- sqrt(variance_beta_age / nrow(X))
19
20 if (!is.finite(beta_age_sd)) {
21   stop("Standard error for the coefficient on age is not finite.")
22 }
23
24 mean_age <- beta_hat[4]
25 simulated_draws <- rnorm(1000, mean = mean_age, sd = beta_age_sd)
26
27 library(ggplot2)
28
29 bootstrap_data <- data.frame(Distribution = "Bootstrap", Values = beta_
     age_bootstrap)
30 asymptotic_data <- data.frame(Distribution = "Asymptotic", Values =
     simulated_draws)
31 combined_data <- rbind(bootstrap_data, asymptotic_data)
32
33 ggplot(combined_data, aes(x = Values, fill = Distribution)) +
34   geom_histogram(aes(y = ..density..),
```

```
                bins = 20, alpha = 1, position = "identity", color = "
  black") +
scale_fill_manual(values = c("Bootstrap" = "grey", "Asymptotic" = "red
  ")) +
labs(title = "Comparison of Bootstrap and Asymptotic Distributions",
     x = "Coefficient on Age", y = "Density") +
theme_minimal() +
theme(plot.title = element_text(hjust = 0.5, size = 14),
      legend.title = element_blank(),
      legend.position = "topright") +
guides(fill = guide_legend(reverse = TRUE))

```

Solution (i).
The Delta Method states that if

$$\sqrt{n}(\hat{\beta} - \beta_0) \xrightarrow{d} N(0, H^{-1}),$$

and $g(\cdot)$ is a continuously differentiable function at β_0, then

$$\sqrt{n}(g(\hat{\beta}) - g(\beta_0)) \xrightarrow{d} N(0, \nabla_\beta g(\beta_0)' H^{-1} \nabla_\beta g(\beta_0)).$$

In our case, $g(\beta) = \gamma_1(\beta)$.
Computing the Gradient $\nabla_\beta \gamma_1(\beta)$:
We have:

$$\gamma_1(\beta) = \Phi(x_2'\beta) - \Phi(x_1'\beta).$$

The gradient with respect to β is:

$$\nabla_\beta \gamma_1(\beta) = \frac{\partial}{\partial \beta} \left[\Phi(x_2'\beta)\right] - \frac{\partial}{\partial \beta} \left[\Phi(x_1'\beta)\right].$$

Since $\frac{d}{dt}\Phi(t) = \phi(t)$, we get:

$$\nabla_\beta \gamma_1(\beta) = \phi(x_2'\beta)x_2 - \phi(x_1'\beta)x_1.$$

Asymptotic Distribution of $\gamma_1(\hat{\beta})$:
Applying the Delta Method at β_0:

$$\sqrt{n}(\gamma_1(\hat{\beta}) - \gamma_1(\beta_0)) \xrightarrow{d} N(0, \nabla_\beta \gamma_1(\beta_0)' H^{-1} \nabla_\beta \gamma_1(\beta_0)).$$

In finite samples, we replace β_0 with $\hat{\beta}$, and H with its estimator $\hat{H}$, thus:

$$\gamma_1(\hat{\beta}) \overset{approx}{\sim} N\left(\gamma_1(\hat{\beta}), \frac{1}{n} \nabla_\beta \gamma_1(\hat{\beta})' \hat{H}^{-1} \nabla_\beta \gamma_1(\hat{\beta})\right),$$

where $\hat{H}$ and $\nabla_\beta \gamma_1(\hat{\beta})$ are computed from the sample and the estimated parameters. This gives us an asymptotic approximation to the finite sample distribution of $\gamma_1(\hat{\beta})$.
To summarize, the asymptotic variance of $\gamma_1(\hat{\beta})$ is:

$$\widehat{\text{Var}}(\gamma_1(\hat{\beta})) = \frac{1}{n} \nabla_\beta \gamma_1(\hat{\beta})' \hat{H}^{-1} \nabla_\beta \gamma_1(\hat{\beta}).$$

Empirical Implementation:

1. Estimate $\hat{\beta}$ using the Probit model.
2. Compute $\nabla_\beta \gamma_1(\hat{\beta}) = \phi(x_2'\hat{\beta})x_2 - \phi(x_1'\hat{\beta})x_1$.
3. Compute $\hat{H}^{-1}$ (the inverse of the estimated H matrix).
4. Approximate:

$$\gamma_1(\hat{\beta}) \sim N\left(\gamma_1(\hat{\beta}), \frac{1}{n} \nabla_\beta \gamma_1(\hat{\beta})' \hat{H}^{-1} \nabla_\beta \gamma_1(\hat{\beta})\right).$$

The comparison shows that both approximations yield similar conclusions: the effect is not statistically significant, and the distributions are roughly symmetric.

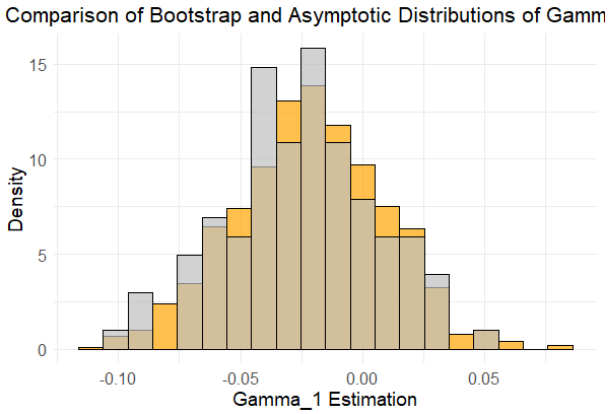

Figure 5: Comparison of Bootstrap and Asymptotic Distributions of Gamma_1

```r
phi_age_60 <- dnorm(sum(x_age_60 * beta_hat))
phi_age_30 <- dnorm(sum(x_age_30 * beta_hat))
grad_g <- phi_age_60 * x_age_60 - phi_age_30 * x_age_30

print(grad_g)

# Compute asymptotic variance
var_gamma_1 <- t(grad_g) %*% V_hat %*% grad_g / nrow(dat_1000)
gamma_1_sd <- sqrt(var_gamma_1)

simulated_gamma <- rnorm(1000, mean = gamma_1, sd = gamma_1_sd)

bootstrap_data2 <- data.frame(Value = gamma_1_bootstrap, Distribution =
    "Bootstrap")
simulated_data2 <- data.frame(Value = simulated_gamma, Distribution = "
    Asymptotic")
```

```r

# Combine data
combined_data2 <- rbind(bootstrap_data2, simulated_data2)

# Create the plot
ggplot(combined_data2, aes(x = Value, fill = Distribution)) +
  geom_histogram(aes(y = ..density..), bins = 20, position = "identity",
    alpha = 0.7, color = "black") +
  scale_fill_manual(values = c("Bootstrap" = "grey", "Asymptotic" = "
   orange")) +
  labs(title = "Comparison of Bootstrap and Asymptotic Distributions of
   Gamma_1",
       x = "Gamma_1 Estimation", y = "Density") +
  theme_minimal() +
  theme(plot.title = element_text(hjust = 0.5, size = 16),
        legend.title = element_blank(),
        legend.position = "topright",
        axis.title = element_text(size = 14),
        axis.text = element_text(size = 12)) +
  guides(fill = guide_legend(reverse = TRUE))

```

Solution (j).

We test

$$H_0 : \gamma_1(\beta) = 0 \quad \text{vs.} \quad H_1 : \gamma_1(\beta) \neq 0.$$

Under the asymptotic approximation,

$$\gamma_1(\hat{\beta}) \approx N\left(\gamma_1(\beta_0), \frac{1}{n}\hat{V}\right),$$

where

$$\hat{V} = \nabla_\beta \gamma_1(\hat{\beta})' \hat{H}^{-1} \nabla_\beta \gamma_1(\hat{\beta}).$$

The t-statistic is:

$$t = \frac{\gamma_1(\hat{\beta})}{\sqrt{\hat{V}/n}}.$$

Empirically, $t \approx -0.68$.

A 95% confidence interval for $\gamma_1(\beta)$ is:

$$\left[\gamma_1(\hat{\beta}) - 1.96\sqrt{\frac{\hat{V}}{n}}, \gamma_1(\hat{\beta}) + 1.96\sqrt{\frac{\hat{V}}{n}}\right].$$

Empirically, the 95% Confidence Interval for $\gamma_1(\beta)$ is: -0.0815 to 0.0396, covering 0.

So, we conclude that the expected probabilities of cash payment for a 30 year-old and a 60 year-old male buying clothes for 500 TRY are not significantly different.

```r
t_statistic <- gamma_1 / gamma_1_sd
critical_value <- qnorm(0.975)  # 1.96 for 95% confidence

if (abs(t_statistic) > critical_value) {
  conclusion <- "Reject the null hypothesis."
} else {
  conclusion <- "Fail to reject the null hypothesis."
}

print(paste("t-statistic:", round(t_statistic, 2)))
print(paste("Conclusion:", conclusion))

lower_bound <- gamma_1 - critical_value * gamma_1_sd
upper_bound <- gamma_1 + critical_value * gamma_1_sd

print(paste("95% Confidence Interval for gamma_1:", round(lower_bound,
    4), "to", round(upper_bound, 4)))

```