

Lecture Notes: Econometrics II

Based on lectures by [Marko Mlikota](#) in Spring semester, 2025

Draft updated on May 22, 2025

This is the lecture note taken in the course *Econometrics II* taught by [Marko Mlikota](#) at Graduate Institute of International and Development Studies, Geneva as part of the International Economics program (Semester II, 2024). The content is partly based on the course notes provided by the professor and supplemented by many other references I read myself. The main reason is that the original notes are found a bit ambiguous and I want to further clarify.

Currently, these are just drafts of the lecture notes. There can be typos and mistakes anywhere. So, if you find anything that needs to be corrected or improved, please inform at jingle.fu@graduateinstitute.ch.

Contents

1.	Review of Econometrics I	1
1.1.	Basic assumptions	1
1.2.	Frisch-Waugh-Lovell Theorem	1
1.3.	Endogeneity	3
1.3.1.	Instrumental Variables and 2SLS	5
1.3.2.	Weak Identification in IV Models	6
2.	Causal Inference	8
2.1.	Potential Outcomes Framework	8
2.1.1.	Identification of Causal Effects	10
3.	Panel Data Analysis	11
3.1.	Incidental Parameters Problem	11
3.1.1.	Pooled OLS Estimation	11
3.1.2.	Asymptotic Normality	12
3.1.3.	One-way error component model	13
3.2.	Random Effects	14
3.2.1.	Basic Assumptions and POLS	14
3.2.2.	From POLS to GLS	15
	RE Consistency	17
	RE Asymptotic Distribution	17
3.2.3.	Comparing POLS and GLS	18
3.3.	Fixed Effects	18
3.3.1.	Within Transformation	19
	FE Consistency	20
	FE Asymptotic Distribution	20
3.3.2.	First Difference Transformation	22
	FE-FD Consistency	22
	FE-FD Asymptotic Distribution	23
3.3.3.	Hausman Test for Random vs. Fixed Effects	25
3.3.4.	FE-IV Estimation	25
4.	Univariate Time Series	28
4.1.	Fundamentals of Time Series Analysis	28
4.1.1.	Stationarity and Strict Stationarity	29
4.1.2.	Transformations of Stationary Processes	30
4.1.3.	Ergodicity	31

4.1.4.	Conditioning on Information Sets	33
4.1.5.	Martingale Difference Sequences	34
4.2.	ARMA Models	36
4.2.1.	White Noise	36
4.2.2.	ARMA(p, q) Models	37
4.2.3.	Autoregressive Models (AR(p))	38
4.2.4.	Moving Average Models (MA(q))	40
4.3.	Inference of Univariate Time Series Models	41
4.3.1.	Estimation of AR(p) Models	41
	OLS estimator	42
	Consistency of OLS Estimator	42
	Asymptotic Normality of OLS Estimator	43
	MLE estimator	45
	Unit-Root	45
4.3.2.	Estimating Regressions with Autocorrelated Errors	46
5.	Multivariate Time Series	48
5.1.	Multivariate Processes	48
5.2.	Vector Autoregressions (VARs)	49
5.2.1.	Stationarity of VAR(p)	50
5.2.2.	Moments & GLP-Representation	51
5.2.3.	Reduced-Form vs. Structural Representation	54
5.3.	Estimation of Reduced-Form VARs	55
5.4.	Estimation of Structural VARs	56
5.5.	Point Restrictions	57
6.	Appendix	59
6.1.	Yule-Walker	59
6.2.	Kronecker Products & Vector Operator	59
	Recommended Resources	60

Review of Econometrics I

1.1 Basic assumptions

Firstly, we recall the basic assumptions of the linear regression model.

Assumption 1.1.1 (Basic Assumptions).

- A0. (Correct Specification). Model is correctly specified: $y_i = x_i'\beta + u_i$
- A1. (Independent Sampling). Observations $z_i = \{y_i, x_i\}_{i=1}^n$ are independent across i .
- A2. (Full rank). The matrix $X'X = \sum x_i x_i'$ is of full rank.
- A3. (Conditional Independence). $\mathbb{E}[u_i|x_i] = 0$.
- A4. (Homoskedasticity). $\mathbb{V}[u_i|x_i] = \sigma^2$ for all i .
 $\mathbb{V}[y_i] = \mathbb{V}[x_i'\beta + u_i|x_i] = \sigma^2$

Under these four basic assumptions, and that x_i is exogenous, giving $\mathbb{E}[x_i u_i] = 0$, then

$$\hat{\beta} = (X'X)^{-1}X'Y \xrightarrow{P} \beta.$$

1.2 Frisch-Waugh-Lovell Theorem

Definition 1.2.1 (Partitioned regression).

We consider a normal linear regression model $Y = X\beta + U$. Let X be partitioned as:

$$X = \begin{bmatrix} X_1 & X_2 \end{bmatrix}$$

where X is $n \times K$, X_1 is $n \times K_1$ and X_2 is $n \times K_2$. And we partition the parameter vector β accordingly:

$$\beta = \begin{bmatrix} \beta_1 \\ \beta_2 \end{bmatrix}$$

where β_1 is $K_1 \times 1$ and β_2 is $K_2 \times 1$. Thus, the model can be written as:

$$Y = X_1\beta_1 + X_2\beta_2 + U$$

where U is the error term.

Also take the following notation:

$$P_1 = X_1(X_1'X_1)^{-1}X_1', \quad M_1 = I - P_1, \quad \tilde{X}_2 = M_1X_2, \quad \tilde{U} = M_1Y$$

thus \tilde{U} is the residual vector from the regression of Y on X_1 , and the k -th column of \tilde{X}_2 is the residual vector from the regression of the corresponding k -th column of X_2 on X_1 .

The OLS estimator $\beta = (\beta_1, \beta_2)$ can be obtained by regression of Y on $X = [X_1, X_2]$, and can be written as:

$$Y = X\hat{\beta} + \hat{U} = X_1\hat{\beta}_1 + X_2\hat{\beta}_2 + \hat{U}$$

We are interested in algebraic expressions for $\hat{\beta}_1$ and $\hat{\beta}_2$.

Let's first focus on $\hat{\beta}_1$. The least squares estimator $\hat{\beta}_1$ is found by the joint minimization:

$$(\hat{\beta}_1, \hat{\beta}_2) = \arg \min_{\beta_1, \beta_2} (Y - X_1\beta_1 - X_2\beta_2)'(Y - X_1\beta_1 - X_2\beta_2)$$

Denote $(Y - X_1\beta_1 - X_2\beta_2)'(Y - X_1\beta_1 - X_2\beta_2)$ as $\text{SSE}(\beta_1, \beta_2)$.

By nested minimization, we can rewrite the above as:

$$\hat{\beta}_1 = \arg \min_{\beta_1} \left(\min_{\beta_2} \text{SSE}(\beta_1, \beta_2) \right)$$

For the inner minimization problem: $\min_{\beta_2} \text{SSE}(\beta_1, \beta_2)$, this is simply the regression of $Y - X_1\beta_1$ on X_2 , with the solution:

$$\arg \min_{\beta_2} \text{SSE}(\beta_1, \beta_2) = (X_2'X_2)^{-1}X_2'(Y - X_1\beta_1)$$

with residuals:

$$Y - X_1\beta_1 - X_2(X_2'X_2)^{-1}X_2'(Y - X_1\beta_1) = (M_2Y - M_2X_1\beta_1) = M_2(Y - X_1\beta_1)$$

where $M_2 = I - X_2(X_2'X_2)^{-1}X_2'$ is the annihilator matrix for X_2 .

So the inner minimization problem has minimized value:

$$\min_{\beta_2} \text{SSE}(\beta_1, \beta_2) = (Y - X_1\beta_1)'M_2'(M_2(Y - X_1\beta_1)) = (Y - X_1\beta_1)'M_2(Y - X_1\beta_1)$$

Substituting this into the outer minimization problem, we have:

$$\begin{aligned} \hat{\beta}_1 &= \arg \min_{\beta_1} (Y - X_1\beta_1)'M_2(Y - X_1\beta_1) \\ &= (X_1'M_2X_1)^{-1}(X_1'M_2Y) \end{aligned}$$

By a similar argument we find

$$\hat{\beta}_2 = (X_2'M_1X_2)^{-1}(X_2'M_1Y)$$

By the previous notation and the fact that M_1 and M_2 are idempotent matrices, we can have:

$$\begin{aligned} \hat{\beta}_2 &= (X_2'M_1X_2)^{-1}(X_2'M_1Y) \\ &= (X_2'M_1M_1X_2)^{-1}(X_2'M_1M_1Y) \\ &= (\tilde{X}_2'\tilde{X}_2)^{-1}(\tilde{X}_2'\tilde{U}) \end{aligned}$$

Thus the coefficient estimator $\hat{\beta}_2$ is algebraically equivalent to the OLS estimator of the regression of \tilde{U} on \tilde{X}_2 . Notice that these two are M_1Y and M_1X_2 , and we know that pre-multiplication by M_1 creates least squares residuals. Therefore, \tilde{U} is simply the least squares residual from a regression of Y on X_1 , and the columns of \tilde{X}_2 are the least squares residuals from a regression of the columns of X_2 on X_1 . From the above steps, we have proven the following theorem.

Theorem 1.2.1 (Frisch-Waugh-Lovell (FWL) theorem).

In the model $Y = X_1\beta_1 + X_2\beta_2 + U$, the OLS estimator of β_2 and the OLS residuals \hat{U} may be computed via the following two steps:

1. Regress Y on X_1 and obtain the residuals $\hat{U}_1 = Y - X_1\hat{\beta}_1$.
2. Regress the columns of X_2 on X_1 and obtain the residuals $\tilde{X}_2 = M_1X_2$.
3. Regress \hat{U}_1 on \tilde{X}_2 : $\hat{U}_1 = \tilde{X}_2b + V$ and obtain the OLS estimator $\hat{\beta}_2 = \hat{b}$ and the residual $\hat{U} = \hat{V}$.

1.3 Endogeneity

We say that there's endogeneity in the linear model

$$y_i = x_i'\beta + u_i$$

if β is the parameter of interest and

$$\mathbb{E}[x_i u_i] \neq 0.$$

This is a core problem in econometrics and largely differentiates the field from statistics.

Endogeneity implies that the least squares estimator is inconsistent for the structural parameter. Indeed, under i.i.d. sampling, least squares is consistent for the projection coefficient.

$$\hat{\beta} \xrightarrow{p} \beta + \left(\mathbb{E}[XX']\right)^{-1} \mathbb{E}[Xu] \neq \beta$$

The inconsistency of least squares is typically referred to as **endogeneity bias** or **estimation bias** due to endogeneity.

Commonly, there are three reasons for endogeneity:

1. Measurement error: x_i is measured with error.

Suppose our true Regression is: $y_i = x_i^{*'}\beta + \varepsilon_i$, $\mathbb{E}[x_i^*\varepsilon_i] = 0$, β is the structural parameter. But, x_i^{*} is not observed. Instead, we observe: $x_i = x_i^* + v_i$, where v_i is the measurement error, independent of x_i^* and ε_i : $\mathbb{E}[x_i^*v_i'] = 0$, $\mathbb{E}[v_i\varepsilon_i] = 0$ ¹.

The model $x_i = x_i^* + v_i$ with x_i^* and v_i uncorrelated, and $\mathbb{E}[v_i] = 0$ is known as the **classical measurement error model**. This means that x_i is a noisy but unbiased estimate of x_i^* . By substitution we can express y_i as a function of the observed variable x_i .

$$y_i = x_i^{*'}\beta + \varepsilon_i = (x_i - v_i)'\beta + \varepsilon_i = x_i'\beta + u_i$$

where $u_i = \varepsilon_i - v_i'\beta$.

This means that (y_i, x_i) satisfy the linear equation $y_i = x_i'\beta + u_i$ with an error u_i . But this error is not a projection error.

$$\begin{aligned} \mathbb{E}[x_i u_i] &= \underbrace{\mathbb{E}[x_i \varepsilon_i]}_0 - \mathbb{E}[x_i v_i']\beta \\ &= -\mathbb{E}[(x_i^* + v_i)v_i']\beta \\ &= -\underbrace{\mathbb{E}[x_i^* v_i']}_0 \beta - \mathbb{E}[v_i v_i']\beta \\ &= -\mathbb{E}[v_i v_i']\beta \neq 0 \end{aligned}$$

if $\mathbb{E}[v_i v_i'] \neq 0$ and $\beta \neq 0$.

¹This is an example of a latent variable model, where “latent” refers to an unobserved structural variable.

Remark (Measurement Error Bias).

Let's rewrite in matrix form: $Y = X^*'\beta + \varepsilon$, $X = X^* + v$, $\mathbb{E}[X^*v] = 0$, $\mathbb{E}[\varepsilon v] = 0$, v is a $k \times 1$ error. We can write $Y = (X - v)'\beta + \varepsilon = X'\beta + u$, where $u = \varepsilon - v'\beta$. And we have: $\mathbb{E}[Xu] = \mathbb{E}[(X^* + v)(\varepsilon - v'\beta)] = -\mathbb{E}[vv']\beta \neq 0$ if $\beta \neq 0$ and $\mathbb{E}[vv'] \neq 0$.

We can calculate the form of the projection coefficient (which is consistently estimated by least squares). For simplicity suppose that $k = 1$, we find:

$$\beta^* = \beta + \frac{\mathbb{E}[Xu]}{\mathbb{E}[X^2]} = \beta \left(1 - \frac{\mathbb{E}[v^2]}{\mathbb{E}[X^2]} \right)$$

Since $\frac{\mathbb{E}[v^2]}{\mathbb{E}[X^2]} < 1$, the projection coefficient shrinks the structural parameter β towards zero. This is called **measurement error bias** or **attenuation bias**.

2. Simultaneity (Reverse causality): Simultaneity arises when at least one of the explanatory variables is determined simultaneously along with y . If, say, x_i is determined partly by y , and x_i and u_i are generally correlated.

$$y_i = x_i'\beta + u_i = x_{i1}'\beta_1 + x_{i2}\beta_2 + u_i, \quad x_i = z_i'\gamma + y_i\delta + v_i.$$

Example 1 (Supply and Demand).

The variables Q and P (quantity and price) are determined jointly by the demand equation:

$$Q = -\beta_1 P + u_1$$

and supply function:

$$Q = \beta_2 P + u_2$$

where u_1 and u_2 are the demand and supply shocks, respectively. Assume that $u = (u_1, u_2)$ satisfy that $\mathbb{E}[u] = 0$ and $\mathbb{E}[uu'] = I_2$ (for simplicity). The question is: If we regress Q on P , what will happen? Let's solve P and Q in error terms:

$$\begin{aligned} \begin{bmatrix} 1 & \beta_1 \\ 1 & -\beta_2 \end{bmatrix} \begin{bmatrix} Q \\ P \end{bmatrix} &= \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \\ \begin{bmatrix} Q \\ P \end{bmatrix} &= \begin{bmatrix} 1 & \beta_1 \\ 1 & -\beta_2 \end{bmatrix}^{-1} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \\ &= \frac{1}{\beta_1 + \beta_2} \begin{bmatrix} \beta_2 & \beta_1 \\ 1 & -1 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \\ &= \begin{bmatrix} \frac{\beta_2 u_1 + \beta_1 u_2}{\beta_1 + \beta_2} \\ \frac{u_1 - u_2}{\beta_1 + \beta_2} \end{bmatrix} \end{aligned}$$

The projection of Q on P yields $Q = \beta^* P + u^*$, where $\mathbb{E}[Pu^*] = 0$ and the projection coefficient is

$$\beta^* = \frac{\mathbb{E}[PQ]}{\mathbb{E}[P^2]} = \frac{\beta_2 - \beta_1}{2}.$$

The OLS estimator satisfies $\hat{\beta} \xrightarrow{p} \beta^*$ and the limit does not equal either β_1 or β_2 . This is called **simultaneity bias** or **simultaneous equation bias**.

This occurs generally when Y and X are jointly determined, as in a market equilibrium. Generally, when both the dependent variable and a regressor are simultaneously determined

then the regressor should be treated as endogenous.

3. Omitted variables: The most prominent cause of endogeneity are omitted variables (OVs).

Suppose the true regression is: $y_i = x_i'\beta + w_i'\delta + \varepsilon_i$, where exogeneity holds: $\mathbb{E}[x_i\varepsilon_i] = 0$, $\mathbb{E}[w_i\varepsilon_i] = 0$.

If we omit w_i and instead estimates:

$$y_i = x_i'\beta + u_i$$

where $u_i = w_i'\delta + \varepsilon_i$, then in this misspecified model, exogeneity is only given if x_i and w_i are uncorrelated, since:

$$\begin{aligned}\mathbb{E}[x_i u_i] &= \mathbb{E}[x_i(w_i'\delta + \varepsilon_i)] \\ &= \mathbb{E}[x_i w_i']\delta + \underbrace{\mathbb{E}[x_i \varepsilon_i]}_0\end{aligned}$$

Since $\hat{\beta} - \beta \xrightarrow{P} \mathbb{E}[x_i x_i']^{-1} \mathbb{E}[x_i u_i]$, we can assess the sign and size of the asymptotic bias based on the signs of correlation between x_i and w_i .

For our general regression model $y_i = x_i'\beta + u_i$, we have $\mathbb{E}[x_i u_i] \neq 0$, thus $\hat{\beta}_{OLS} \xrightarrow{P} \beta$ doesn't hold.

To consistently estimate β , we require additional assumptions. One type of information which is commonly used in economics is the **instruments**.

Definition 1.3.1 (Instrumental Variable).

We take $z_i \in \mathbb{R}^r$ as an instrumental variable if:

$$\begin{aligned}\mathbb{E}[z_i u_i] &= 0 \\ \mathbb{E}[z_i x_i] &\neq 0 \\ \mathbb{E}[z_i z_i'] &> 0 \\ \text{rank}\left(\mathbb{E}[z_i z_i']\right) &= k \leq r \quad 2\end{aligned}$$

We say that the model is just-identified if $k = r$ and over-identified if $k < r$.

1.3.1 Instrumental Variables and 2SLS

Then, we have the 2SLS method:

Definition 1.3.2 (2SLS Method).

1. Estimate: $x_i = z_i'\gamma + e_i \Rightarrow \hat{\gamma} = (Z'Z)^{-1}Z'X \Rightarrow \hat{X} = Z'\hat{\gamma} = P_Z X$;
2. Estimate: $y_i = \hat{x}_i'\beta + u_i^*$.

$$\begin{aligned}\hat{\beta}_{2SLS} &= (\hat{X}'\hat{X})^{-1}\hat{X}'Y \\ &= ((P_Z X)'P_Z X)^{-1}(P_Z X)'Y \\ &= (X'P_Z X)^{-1}X'P_Z Y \\ &= (X'Z(Z'Z)^{-1}Z'X)^{-1}X'Z(Z'Z)^{-1}Z'Y \\ &= \beta + (X'Z(Z'Z)^{-1}Z'X)^{-1}X'Z(Z'Z)^{-1}Z'u \\ &= \beta + (Z'X)^{-1}(Z'Z)(X'Z)^{-1}X'Z(Z'Z)^{-1}Z'u \quad 3 \\ &= \beta + (Z'X)^{-1}Z'u \\ &\xrightarrow{P} \beta.\end{aligned}$$

To compute $\hat{\beta}_{2SLS}$, we need $Z'Z$ to be full rank, which requires us to have more observations than IVs.

Ideally, z_i should be as highly correlated with x_i as possible, but uncorrelated with u_i . To see this, we find the variance of $\hat{\beta}_{2SLS}$

$$\begin{aligned}\mathbb{V}[\hat{\beta}_{2SLS}|X, Z] &= \mathbb{V}[(X'P_ZX)^{-1}X'P_ZU|X, Z] \\ &= (X'P_ZX)^{-1}\mathbb{V}[X'P_ZU|X, Z](X'P_ZX)^{-1} \\ &= (X'P_ZX)^{-1}X'P_Z\mathbb{E}[UU'|X, Z]P_ZX(X'P_ZX)^{-1} \\ &= (X'P_ZX)^{-1}\sigma^2\end{aligned}$$

which holds under homoskedasticity. As we know $\mathbb{V}[\hat{\beta}_{OLS}] = (X'X)^{-1}\sigma^2$,

$$\begin{aligned}\mathbb{V}[\hat{\beta}_{OLS}]^{-1} - \mathbb{V}[\hat{\beta}_{2SLS}]^{-1} &= (\sigma^2)^{-1}X'X - (\sigma^2)^{-1}X'P_ZX \\ &= (\sigma^2)^{-1}X'(I - P_Z)X \\ &= (\sigma^2)^{-1}X'M_ZX \\ &= \sigma^{-2}\underbrace{(M_ZX)'}_{\hat{E}}M_ZX \\ &= \sigma^{-2}SSR_{1SLS} > 0.\end{aligned}$$

This means that the variance of 2SLS estimator is larger than that of the OLS.

By the usual arguments, the asymptotic analysis reveals that:

$$\sqrt{n}(\hat{\beta}_{2SLS} - \beta) \xrightarrow{D} \mathcal{N}(0, V_{2SLS})$$

where

$$V_{2SLS} = Q_{XZ}^{-1}X'Z(Z'Z)^{-1}Z'UU'Z(Z'Z)^{-1}(X'Z)'Q_{XZ}^{-1}$$

where $Q_{XZ} = (Z'X)(Z'Z)^{-1}(X'Z)$

As usual, we can estimate it by replacing u_i with \hat{u}_i and expectation operators with population means. Thereby, it's important to note that $u_i \neq u_i^*$, and to obtain \hat{u}_i , we don't use \hat{x}_i , but x_i :

$$\hat{u}_i = y_i - x_i'\hat{\beta}_{2SLS}$$

Under homoskedasticity, $V_{2SLS} = \sigma^2 Q_{XZ}^{-1}$, which we estimate using $\hat{\sigma}^2 = \frac{1}{n} \sum_i \hat{u}_i^2$.

1.3.2 Weak Identification in IV Models

If the correlation between x_i and z_i is weak, then we say it's a **weak instrument**. Under weak IVs, the finite sample distribution of $\hat{\beta}_{2SLS}$ may not assemble the asymptotic property.

In absence of an asymptotic distribution, we can conduct inference using its numerical approximation via bootstrapping. Or alternatively, we can construct a confidence set for β using the following procedure of Anderson and Rubin (1949).

The method is based on the idea that, for $\beta = \beta_0$, the auxiliary regression $y_i - x_i'\beta = \delta z_i + v_i$ should yield $\delta = 0$, because $y_i - x_i'\beta_0 = u_i$ and u_i is uncorrelated with z_i .

Theorem 1.3.1 (Anderson-Rubin Method).

For a given β_0 , we get:

$$\sqrt{n}\hat{\delta}(\beta_0) = \sqrt{n}(Z'Z)^{-1}Z'(Y - X\beta_0) = (Z'Z)^{-1}\sqrt{n}Z'U \xrightarrow{d} \mathcal{N}\left(0, \frac{\sigma_u^2}{\mathbb{E}(z_i^2)}\right)$$

which allows us to test $\mathcal{H}_0 : \delta = 0$. For many β s, test: $\mathcal{H}_0 : \delta(\beta) = 0$, e.g. using t-test.

$$T_t = \frac{\hat{\delta}(\beta_0)}{se(\hat{\delta}(\beta_0))} = \frac{\hat{\delta}_0}{\sqrt{\hat{\sigma}_u^2/Z'Z}} \xrightarrow{d} \mathcal{N}(0, 1)$$

The 90% CI for β is the set of β s at which $\delta(\beta) = 0$ cannot be rejected at 90% confidence level. A confidence set for β is given by taking all β_0 such that $\mathcal{H}_0 : \delta = 0$ cannot be rejected.

Remark (About Anderson-Rubin (AR) Test). ^a

Consider our model

$$\begin{aligned} y &= X\beta + u, \\ X &= Z\Pi + v, \end{aligned}$$

where X is one-dimensional and test for hypothesis $H_0 : \beta = \beta_0$. Under the null, vector $y - X\beta$ is equal to the error u_t and is uncorrelated with Z (due to exogeneity of instruments). The suggested statistics is:

$$AR(\beta_0) = \frac{(y - X\beta)'P_Z(y - X\beta)}{(y - X\beta)'M_Z(y - X\beta)/(T - k)}.$$

here $P_Z = Z(Z'Z)^{-1}Z'$, $M_Z = I - P_Z$.

The distribution of AR does not depend on μ asymptotically $AR \rightarrow \chi_k^2/k$. The formula may remind you of the J-test for over-identifying restrictions. It would be a J-test if one were to plug in $\hat{\beta}_{TSLS}$. In a more general situation of more than one endogenous variable and/or included exogenous regressors AR statistic is F-statistic testing that all coefficients on Z are zero in the regression of $y - \beta_0 X$ on Z and W .

Note, that one tests all coefficients β simultaneously (as a set) in a case of more than one endogenous regressor. AR confidence set One can construct a confidence set robust towards weak instruments based on the AR test by inverting it. That is, by finding all β which are not rejected by the data. In this case, it is the set :

$$CI = \{\beta_0 : AR(\beta_0) < \chi_{k,1-\alpha}^2\}.$$

The nice thing about this procedure is that solving for the confidence set is equivalent to solving a quadratic inequality. This confidence set can be empty with positive probability (caution!).

^aRetrieved from MIT14.384 Time Series Analysis, Fall 2007 Professor Anna Mikusheva, Lecture 7-8, https://ocw.mit.edu/courses/14-384-time-series-analysis-fall-2013/365cba34145fa204731e9df202d4771e_MIT14_384F13_lec7and8.pdf

Lecture 2.

Causal Inference

Rubin (1975[15]) and Holland (1986[16]) made up the aphorism[1]:

“No causation without manipulation”

Not everybody agrees with this point of view. In our lecture, we’ll define causal effects using the potential outcomes framework (Neyman, 1923[17]; Rubin, 1974[18]).

2.1 Potential Outcomes Framework

In this framework, an experiment, or at least a thought experiment, has a treatment, and we are interested in its effect on an outcome or multiple outcomes. Sometimes, the treatment is also called an intervention or a manipulation.

Firstly, we consider an experiment with n units indexed by $i = 1, 2, \dots, n$. We focus on a treatment with two levels:

$$d_i = \begin{cases} 0 & \text{control} \\ 1 & \text{treatment} \end{cases}$$

We seek to identify the causal effect of treatment d_i on some outcome y_i . For each i , the outcome of interest y_i has two versions:

$$y_i = \begin{cases} y_{0i} & d_i = 0 \\ y_{1i} & d_i = 1 \end{cases}$$

This notation emphasizes that y_{di} is the realization of the outcome y_i that would materialize if unit i received treatment $d_i = d$.

Neyman (1923[17]) first used this notation. It seems intuitive but has some hidden assumptions. Rubin (1980[19]) made the following clarifications on the hidden assumptions.

Assumption 2.1.1 (No interference).

Unit i ’s potential outcomes do not depend on other units’ treatments. This is sometimes called the no-interference assumption.

Assumption 2.1.2 (Consistency).

There are no other versions of the treatment. Equivalently, we require that the treatment levels be well-defined, or have no ambiguity at least for the outcome of interest. This is sometimes called the consistency assumption.

The causal effect of the treatment on the i -th unit is then defined as:

$$\Delta_i = y_{1i} - y_{0i}$$

These potential outcomes are constants at the level of unit i .

Remark (Problem of causal inference).

The fundamental problem in causal inference is that only one treatment can be assigned to a given individual, and so only one of y_{0i} and y_{1i} can be observed. Thus Δ_i can never be observed.

Definition 2.1.1 (Stable Unit Treatment Value Assumption (SUTVA)).

Rubin (1980[19]) called the Assumptions 2.1.1 and 2.1.2 above together the *Stable Unit Treatment Value Assumption (SUTVA)*.

The observed outcome of unit i is a function of the potential outcomes and the treatment indicator, we can write:

$$y_i = d_i y_{1i} + (1 - d_i) y_{0i}$$

In principle, by virtue of being (discrete) RVs, both d_i and y_i each have a distribution function, which, together with their possible realizations, defines various moments. However, their unconditional probabilities and moments at the level of unit i is not of interest. Only the conditional probabilities of y_i given d_i is of interest.

Remark (Rubin (2005[20])).

Under SUTVA, Rubin (2005) called the $n \times 2$ matrix of potential outcomes the Science Table:

i	y_{1i}	y_{0i}
1	y_{11}	y_{01}
2	y_{12}	y_{02}
\vdots	\vdots	\vdots
n	y_{1n}	y_{0n}

Due to the fundamental contributions of Neyman and Rubin to statistical causal inference, the potential outcomes framework is sometimes referred to as the Neyman Model, the Neyman-Rubin Model, or the Rubin Causal Model. Causal effects are functions of the Science Table. Inferring individual causal effects

$$\tau_i = y_{1i} - y_{0i}, \quad (i = 1, \dots, n)$$

is fundamentally challenging because we can only observe either y_{1i} or y_{0i} , for each unit i , that is, we can observe only half of the Science Table.

SUTVA(2.1.1) ensures that the individual treatment effect is well defined.

Now, although Δ_i itself is unobservable, we can (perhaps remarkably) use randomized experiments to learn certain properties of it. The expectations $\mathbb{E}[y_{0i}]$ and $\mathbb{E}[y_{1i}]$ denote the average potential outcomes across unit i in population.

In particular, large randomized experiments let us recover the **Average Treatment Effect (ATE)**:

$$\text{ATE} = \mathbb{E}[y_{1i} - y_{0i}] = \mathbb{E}[y_{1i}] - \mathbb{E}[y_{0i}]$$

For a population, we can define the treatment conditional expectations:

$$\mathbb{E}[y_i | d_i = 1], \mathbb{E}[y_{0i} | d_i = 1], \mathbb{E}[y_{1i} | d_i = 1] = \mathbb{E}[y_i | d_i = 1]$$

that denote the averages of the outcome y_i .

Analogously, we can define the control conditional expectations:

$$\mathbb{E}[y_i|d_i = 0], \mathbb{E}[y_{0i}|d_i = 0] = \mathbb{E}[y_i|d_i = 0], \mathbb{E}[y_{1i}|d_i = 0]$$

for the non-treated subpopulation.

Similar to ATE, we can define the Average Treatment Effect for the Treatment-Group (ATT) and the Average Treatment Effect for the Control-Group (ATC) as distinct objects:

$$\begin{aligned} \text{ATT} &= \mathbb{E}[y_{1i} - y_{0i}|d_i = 1] \\ \text{ATC} &= \mathbb{E}[y_{1i} - y_{0i}|d_i = 0] \\ \mathbb{E}[z] &= \mathbb{E}[z|d = 1]\mathbb{P}[d = 1] + \mathbb{E}[z|d = 0]\mathbb{P}[d = 0] = \mathbb{E}[\mathbb{E}[z|d]] \end{aligned}$$

2.1.1 Identification of Causal Effects

Now, suppose we observe treatments and outcomes over a random sample n from the overall population, $\{d_i, y_i\}_{i=1}^n = \{d_i, y_{d_i}\}_{i=1}^n$, as either $y_I = y_{0i}$, or $y_i = y_{1i}$.

Let $n_w = |\{i : d_i = w\}|$ be the size of sets of units in our sample who received and did not receive treatment, respectively. This means that: while we observe a sample of size n of d_i and y_i from the overall population, we are observing a sample of size n_0 of realizations of y_{0i} from the non-treated subpopulation and a sample of size n_1 of realizations of y_{1i} from the treated subpopulation.

$N = \{i = 1, 2, \dots, n\}$, $N_1 = \{i \in N : d_i = 1\} \leftarrow n_1 = |N_1|$, $N_0 = \{i : d_i = 0\} \leftarrow n_0 = |N_0|$.

Based on this data, we can use the analogy principle to consistently estimate the first term in the ATT formula and the second term in the ATC formula:

$$\begin{aligned} \frac{1}{n_1} \sum_{i \in N_1} y_i &= \frac{1}{n_1} \sum_{i \in N_1} y_{1i} \xrightarrow{p} \mathbb{E}[y_{1i}|d_i = 1] = \mathbb{E}[y_i|d_i = 1] \\ \frac{1}{n_0} \sum_{i \in N_0} y_i &= \frac{1}{n_0} \sum_{i \in N_0} y_{0i} \xrightarrow{p} \mathbb{E}[y_{0i}|d_i = 0] = \mathbb{E}[y_i|d_i = 0] \end{aligned}$$

Without further assumptions, we cannot identify the remaining terms. Firstly, we cannot observe $\mathbb{E}[y_{0i}|d_i = 1]$ and $\mathbb{E}[y_{1i}|d_i = 0]$ because we do not observe y_{0i} for treated units, and we do not observe y_{1i} for non-treated units. Secondly, we can not observe $\mathbb{E}[y_{1i}]$ and $\mathbb{E}[y_{0i}]$ because both N_1 and N_0 are random samples from the overall population. As a result, the ATE is in general not identified from our data!

We can define the the difference-in-means estimator as:

$$\hat{\tau}_{DM} = \frac{1}{n_1} \sum_{i \in N_1} y_i - \frac{1}{n_0} \sum_{i \in N_0} y_i \xrightarrow{p} \mathbb{E}[y_{1i}|d_i = 1] - \mathbb{E}[y_{0i}|d_i = 0] = \text{ATE} = \text{ATT} = \text{ATC}.$$

We define the difference of treated and non-treated as: *Naive Difference*.

$$\begin{aligned} \text{ND} &= \mathbb{E}[y_{1i}|d_i = 1] - \mathbb{E}[y_{0i}|d_i = 0] \\ &= \mathbb{E}[y_{1i}|d_i = 1] - \mathbb{E}[y_{0i}|d_i = 1] + \mathbb{E}[y_{0i}|d_i = 1] - \mathbb{E}[y_{0i}|d_i = 0] \\ &= \text{ATT} + \mathbb{E}[y_{0i}|d_i = 1] - \mathbb{E}[y_{0i}|d_i = 0] \end{aligned}$$

For LRM, $y_i = \beta_0 + \beta_1 d_i + u_i$,

$$\begin{aligned} \text{ND} &= \mathbb{E}[y_i|d_i = 1] - \mathbb{E}[y_i|d_i = 0] \\ &= \mathbb{E}[\beta_0 + \beta_1 + u_i|d_i = 1] - \mathbb{E}[\beta_0 + u_i|d_i = 0] \\ &= \beta_1 + \mathbb{E}[u_i|d_i = 1] - \mathbb{E}[u_i|d_i = 0] \end{aligned}$$

Panel Data Analysis

Economists traditionally use the term **panel data** to refer to data structures consisting of observations on individuals for multiple time periods. There are several distinct advantages of panel data relative to cross-section data:

1. Possibility of controlling for unobserved time-invariant endogeneity without the use of instrumental variables
2. Possibility of allowing for broader forms of heterogeneity
3. Modeling dynamic relationships and effects

It's typical to index observations by both the individual i and the time period, t , thus y_{it} denotes a variable for individual i in time t , where $n = 1, \dots, N$, $t = 1, \dots, T$.

Definition 3.0.1 (Balanced and Unbalanced Panel Data[2]).

When observations are available on all individuals for the same time periods we say that the panel is **balanced**. In this case there are an equal number T of observations for each individual and the total number of observations is $n = NT$.

When different time periods are available for the individuals in the sample we say that the panel is **unbalanced**. This is the most common type of panel data set. It does not pose a problem for applications but does make the notation cumbersome and also complicates computer programming.

3.1 Incidental Parameters Problem

3.1.1 Pooled OLS Estimation

Suppose we are estimating the following panel data regression:

$$y_{it} = \alpha + x'_{it}\beta + u_{it}, \quad \mathbb{E}[u_{it}|x_{it}] = 0, \quad \mathbb{V}[u_{it}|x_{it}] = \sigma^2$$

Omitting the distinction between intercept and slope, we can write the model as:

$$y_{it} = \tilde{x}'_{it}\tilde{\beta} + u_{it}$$

$$\tilde{x}_{it} = \begin{bmatrix} 1 \\ x_{it} \end{bmatrix}, \quad \tilde{\beta} = \begin{bmatrix} \alpha \\ \beta \end{bmatrix}$$

where $i = 1 : N$, $T = 1 : t$.

Or, we can write the model as:

$$\underset{T \times 1}{y_i} = \underset{T \times K}{\tilde{X}_i} \underset{K \times 1}{\tilde{\beta}} + \underset{T \times 1}{u_i}$$

Using OLS method to estimate $\tilde{\beta}$, we have:

$$\min_{\tilde{\beta}} \sum_i \sum_t u_{it}^2 = \min_{\tilde{\beta}} \sum_i u_i' u_i = \min_{\tilde{\beta}} (y_i - \tilde{X}_i \tilde{\beta})' (y_i - \tilde{X}_i \tilde{\beta})$$

The FOC of this equation is:

$$\begin{aligned}
\sum_i -\tilde{X}_i'(y_i - \tilde{X}_i\tilde{\beta}) &= 0 \\
\left(\sum_i \tilde{X}_i'\tilde{X}_i\right)\tilde{\beta} &= \sum_i \tilde{X}_i'y_i \\
\hat{\beta}_{POLS} &= \left(\sum_i \tilde{X}_i'\tilde{X}_i\right)^{-1} \sum_i \tilde{X}_i'y_i \\
&= \left(\sum_i \sum_t \tilde{x}_{it}\tilde{x}_{it}'\right)^{-1} \left(\sum_i \sum_t \tilde{x}_{it}y_{it}\right) \\
&= \tilde{\beta} + \left(\frac{1}{NT} \sum_i \sum_t \tilde{x}_{it}\tilde{x}_{it}'\right)^{-1} \frac{1}{NT} \left(\sum_i \sum_t \tilde{x}_{it}u_{it}\right) \\
&\xrightarrow{p} \tilde{\beta} + \mathbb{E}\left[\sum_t \tilde{x}_{it}\tilde{x}_{it}'\right]^{-1} \mathbb{E}\left[\sum_t \tilde{x}_{it}u_{it}\right] \\
&= \tilde{\beta}
\end{aligned}$$

Hence $\hat{\beta}_{OLS}$ is consistent provided that x_{it} and u_{it} are contemporaneously uncorrelated, as $\mathbb{E}[x_{it}u_{it}] = 0, \forall t$. The regressors are allowed to be correlated with the past, and future u_{it} . This occurs when there's feedback loop by which $y_{i,t-1}$ affects x_{it} .

In this proof, we show that either $N \rightarrow \infty$ or $T \rightarrow \infty$ is sufficient for consistency of $\hat{\beta}_{POLS}$. However, most panel data applications have a large N and small T dimension, so standard panel data features T fixed and $n \rightarrow \infty$.

3.1.2 Asymptotic Normality

From the analysis of consistency, we know that:

$$\hat{\beta}_{POLS} = \left(\sum_i \tilde{X}_i'\tilde{X}_i\right)^{-1} \sum_i \tilde{X}_i'y_i$$

Hence:

$$\begin{aligned}
\sqrt{n}(\hat{\beta}_{POLS} - \tilde{\beta}) &= \left(\frac{1}{n} \sum_i \tilde{X}_i'\tilde{X}_i\right)^{-1} \left(\frac{1}{\sqrt{n}} \sum_i \tilde{X}_i'u_i\right) \\
&\xrightarrow{p} \mathbb{E}[\tilde{X}_i'\tilde{X}_i]^{-1} \xrightarrow{d} \mathcal{N}\left(0, \mathbb{E}\left[\left(\tilde{X}_i'u_i\right)\left(\tilde{X}_i'u_i\right)'\right]\right) \\
&\xrightarrow{d} \mathcal{N}\left(0, \mathbb{E}\left[\tilde{X}_i'\tilde{X}_i\right]^{-1} \mathbb{E}\left[\tilde{X}_i'u_i u_i' \tilde{X}_i\right] \mathbb{E}\left[\tilde{X}_i'\tilde{X}_i\right]^{-1}\right)
\end{aligned}$$

The above model is homogeneous, which is unattractive, as the data generating process would differ across i , with some units having a higher level of the outcome variable y_{it} than others, regardless of covariates x_{it} (with a higher intercept α) or a stronger effect of some covariates $x_{it,k}$ on y_{it} than others.

At the other extreme, we assume the fully heterogenous estimation:

$$y_{it} = \alpha_i + x_{it}'\beta + u_{it}, \quad \mathbb{E}[u_{it}x_{it}] = 0, \quad \mathbb{V}[u_{it}|x_{it}] = \sigma_i^2.$$

Under $T = 1$, we run $y_i = \beta_0 + x_i'\beta + v_i$, where $v_i = u_i + \underbrace{\alpha_i - \beta_0}_{\tilde{\alpha}_i}$ and $\mathbb{E}[v_i] = 0$.

Under $T > 1$, we run:

$$\begin{aligned} y_i &= x'_i \beta + \sum_{j=1}^n \alpha_j \mathbf{1}\{i = j\} + u_{it} \\ &= \tilde{x}'_{it} \tilde{\beta} + u_{it} \\ \tilde{x}_{it} &= \begin{bmatrix} x_{it} \\ \mathbf{1}\{i = 1\} \\ \mathbf{1}\{i = 2\} \\ \vdots \\ \mathbf{1}\{i = n\} \end{bmatrix}, \quad \tilde{\beta} = \begin{bmatrix} \beta \\ \alpha_1 \\ \alpha_2 \\ \vdots \\ \alpha_n \end{bmatrix} \end{aligned}$$

In a similar way, we can write the regression as

$$y_i = \tilde{X}_i \tilde{\beta}_i + u_i$$

with $\tilde{\beta}_i$ is specific for each i . We have n separate time series regressions, one for each unit i .

Following the same analyzing process, we can get:

$$\hat{\beta}_{i,OLS} = \left(\sum_i \tilde{X}_i' \tilde{X}_i \right) \sum_i \tilde{X}_i' y_i = \left(\sum_t \tilde{x}_{it} \tilde{x}_{it}' \right)^{-1} \left(\sum_t \tilde{x}_{it} y_{it} \right),$$

which obviously shows that $\hat{\beta}$ is consistent $\Leftrightarrow T \rightarrow \infty$.

3.1.3 One-way error component model

With the fully homogeneous specification unattractive and the fully heterogeneous specification infeasible, researchers usually go for a compromise and let intercepts (and error term variances) be unit-specific.

Definition 3.1.1 (One-way error component model).

$$y_{it} = \alpha_i + x'_{it} \beta + u_{it}, \quad \mathbb{E}[u_{it} x_{it}] = 0, \quad \mathbb{V}[u_{it} | x_{it}] = \sigma^2, \quad (3.1)$$

where α_i is an individual-specific effect, and u_{it} are idiosyncratic(i.i.d.) errors.

In any case, the equation above makes clear that α_i contains all factors that affect y_{it} , that are not included in x_{it} and that are fixed over time (the time-varying factors are in u_{it}).

Suppose the model is correctly specified, and we have a cross-sectional dataset available, i.e. $T = 1$. Then, we would estimate:

$$y_{it} = \beta_0 + x'_{it} \beta + v_i, \text{ for } t = 1,$$

where $v_i = \alpha_i + u_{it} - \beta_0$.

If the unobserved heterogeneity α_i is correlated with the covariate x_{it} , our standard OLS estimator is biased and inconsistent.

If we have a panel dataset, i.e. $T > 1$, we can write the above model into a regression of $k + n$ regressors:

$$y_{it} = x'_{it} \beta + \sum_{j=1}^N \mathbf{1}\{i = j\} \alpha_j + u_{it} = x_{it}^{*'} \beta^* + u_{it},$$

where $x_{it}^* = (x'_{it}, \mathbf{1}\{i = 1\}, \dots, \mathbf{1}\{i = N\})'$, and $\beta^* = (\beta', \alpha_1, \dots, \alpha_N)'$.

This leads to the pooled OLS estimator for β^* :

$$\hat{\beta}^* = \left(\sum_i \sum_t x_{it}^* x_{it}^{*'} \right) \sum_i \sum_t x_{it}^* y_{it}.$$

However, the estimator suffers from the so-called **IPP problem**, as the number of parameters increase with $N \rightarrow \infty$, the limit of $\frac{1}{N} \sum_i x_{it}^* x_{it}^{*'}$ is not well-defined and as a result, we can't establish consistency of $\hat{\beta}_{OLS}$.

3.2 Random Effects

As with pooled OLS, a random effects analysis puts α_i into the error term. In fact, random effects analysis imposes more assumptions than those needed for pooled OLS: **strict exogeneity** in addition to orthogonality between α_i and x_{it} .

3.2.1 Basic Assumptions and POLS

Stating the assumption in terms of conditional means, we have:

Assumption 3.2.1 (Random Effect).

(a) $\mathbb{E}[u_{it} | \tilde{X}_i, \tilde{\alpha}_i] = 0, \forall t.$

(b) $\mathbb{E}[\tilde{\alpha}_i | \tilde{X}_i] = \mathbb{E}[\tilde{\alpha}_i] = 0.$

where $\tilde{X}_i = (x_{i1}, \dots, x_{iT})$.

Assumption 3.2.1(a) is the strict exogeneity condition and Assumption 3.2.1(b) is how we will state the orthogonality.

Remark (Why Strict Exogeneity?[3]).

Why do we maintain Assumption 3.2.1(a) when it is more restrictive than needed for a pooled OLS analysis? Because the random effects approach exploits the serial correlation in the composite error, $v_{it} = \alpha_i + u_{it}$, in a generalized least squares (GLS) framework. In order to ensure that feasible GLS is consistent, we need some form of strict exogeneity between the explanatory variables and the composite error.

Under this assumption, we can write:

$$\begin{aligned} y_{it} &= x_{it}' \beta + v_{it} \\ \mathbb{E}[v_{it} | X_i] &= 0, t = 1, \dots, T \end{aligned}$$

The conditions shows that our model satisfies the GLS assumption, which confirms that we can apply GLS methods that account for the particular error structure $v_{it} = \alpha_i + u_{it}$.

By defining $v_{it} = u_{it} + \alpha_i - \beta_0$, we can transform the random effect model to the following:

$$\begin{aligned} y_{it} &= \alpha_i + x_{it}' \beta + u_{it} \\ &= \underbrace{\beta_0 + x_{it}' \beta}_{\tilde{x}_{it}' \tilde{\beta}} + \underbrace{u_{it} + \alpha_i - \beta_0}_{\equiv v_{it}} \end{aligned}$$

Defining again $\tilde{x}_{it} = (1, x'_{it})'$, $\tilde{\beta} = (\beta_0, \beta')'$, we can rewrite the model as:

$$\begin{aligned} y_{it} &= \tilde{x}'_{it}\tilde{\beta} + v_{it} \Leftrightarrow y_i = \tilde{X}'_i\tilde{\beta} + v_i \\ \rightarrow \hat{\tilde{\beta}} &= \left(\sum_i \tilde{X}'_i\tilde{X}_i \right)^{-1} \sum_i \tilde{X}'_i y_i \end{aligned}$$

With this intercept β_0 , $\mathbb{E}[v_i] = 0$ is guaranteed to hold. Define $\tilde{\alpha}_i = \alpha_i - \beta_0$ as the mean-zero unit-specific heterogeneity so that $v_i = u_i + \tilde{\alpha}_i$.

Note (POLS).

Homogenous spec: $y_{it} = \alpha + x'_{it}\beta + u_{it} = \tilde{x}'_{it}\tilde{\beta} + v_{it}$. $\hat{\tilde{\beta}}$ is consistent if $\mathbb{E}[v_{it}x_{it}] = 0, \forall t$.

Using pooled OLS to estimate $\hat{\tilde{\beta}}$,

$$\begin{aligned} \hat{\tilde{\beta}}_{RE-OLS/POLS} &= \left(\frac{1}{n} \sum_i \tilde{X}'_i\tilde{X}_i \right)^{-1} \frac{1}{n} \sum_i \tilde{X}'_i y_i \\ &= \tilde{\beta} + \left(\frac{1}{n} \sum_i \tilde{X}'_i\tilde{X}_i \right)^{-1} \frac{1}{n} \sum_i \tilde{X}'_i v_i \\ &\xrightarrow{p} \tilde{\beta} + \mathbb{E}[\tilde{X}'_i\tilde{X}_i]^{-1} \mathbb{E}[\tilde{X}'_i v_i] \\ \text{where } \mathbb{E}[\tilde{X}'_i v_i] &= \mathbb{E} \left[\sum_t \tilde{x}'_{it} v_{it} \right] \\ &= \sum_t \mathbb{E} [\tilde{x}'_{it} v_{it}] \\ &= \sum_t \mathbb{E} [\tilde{x}_{it}(u_{it} + \alpha_i - \beta_0)] \end{aligned}$$

Here, the error term v_i is not equal to the original error term u_{it} .

Note.

Under the random effect, you have to use the heteroskedasticity-robust methods. Because even if we assume u_{it} to be homoskedastic, v_{it} is not, as it includes also the unit-specific heterogeneity α_i .

3.2.2 From POLS to GLS

So, to obtain consistency, we need to assume that:

Assumption 3.2.2 (Random Effect Independence). (a) $\mathbb{E}[u_{it}|\tilde{x}_{it}, \tilde{\alpha}_i] = 0, \forall t$.

(b) $\mathbb{E}[\tilde{\alpha}_i|\tilde{x}_{it}] = 0, \forall t$.

And, we are also obliged to use HAC-robust standard error because:

$$\Omega \equiv \mathbb{E}[v_i v'_i | \tilde{X}_i] = \mathbb{E}[(\alpha_i \mathbf{1}_i + u_i)(\tilde{\alpha}_i \mathbf{1}_i + u_i)' | \tilde{X}_i] = \mathbb{E}[\tilde{\alpha}_i^2 \mathbf{1}_i \mathbf{1}'_i | \tilde{X}_i] + \mathbb{E}[u_i u'_i | \tilde{X}_i]$$

is not diagonal.

Assumption 3.2.3 (Random Effect Rank).

$$\text{rank } \mathbb{E}[X'_i \Omega^{-1} X_i] = K$$

We know that both GLS and feasible GLS estimator would be consistent under Assumption 3.2.1 and 3.2.3. A general FGLS analysis, using an unrestricted variance estimator Ω , is consistent and asymptotically normal as $N \rightarrow \infty$.

But, we won't exploit the unobserved effects structure v_{it} . A standard random effects analysis adds assumptions on the idiosyncratic errors that give Ω a special form. The first assumption is that the idiosyncratic errors u_{it} have a constant unconditional variance across t :

Assumption 3.2.4 (RE-Homoskedasticity).

$$\mathbb{E}[u_{it}^2] = \sigma_u^2, \forall t$$

The second assumption is that the idiosyncratic errors are serially uncorrelated:

Assumption 3.2.5 (RE-Serial Uncorrelated).

$$\mathbb{E}[u_{it}u_{is}] = 0, \forall t \neq s$$

Under these two assumptions, we can derive the variances and covariances of the elements of v_i . Given the error structure the natural estimator for β is GLS. The GLS estimator for β is:

$$\hat{\beta}_{RE-GLS} = \left(\sum_i \tilde{X}_i' \Omega^{-1} \tilde{X}_i \right)^{-1} \sum_i \tilde{X}_i' \Omega^{-1} y_i$$

where $\Omega^{-\frac{1}{2}} y_i = \Omega^{-\frac{1}{2}} \tilde{X}_i' \beta + \Omega^{-\frac{1}{2}} v_i$.

$$\begin{aligned} \Omega &= \mathbb{E}[v_i v_i' | \tilde{X}_i] = \mathbb{E} \left[\begin{bmatrix} v_{i1} \\ v_{i2} \\ \vdots \\ v_{iT} \end{bmatrix} \begin{bmatrix} v_{i1} & v_{i2} & \cdots & v_{iT} \end{bmatrix} | \tilde{X}_i \right] \\ &= \mathbb{E} \begin{bmatrix} \mathbb{E}[v_{i1}^2 | \tilde{X}_i] & \mathbb{E}[v_{i1}v_{i2} | \tilde{X}_i] & \cdots & \mathbb{E}[v_{i1}v_{iT} | \tilde{X}_i] \\ \mathbb{E}[v_{i2}v_{i1} | \tilde{X}_i] & \mathbb{E}[v_{i2}^2 | \tilde{X}_i] & \cdots & \mathbb{E}[v_{i2}v_{iT} | \tilde{X}_i] \\ \vdots & \vdots & \ddots & \vdots \\ \mathbb{E}[v_{iT}v_{i1} | \tilde{X}_i] & \mathbb{E}[v_{iT}v_{i2} | \tilde{X}_i] & \cdots & \mathbb{E}[v_{iT}^2 | \tilde{X}_i] \end{bmatrix} \\ &= \begin{bmatrix} \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{i1}^2 | \tilde{X}_i] & \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{i1}u_{i2} | \tilde{X}_i] & \cdots & \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{i1}u_{iT} | \tilde{X}_i] \\ \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{i2}u_{i1} | \tilde{X}_i] & \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{i2}^2 | \tilde{X}_i] & \cdots & \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{i2}u_{iT} | \tilde{X}_i] \\ \vdots & \vdots & \ddots & \vdots \\ \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{iT}u_{i1} | \tilde{X}_i] & \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{iT}u_{i2} | \tilde{X}_i] & \cdots & \mathbb{E}[\alpha_i^2 | \tilde{X}_i] + \mathbb{E}[u_{iT}^2 | \tilde{X}_i] \end{bmatrix} \\ &= \begin{bmatrix} \sigma_u^2 + \sigma_\alpha^2 & \sigma_\alpha^2 & \cdots & \sigma_\alpha^2 \\ \sigma_\alpha^2 & \sigma_u^2 + \sigma_\alpha^2 & \cdots & \sigma_\alpha^2 \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_\alpha^2 & \sigma_\alpha^2 & \cdots & \sigma_u^2 + \sigma_\alpha^2 \end{bmatrix} \\ &= \sigma_\alpha^2 \mathbf{1}_i \mathbf{1}_i' + \sigma_u^2 I \\ &\text{because } \mathbb{V}[\tilde{\alpha}_i | \tilde{X}_i] = \sigma_{\alpha_i}^2 = \sigma_\alpha^2 \\ &\mathbb{V}[u_{it} | \tilde{X}_i] = \sigma_u^2, \forall i. \end{aligned}$$

where I is an identity matrix of dimension T_i . Under the assumption $\mathbb{E}[u_{it}u_{is}] = 0$, we now describe some statistical properties of $\hat{\beta}_{RE-GLS}$.

RE Consistency

$$\begin{aligned}
\hat{\beta}_{RE-GLS} - \tilde{\beta} &= \left(\sum_i \tilde{X}_i' \Omega^{-1} \tilde{X}_i \right)^{-1} \left(\sum_i \tilde{X}_i' \Omega^{-1} v_i \right) \\
&\rightarrow \mathbb{E} \left[\sum_i \tilde{X}_i' \Omega^{-1} \tilde{X}_i \right] \mathbb{E} \left[\sum_i \tilde{X}_i' \Omega^{-1} v_i \right] \\
\text{where } \mathbb{E} \left[\sum_i \tilde{X}_i' \Omega^{-1} v_i \right] &= \sum_i \mathbb{E} \left[\tilde{X}_i' \Omega^{-1} v_i \right] \\
&= \sum_i \tilde{X}_i' \Omega^{-1} \mathbb{E}[v_i | \tilde{X}_i] \\
&= \sum_i \tilde{X}_i' \Omega^{-1} \mathbb{E}[u_i + \tilde{\alpha}_i | \tilde{X}_i] \\
&= 0
\end{aligned}$$

Thus, $\hat{\beta}_{RE-GLS}$ is conditionally unbiased for $\tilde{\beta}$. The conditional variance of $\hat{\beta}_{RE-GLS}$ is:

$$\mathbb{V}[\hat{\beta}_{RE-GLS}] = \left(\sum_i \tilde{X}_i' \Omega^{-1} \tilde{X}_i \right)^{-1} \sigma_u^2$$

RE Asymptotic Distribution

The asymptotic variance of $\hat{\beta}_{RE-GLS}$ is:

$$\begin{aligned}
\sqrt{n} \left(\hat{\beta}_{RE-GLS} - \tilde{\beta} \right) &\xrightarrow{d} \mathcal{N}(0, V) \\
\text{where } V_{GLS} &= \mathbb{E} \left[\tilde{X}_i' \Omega^{-1} \tilde{X}_i \right]^{-1} \mathbb{E} \left[\tilde{X}_i' \Omega^{-1} v_i v_i' \Omega^{-1} \tilde{X}_i \right] \mathbb{E} \left[\tilde{X}_i' \Omega^{-1} \tilde{X}_i \right]^{-1} \\
&= \mathbb{E} \left[\tilde{X}_i' \Omega^{-1} \tilde{X}_i \right]^{-1} \underbrace{\mathbb{E}[v_i v_i' | \tilde{X}_i]}_{\equiv \Omega}
\end{aligned}$$

Because we do not know Ω , the RE-GLS estimator is infeasible.

The previous assumptions 3.2.4 and 3.2.5 are special to random effects. For efficiency of feasible GLS, we assume that the variance matrix of v_i conditional on \tilde{X}_i is constant:

$$\mathbb{E}[v_i v_i' | \tilde{X}_i] = \mathbb{E}[v_i v_i'].$$

The two conditions are also implied by the following stronger version of assumption:

Assumption 3.2.6 (RE General Homoskedasticity and Serial Uncorrelation).

- (a) $\mathbb{E}[u_i u_i' | \tilde{X}_i, \tilde{\alpha}_i] = 0$
- (b) $\mathbb{E}[\tilde{\alpha}_i^2 | \tilde{X}_i] = \sigma_\alpha^2$

Under assumption 3.2.6(a), $\mathbb{E}[u_{it}^2 | \tilde{X}_i, \tilde{\alpha}_i] = \sigma_u^2$, which implies assumption 3.2.4, and $\mathbb{E}[u_{it} u_{is} | \tilde{X}_i, \tilde{\alpha}_i] = 0$ which implies assumption 3.2.5. But this new assumption is stronger because it implies that the conditional variances are constant and the conditional covariances are zero.

Together with assumption 3.2.2(b), assumption 3.2.6(b) is the same as: $\mathbb{V}[\tilde{\alpha}_i | \tilde{X}_i] = \mathbb{V}[\tilde{\alpha}_i]$, which is a homoskedasticity assumption for the unobserved effects $\tilde{\alpha}_i$.

A feasible version replaces Ω with an estimator $\hat{\Omega}$. To implement an FGLS procedure, define:

$$\sigma_v^2 = \sigma_u^2 + \sigma_\alpha^2$$

then we can obtain: $\hat{\Omega} = \hat{\sigma}_\alpha^2 \mathbf{1}_i \mathbf{1}_i' + \hat{\sigma}_u^2 I_T$, a $T \times T$ matrix that we assume to be positive definite. In a panel data context, the FGLS estimator that uses this variance matrix is what is known as the **random effects estimator**.

$$\hat{\beta}_{RE} = \left(\sum_i \tilde{X}_i' \hat{\Omega}^{-1} \tilde{X}_i \right)^{-1} \sum_i \tilde{X}_i' \hat{\Omega}^{-1} y_i$$

Hence, the motivation for using GLS is different than under a cross-sectional regression with heteroskedasticity. We use GLS because of the autocorrelation in v_{it} induced by the presence of time variant α_i .

3.2.3 Comparing POLS and GLS

Now, let's compare the $\hat{\beta}_{RE-GLS}$ with the pooled estimator $\hat{\beta}_{POLS}$.

Under the assumptions of the random effects model, POLS estimator is also unbiased for β and has conditional variance:

$$V_{POLS} = \left(\sum_i X_i' X_i \right)^{-1} \left(\sum_i X_i' \Omega X_i \right) \left(\sum_i X_i' X_i \right)^{-1}$$

Using the algebra of the Gauss-Markov Theorem we deduce that:

$$V_{RE-GLS} \leq V_{POLS}$$

and thus the random effects estimator $\hat{\beta}_{RE-GLS}$ is more efficient than $\hat{\beta}_{POLS}$ under the strict exogeneity assumption 3.2.1. The two variance matrices are identical when there is no individual-specific effect $\sigma_\alpha^2 = 0$ for then $V_{RE-GLS} = V_{POLS} = (X'X)^{-1} \sigma_u^2$.

Under the assumption that the random effects model is a useful approximation but not literally true, we may use the cluster-robust covariance matrix estimator such as:

$$\hat{V}_{RE-GLS} = \left(\sum_i X_i' \Omega^{-1} X_i \right)^{-1} \left(\sum_i X_i' \Omega^{-1} \hat{v}_i \hat{v}_i' \Omega^{-1} X_i \right) \left(\sum_i X_i' \Omega^{-1} X_i \right)^{-1}$$

where $\hat{v}_i = y_i - X_i \hat{\beta}_{RE-GLS}$, This may be re-scaled by a degree of freedom adjustment if desired.

3.3 Fixed Effects

In the econometrics literature if the stochastic structure of α_i is treated as unknown and possibly correlated with x_{it} , then α_i is called a **fixed effect**.

Correlation between α_i and x_{it} will cause both pooled and random effect estimators to be biased. This is due to the classic problems of omitted variables bias and endogeneity.

The presence of the unstructured individual effect α_i means that it is not possible to identify β under a simple projection assumption such as $\mathbb{E}[u_{it} x_{it}] = 0$. It turns out that a sufficient condition for identification is the following.

Definition 3.3.1 (Strictly exogeneity).

A regressor x_{it} is said to be strictly exogeneity if $\mathbb{E}[x_{it} u_{is}] = 0, \forall t, s = 1, \dots, T$.

Strict exogeneity is a strong projection condition, meaning that is a $X_{is}, s \neq t$ is added into the regression

model, it would have a zero coefficient. Strict exogeneity is a projection analog of the **strict mean independence**:

$$\mathbb{E}[u_{it}|X_i] = 0$$

which implies the strict exogeneity but not vice versa.

The strict exogeneity assumption 3.3.1 is sufficient for identification and asymptotic theory, we'll also use the strict mean independence assumption for finite sample analysis.

Remark (About strict exogeneity[2]).

Strict exogeneity (assumption 3.3.1) is typically inappropriate in dynamic models.

3.3.1 Within Transformation

In previous steps, we showed that if x_{it} and α_i are correlated, then pooled OLS and RE-GLS estimator would be biased and inconsistent. If we leave the relationship between α_i and x_{it} fully unstructured, then the only way to consistently estimate the coefficient β is by an estimator which is invariant to α_i .

The first fixed effects (FE) assumption is strict exogeneity of the explanatory variables conditional on α_i :

Assumption 3.3.1 (FE Strict Exogeneity).

$$\mathbb{E}[u_{it}|X_i, \alpha_i] = 0, \forall t = 1, \dots, T$$

This assumption is identical to the assumption 3.2.1(a), we maintain strict exogeneity of $x_{it}, t = 1, \dots, T$ conditional on the unobserved effect. The key difference is that *we do not assume assumption 3.2.1(b), which means that, for FE analysis, $\mathbb{E}[\alpha_i|X_i]$ can be any function of X_i .*

By relaxing assumption 3.2.1(b), we can consistently estimate partial effects in the presence of time-consistent omitted variables that can be arbitrarily related to unobserved variables x_{it} . *Therefore, FE analysis is more robust than RE analysis.*

But this robustness has a cost: we can not include any time-constant variables in x_{it} without further assumptions. The reason is simple: if α_i can be arbitrarily correlated with each element of x_{it} , then there's no way to distinguish the effect of time-constant observables from the time-constant unobservable α_i .

We transform the equation to get rid of α_i : $y_{it} = \alpha_i + x'_{it}\beta + u_{it}$. The first transformation is the **within transformation**. Define the mean of a variable for a given individual as

$$\begin{aligned}\bar{y}_i &= \frac{1}{T} \sum_t y_{it} \\ \bar{x}_i &= \frac{1}{T} \sum_t x_{it} \\ \bar{u}_i &= \frac{1}{T} \sum_t u_{it}\end{aligned}$$

We call this the **individual-specific mean** since it is the mean of a given individual.¹

Then, subtracting the individual-specific mean from the variable we obtain the deviations:

$$\begin{aligned}(y_{it} - \bar{y}_i) &= (x_{it} - \bar{x}_i)' \beta + (u_{it} - \bar{u}_i) + (\alpha_i - \alpha_i) \\ \ddot{y}_{it} &= \ddot{x}'_{it} \beta + \ddot{u}_{it} \text{ or at individual level: } \ddot{y}_i = \ddot{X}'_i \beta + \ddot{u}_i\end{aligned}$$

¹Some authors call this the **time-average** or **time-mean** since it is the average over the time periods.

This is the **within transformation**. We also refer to \ddot{y}_{it} as the **demeaned values** or **deviation from individual means**. *What is important is that the demeaning has occurred at the individual level.*

Denote the time-averages method by $\hat{\beta}_{FE-W}$, in order to ensure that the FE estimator is consistent and well behaved asymptotically, we need a standard rank condition on the matrix of time-demeaned explanatory variables:

Assumption 3.3.2 (FE full rank).

$$\text{rank} \sum_t \mathbb{E}[\ddot{x}'_{it} \ddot{x}_{it}] = \text{rank} \mathbb{E}[\ddot{X}'_i \ddot{X}_i] = K$$

FE Consistency

$$\begin{aligned} \hat{\beta}_{FE-W} &= \left(\sum_i \ddot{X}'_i \ddot{X}_i \right)^{-1} \left(\sum_i \ddot{X}'_i \ddot{y}_i \right) \\ &= \left(\sum_i \sum_t \ddot{x}_{it} \ddot{x}'_{it} \right)^{-1} \sum_i \sum_t \ddot{x}_{it} \ddot{y}_{it} \\ &= \beta + \left(\sum_i \sum_t \ddot{x}_{it} \ddot{x}'_{it} \right)^{-1} \sum_i \sum_t \ddot{x}_{it} \ddot{u}_{it} \\ &\xrightarrow{p} \beta + \mathbb{E} \left[\sum_t \ddot{x}_{it} \ddot{x}'_{it} \right]^{-1} \mathbb{E} \left[\sum_t \ddot{x}_{it} \ddot{u}_{it} \right] \end{aligned}$$

where $\mathbb{E} \left[\sum_t \ddot{x}_{it} \ddot{u}_{it} \right] = \sum_t \mathbb{E} [\ddot{x}_{it} \ddot{u}_{it}]$

$$\begin{aligned} \mathbb{E} [\ddot{x}_{it} \ddot{u}_{it}] &= \mathbb{E} \left[\left(x_{it} - \frac{1}{T} \sum_t x_{it} \right) \left(u_{it} - \frac{1}{T} \sum_t u_{it} \right)' \right] \\ &= 0 \quad \text{if } u_{it} \perp\!\!\!\perp x_{is}, \forall t, s = 1, \dots, T. \end{aligned}$$

Then, let $\Sigma_i = \mathbb{E}[u_i u'_i | X_i]$ denote the $T_i \times T_i$ covariance matrix of the idiosyncratic errors. The variance of $\hat{\beta}_{FE-W}$ is:

$$V_{FE-W} = \mathbb{V}[\hat{\beta}_{FE-W} | X_i] = \left(\sum_i \ddot{X}'_i \ddot{X}_i \right)^{-1} \left(\sum_i \ddot{X}'_i \Sigma_i \ddot{X}_i \right) \left(\sum_i \ddot{X}'_i \ddot{X}_i \right)^{-1}$$

This expression simplifies when the idiosyncratic errors are homoskedastic and serially uncorrelated:

Assumption 3.3.3 (FE homoskedasticity and Serial Uncorrelation).

- (a) $\mathbb{E}[u_{it}^2 | X_i] = \sigma_u^2$
- (b) $\mathbb{E}[u_{it} u_{is} | X_i] = 0, \forall s \neq t.$

FE Asymptotic Distribution

In this case, $\Sigma_i = \sigma_u^2 I_i$ and V_{FE-W} simplifies to:

$$V_{FE-W}^0 = \sigma_u^2 \left(\sum_i \ddot{X}'_i \ddot{X}_i \right)^{-1}$$

We can also write the asymptotic distribution as below

$$\begin{aligned}
 \sqrt{n}(\hat{\beta}_{FE-W} - \beta) &= \left(\frac{1}{N} \sum_i \ddot{X}_i' \ddot{X}_i \right)^{-1} \left(\frac{1}{\sqrt{N}} \sum_i \ddot{X}_i' \ddot{u}_i \right) \\
 &= \left(\frac{1}{N} \sum_i \ddot{X}_i' \ddot{X}_i \right)^{-1} \left(\frac{1}{\sqrt{N}} \sum_i \ddot{X}_i' u_i \right)^2 \\
 &\rightarrow \mathbb{E}[\ddot{X}_i' \ddot{X}_i]^{-1} \cdot \mathcal{N}(0, \mathbb{V}[\ddot{X}_i' u_i]) \\
 &\sim \mathcal{N}(0, V_{FE-W}) \\
 \text{where } V_{FE-W} &= \sigma_u^2 \left(\sum_i \ddot{X}_i' \ddot{X}_i \right)^{-1}
 \end{aligned}$$

Remark (About FE Asymptotic Distribution).

Actually, the asymptotic distribution of $\hat{\beta}_{FE-W}$ is not as obvious as it seems. We have to restress a few assumptions to guarantee its validity:

$$\frac{1}{\sqrt{N}} \sum_i \ddot{X}_i' u_i \xrightarrow{d} \mathcal{N}(0, \mathbb{V}[\ddot{X}_i' u_i]).$$

Assumption 3.3.4.

- (1) Variables $(u_i, X_i), i = 1, \dots, N$ are independent and identically distributed.
- (2) $\mathbb{E}[X_{it}u_{is}] = 0, \forall t = 1, \dots, T$.
- (3) $Q_T = \mathbb{E}[\ddot{X}_i' \ddot{X}_i] > 0$.
- (4) $\mathbb{E}[u_{it}^4] < \infty$.
- (5) $\mathbb{E}[|X_{it}|^4] < \infty$.

Assumption 3.3.4(2) implies that:

$$\mathbb{E}[\ddot{X}_i u_i] = \sum_t \mathbb{E}[\ddot{X}_{it} u_{it}] = \sum_t \mathbb{E}[X_{it} u_{it}] - \sum_t \sum_{s=1}^T \mathbb{E}[X_{is} u_{it}] = 0$$

so they are mean zero. Assumption 3.3.4(4) and (5) imply that $\ddot{X}_i u_i$ has a finite covariance matrix Ω . The assumptions for the CLT hold, thus we have the result.

Remark (FE VS. POLS).

²From the regression model $\ddot{y}_i = \ddot{X}_i \beta + \ddot{u}_i$, where \ddot{y}_i is $T \times 1$, \ddot{X}_i is $T \times K$, and \ddot{u}_i is $T \times 1$, We can write the individual-specific mean as $\bar{y}_i = (\mathbf{1}_i' \mathbf{1}_i)^{-1} \mathbf{1}_i' y_i$. Then, we can define a **individual-specific demeaning operator**:

$$M_i = I_i - \mathbf{1}_i (\mathbf{1}_i' \mathbf{1}_i)^{-1} \mathbf{1}_i',$$

giving that

$$\ddot{y}_i = y_i - \mathbf{1}_i \bar{y}_i = y_i - \mathbf{1}_i (\mathbf{1}_i' \mathbf{1}_i)^{-1} \mathbf{1}_i' y_i = M_i y_i.$$

Notice that M_i is idempotent ($M_i M_i = M_i, M_i' = M_i$). Similarly for \ddot{X}_i and \ddot{u}_i . Thus, we have:

$$\ddot{X}_i' \ddot{u}_i = X_i' M_i M_i u_i = X_i' M_i u_i = \ddot{X}_i' u_i.$$

It is instructive to compare the variances of the fixed-effects and pooled estimators under

$$\begin{aligned}\mathbb{E}[u_{it}^2|X_i] &= \sigma_u^2 \\ \mathbb{E}[u_{it}u_{is}|X_i] &= 0, \forall s \neq t.\end{aligned}$$

and the assumption that there is no individual-specific effect, $\alpha_i = 0$. In this case, we can see that:

$$V_{FE-W}^0 = \sigma_u^2 \left(\sum_i \ddot{X}_i' \ddot{X}_i \right)^{-1} \geq \sigma_u^2 \left(\sum_i X_i' X_i \right)^{-1} = V_{POLS}.$$

The inequality holds since the demeaned variables \ddot{X}_i have reduced variation compared to the original observations X_i .

This shows the cost of using fixed effects relative to pooled estimation. The estimation variance increases due to reduced variation in the regressors. This reduction in efficiency is a necessary by-product of the robustness of the estimator to the individual effects α_i .

3.3.2 First Difference Transformation

Another important transformation which does the same as within transformation is **first-differencing**. *This can be applied to all but the first observation (which is essentially lost).*

$$\begin{aligned}y_{it} - y_{i,t-1} &= (x_{it} - x_{i,t-1})' \beta + (u_{it} - u_{i,t-1}) \\ \Delta y_{it} &= \Delta x_{it}' \beta + \Delta u_{it} \text{ or at individual level } \Delta y_i = \Delta X_i \beta + \Delta u_i, \quad i = 1 \dots N, t = 2 \dots T\end{aligned}$$

We can see that the individual effect α_i has been eliminated.

Denote the first difference method by $\hat{\beta}_{FE-FD}$, the fixed effect estimator is consistent and asymptotically normal based on two assumptions.

Assumption 3.3.5 (FD Strict exogeneity).

It's the same as FE's assumption 3.3.1.

Assumption 3.3.6 (FD Full rank).

$$\text{rank} \sum_{t=2}^T \mathbb{E}[\Delta x_{it}' \Delta x_{it}] = K$$

FE-FD Consistency

$$\begin{aligned}\hat{\beta}_{FE-FD} &= \left(\sum_i \sum_t \Delta x_{it} \Delta x_{it}' \right)^{-1} \sum_i \sum_t \Delta x_{it} \Delta y_{it} \\ &= \beta + \left(\frac{1}{NT} \sum_i \sum_t \Delta x_{it} \Delta x_{it}' \right)^{-1} \frac{1}{NT} \sum_i \sum_t \Delta x_{it} \Delta u_{it} \\ &\xrightarrow{p} \beta + \mathbb{E} \left[\sum_t \Delta x_{it} \Delta x_{it}' \right]^{-1} \mathbb{E} \left[\sum_t \Delta x_{it} \Delta u_{it} \right] \\ \text{where } \mathbb{E} \left[\sum_t \Delta x_{it} \Delta u_{it} \right] &= \sum_t \mathbb{E} [\Delta x_{it} \Delta u_{it}] \\ \mathbb{E} [\Delta x_{it} \Delta u_{it}] &= \mathbb{E} [(x_{it} - x_{i,t-1}) (u_{it} - u_{i,t-1})'] \\ &= 0 \quad \text{if } x_{it} \perp\!\!\!\perp (u_{it}, u_{i,t-1}), \forall t.\end{aligned}$$

For $T = 2$, $\hat{\beta}_{FE-FD} = \hat{\beta}_{FE-W}$, equals the fixed effects estimator and they differ however, for $T > 2$ (See Hanse, 2022[2]).

FE-FD Asymptotic Distribution

We just use the standard calculation:

$$\sqrt{n}(\hat{\beta}_{FE-FD} - \beta) \xrightarrow{d} \mathcal{N}(0, V_{FE-FD})$$

where

$$V_{FE-FD} = \mathbb{E} \left[\sum_{t=2}^T \Delta x_{it} \Delta x'_{it} \right]^{-1} \mathbb{E} \left[\left(\sum_{t=2}^T \Delta x_{it} \Delta u_{it} \right) \left(\sum_{s=2}^T \Delta x_{is} \Delta u_{is} \right)' \right] \mathbb{E} \left[\sum_{t=2}^T \Delta x_{it} \Delta x'_{it} \right]^{-1}$$

If we still assume that the first-difference error term Δu_{it} is homoskedastic:

Assumption 3.3.7 (FD homoskedasticity).

Denote $e_{it} \equiv \Delta u_{it}$, e_i is the stack of e_{it} for $t = 2, \dots, T$. $\mathbb{E}[e_i e_i' | X_i, \alpha_i] = \sigma_e^2 I$.

then, we can write:

$$AV[\hat{\beta}_{FE-FD}] = \hat{\sigma}_e^2 \left(\sum_i \Delta X_i' \Delta X_i \right)^{-1}$$

where $\hat{\sigma}_e^2$ is a consistent estimator of σ_e^2 , and the simplest estimator is obtained by computing the OLS residuals:

$$\hat{e}_{it} = \Delta y_{it} - \Delta x_{it} \hat{\beta}_{FE-FD}$$

from the pooled regression.

If the assumption 3.3.7 is violated, replacing expectations with sample means and $\Delta u_{it}(e_{it})$ with $\widehat{\Delta u_{it}}(\hat{e}_{it})$ yields the HAC-robust variance estimator \hat{V}_{FE-FD} .

$$\hat{V}_{FE-FD} = \left(\sum_i \Delta X_i' \Delta X_i \right)^{-1} \left(\sum_i \Delta X_i' \hat{e}_{it} \hat{e}_{it}' \Delta X_i \right) \left(\sum_i \Delta X_i' \Delta X_i \right)^{-1}$$

Remark (About FE-W and FE-FD (Hansen, 2022 [2])).

The FD method is not as strong as the within method, because it only requires that the variable is uncorrelated with the error term in the same period and the previous period.

If there is a correlation between the error term in current period and two periods ago, there is a problem of feedback loop, which we will imply the correlated random effect model.

Matrix Notation for FE-FD

The first-differencing transformation is $\Delta Y_{it} = Y_{it} - Y_{i,t-1}$. This can be applied to all but the first observation (which is essentially lost). At the level of the individual this can be written as:

$$\Delta Y_i = D_i Y_i$$

where D_i is the $(T_i - 1) \times T_i$ difference matrix differencing operator:

$$D_i = \begin{bmatrix} -1 & 1 & 0 & \cdots & 0 & 0 \\ 0 & -1 & 1 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \cdots & -1 & 1 \end{bmatrix}$$

Applying the transformation Δ to $Y_{it} = \alpha_i + X_{it}\beta + u_{it}$ or

$$\Delta Y_i = \Delta X_i \beta + \Delta u_i$$

We can see that the individual effect u_i has been eliminated.

Least squares applied to the differenced equation gives:

$$\begin{aligned} \hat{\beta}_\Delta &= \left(\sum_i \sum_{t=2}^T \Delta X'_{it} \Delta X_{it} \right)^{-1} \left(\sum_i \sum_{t=2}^T \Delta X'_{it} \Delta Y_{it} \right) \\ &= \left(\sum_i \Delta X'_i \Delta X_i \right)^{-1} \left(\sum_i \Delta X'_i \Delta Y_i \right) \\ &= \left(\sum_i X'_i D'_i D_i X_i \right)^{-1} \left(\sum_i X'_i D'_i D_i Y_i \right) \end{aligned}$$

is called the differenced estimator.

When the errors u_{it} are serially uncorrelated and homoskedastic, then the error $\Delta u_i = D_i u_i$ has the covariance matrix $H\sigma_u^2$, where

$$H = D_i D'_i = \begin{bmatrix} 2 & -1 & 0 & 0 \\ -1 & 2 & \ddots & 0 \\ 0 & \ddots & \ddots & -1 \\ 0 & 0 & -1 & 2 \end{bmatrix}$$

We can reduce estimation variance by using GLS. When errors are i.i.d. (serially uncorrelated and homoskedastic), this is:

$$\begin{aligned} \tilde{\beta}_\Delta &= \left(\sum_i \Delta X'_i H^{-1} \Delta X_i \right)^{-1} \left(\sum_i \Delta X'_i H^{-1} \Delta Y_i \right) \\ &= \left(\sum_i X'_i D'_i (D_i D'_i)^{-1} D_i X_i \right)^{-1} \left(\sum_i X'_i D'_i (D_i D'_i)^{-1} D_i Y_i \right) \\ &= \left(\sum_i X'_i M_i X_i \right)^{-1} \left(\sum_i X'_i M_i Y_i \right) \end{aligned}$$

where $M_i = D'_i (D_i D'_i)^{-1} D_i$. Recall that D_i is $(T_i - 1) \times T_i$ with rank $T_i - 1$ and is orthogonal to the vector of ones $\mathbf{1}_i$. This means that M_i projects orthogonally to $\mathbf{1}_i$ and thus equals the within transformation matrix. Hence $\tilde{\beta}_\Delta = \hat{\beta}_{FE-W}$.

What we have shown is that under i.i.d. errors, GLS applied to the first-differenced equation precisely equals the fixed effects estimator.

Since the Gauss-Markov theorem shows that GLS has lower variance than least squares, this means that the fixed effects estimator is more efficient than first differencing under the assumption that u_{it} is i.i.d.

3.3.3 Hausman Test for Random vs. Fixed Effects

Even if strict exogeneity is satisfied, the consistency of FE estimators comes at an efficiency loss compared to the RE-GLS and POLS estimators. This is easiest seen in the FD-transformation, in which we lose the n observations pertaining to the first time period $t = 1$.

The efficiency loss of the Within-estimator is somewhat more subtle. It arises because the Within-estimator only exploits variation across time and disregards the time-constant variation across cross-sectional units.

As a result, if the core RE assumption of X_i and α_i being uncorrelated is indeed satisfied, we prefer the RE-estimators. If instead it is violated, we of course prefer the less efficient but consistent FE estimators.

Theorem 3.3.1 (Hausman-Test).

$$\mathcal{H}_0: \hat{\beta}_{RE-GLS} - \hat{\beta}_{FE-W} = 0$$

We define:

$$T_{Hausman} = n(\hat{\beta}_{FE} - \hat{\beta}_{RE})' \left(A\mathbb{V}[\hat{\beta}_{FE}] - A\mathbb{V}[\hat{\beta}_{RE}] \right)^{-1} (\hat{\beta}_{FE} - \hat{\beta}_{RE}) \rightarrow \chi_k^2$$

If \mathcal{H}_0 is accepted, the difference between $\hat{\beta}_{RE-GLS}$ and $\hat{\beta}_{FE-W}$ is small enough to suggest that both estimators are consistent. that X_i and α_i are indeed uncorrelated, and therefore, we should use the more efficient estimator $\hat{\beta}_{RE-GLS}$.

If the test rejects this is evidence that the individual effect α_i is correlated with the regressors so the random effects model is not appropriate.

Remark (Random Effects or Fixed Effects?(Hansen, 2022[2])).

We have presented the random effects and fixed effects estimators of the regression coefficients. Which should be used in practice? How should we view the difference?

The basic distinction is that the random effects estimator requires the individual error $\tilde{\alpha}_i$ to satisfy the conditional mean assumption $\mathbb{E}[\tilde{\alpha}_i | \tilde{X}_i] = 0$. The fixed effects estimator does not require this condition, and is robust to its violation.

In particular, the individual effect $\tilde{\alpha}_i$ can be arbitrarily correlated to the regressors. On the other hand the random effects estimator is efficient under random effects.

Current econometric practice is to prefer robustness over efficiency. Consequently, current practice is (nearly uniformly) to use the fixed effects estimator for linear panel data models. Random effects estimators are only used in contexts where fixed effects estimation is unknown or challenging (which occurs in many nonlinear models).

The labels “random effects” and “fixed effects” are misleading. These are labels which arose in the early literature and we are stuck with these labels today. In a previous era regressors were viewed as “fixed”. Viewing the individual effect as an unobserved regressor leads to the label of the individual effect as “fixed”. Today, we rarely refer to regressors as “fixed” when dealing with observational data. We view all variables as random. Consequently describing α_i as “fixed” does not make much sense and it is hardly a contrast with the “random effect” label since under either assumption α_i is treated as random. Once again, the labels are unfortunate but the key difference is whether α_i is correlated with the regressors.

3.3.4 FE-IV Estimation

1. Contemporaneous exogeneity: $\mathbb{E}[x_{it}u_{it}] = 0, \forall t$.

2. Strict exogeneity: $\mathbb{E}[x_{it}u_{is}] = 0, \forall t, s$.
3. Sequential exogeneity: $\mathbb{E}[x_{it}u_{is}] = 0, \forall t, s \geq t$.

Definition 3.3.2 (Predetermined variables(Or Sequential Exogeneity)).

Predetermined variables are variables that were determined prior to the current period. In econometric models this implies that the current period error term is uncorrelated with current and lagged values of the predetermined variable but may be correlated with future values. This is a weaker restriction than strict exogeneity, which requires the variable to be uncorrelated with past, present, and future shocks.

The models we have discussed so far have been static with no dynamic relationships. In many economic contexts it is natural to expect that behavior and decisions are dynamic, explicitly depending on past behavior.

The workhorse dynamic model in a panel framework is the p -th order autoregression with regressors and a one-way error component structure(see 3.1.3). This is:

$$y_{it} = \alpha_1 y_{i,t-1} + \dots + \alpha_p y_{i,t-p} + x'_{it} \beta + \alpha_i + u_{it} \quad (3.2)$$

where α_j are the autoregressive coefficients. x_{it} is a k -vector of regressors, α_i is an individual effect and u_{it} is an idiosyncratic error. It's conventional to assume that u_{it} and α_i are mutually independent and the u_{it} are serially uncorrelated and mean zero. For the present we will assume that the regressors x_{it} are strictly exogenous(assumption 3.3.1). Currently, we focus on the AR(1) model:

$$y_{it} = \alpha_i + u_{it} + \beta_1 y_{i,t-1} + x'_{it} \beta_{-1}$$

where β_{-1} is a $k-1$ vector of coefficients on all other regressors.

Definition 3.3.3 (Anderson and Hsiao(1981)).

Anderson and Hsiao (1982) made an important breakthrough by showing that a simple instrumental variables estimator is consistent for the parameters of 3.2. The method first eliminates the individual effect α_i by first differencing:

$$\begin{aligned} y_{it} &= \alpha_i + x'_{it} \beta + u_{it} \\ &= \alpha_i + \beta_1 y_{i,t-1} + \tilde{x}'_{it} \beta_{-1} + u_{it} \\ \Rightarrow \Delta y_{it} &= \beta_1 \Delta y_{i,t-1} + \Delta x'_{it} \beta + \Delta u_{it} \end{aligned}$$

The challenge is that first-differencing induces correlation between $\Delta y_{i,t-1}$ and Δu_{it} :

$$\mathbb{E}[\Delta y_{i,t-1} \Delta u_{it}] = \mathbb{E}[(y_{i,t-1} - y_{i,t-2})(u_{it} - u_{i,t-1})] = -\sigma_u^2.$$

The other regressors are not correlated with Δu_{it} . For $s > 1$, $\mathbb{E}[\Delta y_{i,t-s} \Delta u_{it}] = 0$ and x_{it} is strictly exogenous $\mathbb{E}[\Delta x_{it} \Delta u_{it}] = 0$. The correlation between $\Delta y_{i,t-1}$ and Δu_{it} is endogeneity. One solution to endogeneity is to use an instrument. Anderson-Hsiao pointed out that $y_{i,t-2}$ is a valid instrument because it is correlated with $\Delta y_{i,t-1}$ yet uncorrelated with Δu_{it} . Under sequential exogeneity, instrument-exogeneity is satisfied: $\mathbb{E}[y_{i,t-2} \Delta u_{it}] = \mathbb{E}[y_{i,t-2} \Delta u_{it}] - \mathbb{E}[y_{i,t-2} \Delta u_{it-1}] = 0$.

This is the IV using the instruments $(y_{i,t-2}, \dots, y_{i,t-s-1})$ for $(\Delta y_{i,t-1}, \dots, \Delta y_{i,t-s})$. The estimator requires $T \geq s + 2$.

$$\mathbb{E}[y_{is} \Delta u_{it}] = 0, \forall s \leq t - 2.$$

The Anderson-Hsiao estimator is IV using $Y_{i,t-2}$ as an instrument for $\Delta Y_{i,t-1}$. Equivalently, this is IV using the instruments $(Y_{i,t-2}, \dots, Y_{i,t-p-1})$ for $(\Delta Y_{i,t-1}, \dots, \Delta Y_{i,t-p})$. The estimator requires $T \geq p + 2$.

To show that this estimator is consistent, for simplicity assume we have a balanced panel with $T = 3$, $p = 1$, and no regressors. In this case the Anderson-Hsiao IV estimator is

$$\hat{\alpha}_{iv} = \left(\sum_{i=1}^N Y_{i1} \Delta Y_{i2} \right)^{-1} \left(\sum_{i=1}^N Y_{i1} \Delta Y_{i3} \right) = \alpha + \left(\sum_{i=1}^N Y_{i1} \Delta Y_{i2} \right)^{-1} \left(\sum_{i=1}^N Y_{i1} \Delta \varepsilon_{i3} \right).$$

Under the assumption that ε_{it} is serially uncorrelated, (17.88) shows that $\mathbb{E}[Y_{i1} \Delta \varepsilon_{i3}] = 0$. In general, $\mathbb{E}[Y_{i1} \Delta Y_{i2}] \neq 0$. As $N \rightarrow \infty$

$$\hat{\alpha}_{iv} \xrightarrow{p} \alpha - \frac{\mathbb{E}[Y_{i1} \Delta \varepsilon_{i3}]}{\mathbb{E}[Y_{i1} \Delta Y_{i2}]} = \alpha.$$

Thus the IV estimator is consistent for α .

The Anderson-Hsiao IV estimator relies on two critical assumptions. First, the validity of the instrument (uncorrelatedness with the equation error) relies on the assumption that the dynamics are correctly specified so that ε_{it} is serially uncorrelated. For example, many applications use an AR(1). If instead the true model is an AR(2) then $Y_{i,t-2}$ is not a valid instrument and the IV estimates will be biased. Second, the relevance of the instrument (correlatedness with the endogenous regressor) requires $\mathbb{E}[Y_{i1} \Delta Y_{i2}] \neq 0$. This turns out to be problematic and is explored further in Section 17.40. These considerations suggest that the validity and accuracy of the estimator are likely to be sensitive to these unknown features.

Figure 3.1: Anderson and Hsiao(1981)

Using similar reasoning, other approaches use sequential exogeneity to circumvent FE methods altogether rather than to save their consistency. For example, Blundell and Bond (1998) start from the original specification:

$$y_{it} = x'_{it} \beta + \alpha_i + u_{it},$$

where correlation between α_i and x_{it} is suspected to be due to $y_{i,t-1}$, contained in x_{it} .

Definition 3.3.4 (Blundell and Bond(1998)).

$$\begin{aligned} y_{it} &= \alpha_i + \beta_1 y_{i,t-1} + u_{it} \\ &= \beta_1 y_{i,t-1} + (u_{it} + \alpha_i) \end{aligned}$$

Use $\Delta y_{i,t-1}$ as the IV for $y_{i,t-1}$

Univariate Time Series

4.1 Fundamentals of Time Series Analysis

A **time series** $Y_t \in \mathbb{R}^m$ is a process which is sequentially ordered over time. The time series is univariate if $m = 1$ and multivariate if $m > 1$.

Most economic time series are recorded at discrete intervals such as annual, quarterly, monthly, weekly, or daily. The number of observed periods s per year is called the **frequency**. In most cases we will denote the observed sample by the periods $t = 1, \dots, n$.

Recall that cross-sectional observations are conventionally treated as random draws from an underlying population. This is not an appropriate model for time series processes due to serial dependence. Instead, we treat the observed sample $\{y_1, \dots, y_n\}$ as a realization of a dependent stochastic process. It is often useful to view $\{y_1, \dots, y_n\}$ as a subset of an underlying doubly-infinite sequence $\{\dots, y_{t-1}, y_t, y_{t+1}, \dots\}$.

A random vector Y_t can be characterized by its distribution. A set such as $\{y_t, y_{t+1}, \dots, y_{t+l}\}$ can be characterized by its joint distribution. Important features of these distributions are their **means, variances, and covariances**.

Remark.

Time series theory is a mixture of probabilistic and statistical concepts. The probabilistic part is to study and characterize probability distributions of sets of variables y_t that will typically be dependent. The statistical problem is to determine the probability distribution of the time series given observations y_1, \dots, y_n at times $1, 2, \dots, n$. The resulting stochastic model can be used in two ways:

- understanding the stochastic system;
- predicting the “future”, i.e. y_{n+1}, y_{n+2}, \dots

As mentioned above, under time series data, we care about the joint distribution of random variable y_t . We give the following definitions of the mean, variance, and covariance of a random variable y_t .

Definition 4.1.1 (Mean function).

The mean function of a random variable y_t is defined as

$$\mu_t = \mathbb{E}[y_t] = \int y_t f_t(y_t) dy_t,$$

where $f_t(y_t)$ is the probability density function (PDF) of y_t .

Definition 4.1.2 (Autocovariance function).

The autocovariance function of a random variable y_t is defined as

$$\gamma_y(r, s) = \text{Cov}(y_r, y_s) = \mathbb{E}\left[(y_r - \mu_r)(y_s - \mu_s)\right].$$

Definition 4.1.3 (Autocorrelation function).

The autocorrelation function (ACF) of a random variable y_t is defined as

$$\rho_y(r, s) = \frac{\text{Cov}(y_r, y_s)}{\sqrt{\text{Var}(y_r) \text{Var}(y_s)}} = \frac{\gamma_y(r, s)}{\sqrt{\gamma_y(r, r) \gamma_y(s, s)}}.$$

4.1.1 Stationarity and Strict Stationarity**Definition 4.1.4** (Stationarity).

The time series $\{y_t, t \in \mathbb{Z}\}$, is said to be stationary if:

- (i) $\mathbb{E}[|y_t|^2] < \infty$, for all t ,
- (ii) $\mathbb{E}[y_t] = \mu$, for all t ,
- (iii) $\gamma_y(r, s) = \gamma_y(r + t, s + t)$, for all $r, s, t \in \mathbb{Z}$.

Remark (About Stationarity).

Stationarity as just defined is frequently referred to in the literature as weak stationarity, covariance stationarity, stationarity in the wide sense or second-order stationarity. For us however the term stationarity, without further qualification, will always refer to the properties specified by Definition 4.1.4, that is, when we say stationary, we mean weak stationarity.

If $\{y_t, t \in \mathbb{Z}\}$ is **stationary**, then $\gamma_y(r, s) = \gamma_y(r - s, 0)$ for all $r, s \in \mathbb{Z}$. It is therefore convenient to redefine the autocovariance function of a stationary process as the function of just one variable,

$$\gamma_y(h) \equiv \gamma_y(h, 0) = \text{Cov}(y_{t+h}, y_t), \quad \forall t, h \in \mathbb{Z}.$$

The function $\gamma_y(\cdot)$ will be referred to as the autocovariance function of $\{y_t\}$ and $\gamma_y(h)$ as its value at lag h . The autocorrelation function (ACF) of $\{y_t\}$ is defined analogously as the function whose value at lag h is given by

$$\rho_y(h) = \frac{\gamma_y(t + h, t)}{\sqrt{\gamma_y(t + h, t) \gamma_y(t, t)}} = \frac{\gamma_y(h)}{\gamma_y(0)} = \text{Corr}(y_{t+h}, y_t).$$

The auto-covariance and auto-correlation are functions $\gamma_y : \mathbb{Z} \rightarrow \mathbb{R}$ and $\rho_y : \mathbb{Z} \rightarrow [-1, 1]$. Together with the mean $\mu = \mathbb{E}[y_t]$, they determine the first and second moments of the stationary time series.

Definition 4.1.5 (Strict Stationarity).

The time series $\{y_t, t \in \mathbb{Z}\}$ is said to be strictly stationary if the joint distribution of $(y_{t_1}, y_{t_2}, \dots, y_{t_k})$ and $(y_{t_1+h}, y_{t_2+h}, \dots, y_{t_k+h})$ are the same for all h and k and t_1, \dots, t_k .

This is the natural generalization of the cross-section definition of identical distributions. Strict stationarity implies that the (marginal) distribution of y_t doesn't vary over time. It also implies that the bivariate distributions of (y_t, y_{t+1}) and multivariate distributions of (y_t, \dots, y_{t+l}) are stable over time.

The Relation Between Stationarity and Strict Stationarity

If $\{y_t\}$ is strict stationary, it immediately follows, on taking $k = 1$ in Definition 4.1.5, that y_t has the same distribution for all t . If $\mathbb{E}[|y_t|^2] < \infty$, this implies in particular that $\mathbb{E}[y_t]$ and $\mathbb{V}[y_t]$ are constant.

Moreover, taking $k = 2$, we find that y_{t+h} and y_t have the same joint distribution, and hence the same covariance for all h . Thus a strictly stationary process with finite second moments is stationary.

The converse of the previous statement is not true. For example if $\{y_t\}$ is a sequence of independent RVs s.t. y_t is exponentially distributed with mean 1 when t is odd and normally distributed with mean 1 and variance 1 when t is even, then $\{y_t\}$ is stationary with $\gamma_y(0) = 1$ and $\gamma_y(h) = 0$ for $h \neq 0$. However since y_1 and y_2 have different distributions, $\{y_t\}$ cannot be strictly stationary.

Note.

There is one important case however in which stationarity does imply strict stationarity.

Definition 4.1.6 (Gaussian Time Series).

The process $\{y_t\}$ is a Gaussian time series if and only if the distribution functions of $\{y_t\}$ are all multivariate normal.

If $\{y_t\}$ is a stationary Gaussian time series, then for all $k \in \{1, 2, \dots\}$ and for all h, t_1, t_2, \dots , the random vectors $(y_{t_1}, \dots, y_{t_k})'$ and $(y_{t_1+h}, \dots, y_{t_k+h})'$ have the same mean and covariance matrix, and hence the same distribution. $\{y_t\}$ is strictly stationary.

A straightforward but essential relationship is that an i.i.d. process is strictly stationary.

Theorem 4.1.1 (IID Strict Stationary).

If $\{y_t\}$ is an i.i.d. process, then $\{y_t\}$ is strictly stationary.

4.1.2 Transformations of Stationary Processes

One of the important properties of strict stationarity is that it is preserved by transformation. That is, transformations of strictly stationary processes are also strictly stationary. This includes transformations which include the full history of y_t .

Theorem 4.1.2 (transformation Invariance).

If $\{y_t\}$ is strict stationary, and $X_t = \phi(y_t, y_{t-1}, \dots) \in \mathbb{R}^q$ is a random vector then X_t is also strict stationary.

A transformation which includes the full past history is an infinite-order moving average. For scalar y and coefficients a_j define the vector process

$$X_t = \sum_{j=0}^{\infty} a_j y_{t-j}.$$

Many time-series models involve representations and transformations of this form.

This infinite series exists if it is convergent, meaning that the sequence $\sum_{j=0}^N a_j y_{t-j}$ has a finite limit as $N \rightarrow \infty$. Since the inputs y_t are random we define this as a probability limit.

Definition 4.1.7 (Convergence).

The infinite series $\{X_t\}$ converges almost surely if $\sum_{j=0}^N a_j y_{t-j}$ has a finite limit as $N \rightarrow \infty$ with probability 1. In this case we describe X_t as convergent.

Theorem 4.1.3 (Convergence-Stationary).

If $\{y_t\}$ is a strict stationary process, $\mathbb{E}[y] < \infty$ and $\sum_{j=0}^{\infty} |a_j| < \infty$, then X_t converges almost surely, and the process $\{X_t\}$ is strict stationary.

4.1.3 Ergodicity

Stationarity alone is not sufficient for the weak law of large numbers as there are strictly stationary processes with no time series variation.

Example 2.

If the stationary process is $y_t = Z$ for some RV Z . This is random but constant over all time. An implication is that the sample mean of $y_t = Z$ will be inconsistent for the population expectation.

We now motivate the concept of **ergodicity**¹. Conceptionally, this is more difficult to understand than the mean and variance. But it is a very helpful tool when analysing estimators. It allows one to simply replace the sample mean by its expectation without the need to evaluating a variance, which is extremely useful in some situations.

It can be difficult to evaluate the mean and variance of an estimator. Therefore, we may want an alternative form of convergence (instead of the mean squared error). To see whether this is possible we recall that for i.i.d random variables we have the very useful law of large numbers

$$\frac{1}{n} \sum_{t=1}^n y_t \xrightarrow{a.s.} \mathbb{E}[y_t] = \mu.$$

and in general $\frac{1}{n} \sum_{t=1}^n g(y_t) \xrightarrow{a.s.} \mathbb{E}[g(y)]$ if $\mathbb{E}[g(y)] < \infty$. Does such a result exist in time series? It does, but we require the slightly stronger condition that a time series is ergodic (which is a slightly stronger condition than the strictly stationary).

A useful intuition is that if y_t is ergodic then its sample paths will pass through all parts of the sample space never getting “stuck” in a subregion.

Definition 4.1.8 (Ergodicity: Formal Definition*).

Let (Ω, \mathcal{F}, P) be a probability space. A transformation $T : \Omega \rightarrow \Omega$ is said to be a measure preserving if for every set $A \in \mathcal{F}$, $P(T^{-1}(A)) = P(A)$. Moreover, it is said to be an ergodic transformation if $T^{-1}A = A$ implies $P(A) = 0$ or 1.

It is not obvious what this has to do with stochastic processes, but we attempt to make a link. Let us suppose that $y = \{y_t\}$ is a strictly stationary process defined on the probability space (Ω, \mathcal{F}, P) . By strict stationarity the transformation (shifting a sequence by one)

$$T(y_1, y_2, \dots) = (y_2, y_3, \dots)$$

is a measure preserving transformation. To understand ergodicity we define the set A , where

$$A = \{\omega : (y_1(\omega), y_0(\omega), \dots) \in H\} = \{\omega : (y_{-1}(\omega), y_{-2}(\omega), \dots) \in H\}.$$

The stochastic process is said to be ergodic, if the only sets which satisfies the above are such that $P(A) = 0$ or 1. Roughly, this means there cannot be too many outcomes ω which generate sequences

¹In the late 1800s, the physicist Ludwig Boltzmann needed a word to express the idea that if you took an isolated system at constant energy and let it run, any one trajectory, continued long enough, would be representative of the system as a whole. Being a highly-educated nineteenth century German-speaker, Boltzmann knew far too much ancient Greek, so he called this the “ergodic property”, from *ergon* “energy, work” and *hodos* “way, path”. The name stuck.

which ‘repeat’ itself (are periodic in some sense).

An equivalent definition is given later. From this definition it can be seen why ‘repeats’ are a bad idea. If a sequence repeats, the time average is unlikely to converge to the mean.

The definition of ergodicity, given above, is quite complex and is rarely used in time series analysis. However, one consequence of ergodicity is the ergodic theorem, which is extremely useful in time series.

Theorem 4.1.4 (Ergodic Theorem).

If $\{y_t\}$ is a strict stationary process, and $\mathbb{E}[y] = \mu < \infty$, then as $n \rightarrow \infty$:

$$\mathbb{E} \left[\frac{1}{n} \sum_{t=1}^n y_t - \mu \right] \rightarrow 0$$

and

$$\frac{1}{n} \sum_{t=1}^n y_t \xrightarrow{a.s.} \mu.$$

Proposition 4.1.1 (CLT for SS & Ergodic Processes).

If the process $\{y_t\}$ is strict stationary and ergodic, and $\mathbb{E}[y] = \mu < \infty$, $\mathbb{V}[y_t] = \sigma^2 < \infty$, and $\bar{\sigma}_T^2 = \mathbb{V}[\frac{1}{\sqrt{T}} \sum_{t=1}^T y_t] \xrightarrow{P} \bar{\sigma}^2 < \infty$, then

$$\frac{1}{\sqrt{T}} \sum_{t=1}^T y_t \xrightarrow{d} \mathcal{N}(\mu, \bar{\sigma}^2).$$

Definition 4.1.9 (Ergodicity: Plain Definition).

In general, for any shift τ_1, \dots, τ_k , and function $g : \mathbb{R}^{k+1} \rightarrow \mathbb{R}$, we have:

$$\frac{1}{n} \sum_{t=1}^n g(y_t, y_{t+\tau_1}, \dots, y_{t+\tau_k}) \xrightarrow{a.s.} \mathbb{E}[g(y, y_{t+\tau_1}, \dots, y_{t+\tau_k})].$$

This is mostly used as the definition of ergodicity, as it is an iff with the ergodic definition. This result generalises the strong law of large numbers (which shows almost sure convergence for i.i.d. random variables) to dependent random variables.

Definition 4.1.9 also gives us an idea of what constitutes an ergodic process. Suppose that $\{\varepsilon_t\}$ is an ergodic process (a classical example are i.i.d. random variables). Then any reasonable (meaning measurable) function of $\{y_t\}$ is also ergodic if y_t is defined as:

$$y_t = h(\dots, \varepsilon_{t-1}, \varepsilon_t, \varepsilon_{t+1}, \dots)$$

where $\{\varepsilon_t\}$ are i.i.d. RVs and h is a measurable function.

Remark.

Definition 4.1.9 is roughly equivalent to:

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{l=1}^n \mathbb{P}[A_l \cap B] = \mathbb{P}[A] \mathbb{P}[B],$$

which is a kind of asymptotic independence-on-average.

Many standard time series processes can be shown to be ergodic. A useful starting point is the observation that an i.i.d. sequence is ergodic.

Theorem 4.1.5 (IID Ergodic).

If $\{y_t\}$ is an i.i.d. process, then $\{y_t\}$ is strict stationary and ergodic.

Second, ergodicity, like stationarity, is preserved by transformation.

Theorem 4.1.6 (Ergodicity transformation Invariance).

If $\{y_t\} \in \mathbb{R}^m$ is strict stationary and ergodic, and $X_t = \phi(y_t, y_{t-1}, \dots) \in \mathbb{R}^q$ is a random vector, then X_t is also strict stationary and ergodic.

As an example, the infinite-order moving average transformation is ergodic if the input is ergodic and the coefficients are absolutely convergent.

Theorem 4.1.7.

If y_t is strict stationary and ergodic, $\mathbb{E}[y] < \infty$, and $\sum_{j=0}^{\infty} |a_j| < \infty$, then $X_t = \sum_{j=0}^{\infty} a_j y_{t-j}$ is strict stationary and ergodic.

We now present a useful property. It is that the Cesàro sum of the autocovariances limits to zero.

Theorem 4.1.8 (Cesàro Sum of Autocovariances).

If $\{y_t\} \in \mathbb{R}$ is strict stationary and ergodic, and $\mathbb{E}[y^2] < \infty$, then the Cesàro sum of the autocovariances converges to zero:

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{l=1}^n \text{Cov}[y_t, y_{t+l}] = 0.$$

4.1.4 Conditioning on Information Sets

In the past few sections we have introduced the concept of the infinite histories. We now consider conditional expectations given infinite histories.

Recall from probability theory that an **outcome** is an element of a sample space. An **event** is a set of outcomes.

Now we wish to define a conditional expectation given an infinite past history. Specifically, we wish to define

$$\mathbb{E}_{t-1}[y_t] = \mathbb{E}[y_t | y_{t-1}, y_{t-2}, \dots].$$

the expected value of y_t given the infinite past history $\tilde{y}_{t-1} = (y_{t-1}, y_{t-2}, \dots)$ up to time t . Intuitively, $\mathbb{E}_{t-1}[y_t]$ is the mean of the conditional distribution, the latter reflecting the information in the history.

Formally, in time series literature, we should follow the measure-theoretic approach to define $\mathbb{E}_{t-1}[y_t]$ as the conditional expectation given a σ -field:

$$\mathbb{E}_{t-1}[y_t] = \mathbb{E}[y_t | \mathcal{F}_{t-1}]$$

where \mathcal{F}_{t-1} is the σ -field generated by the infinite past history \tilde{y}_{t-1} .²

The σ -field generated by a random variable Y is the collection of measurable events involving Y . Similarly, the σ -field generated by an infinite history is the collection of measurable events involving this history.

²A σ -field is a collection of subsets of a sample space which is closed under complementation and countable unions.

Intuitively, \mathcal{F}_{t-1} contains all the information available in the history \tilde{y}_{t-1} . Consequently, economists typically call \mathcal{F}_{t-1} an **information set** rather than a σ -field.

We now describe some properties about information sets \mathcal{F}_t :

- (i) $\mathcal{F}_{t-1} \subset \mathcal{F}_t$. This means that information accumulates over time. Information is not lost.
- (ii) It is important to be precise about which variables are contained in the information set. For example, the information sets $\mathcal{F}_{1t} = \sigma(y_t, y_{t-1}, \dots)$ and $\mathcal{F}_{2t} = \sigma(y_t, x_t, y_{t-1}, x_{t-1}, \dots)$ are distinct even though they are both dated at time t .
- (iii) the conditional expectations (14.10) follow the law of iterated expectations and the conditioning theorem, thus:

$$\begin{aligned}\mathbb{E}[\mathbb{E}[y_t|\mathcal{F}_{t-1}|\mathcal{F}_{t-2}] &= \mathbb{E}[y_t|\mathcal{F}_{t-2}], \\ \mathbb{E}[\mathbb{E}[y_t|\mathcal{F}_{t-1}]] &= \mathbb{E}[y_t].\end{aligned}$$

and

$$\mathbb{E}[y_{t-1}y_t|\mathcal{F}_{t-1}] = y_{t-1}\mathbb{E}[y_t|\mathcal{F}_{t-1}].$$

4.1.5 Martingale Difference Sequences

An important concept in economics is unforecastability, meaning that the conditional expectation is the unconditional expectation. This is similar to the properties of a regression error. An unforecastable process is called a **martingale difference sequence (MDS)**.

A MDS y_t is defined with respect to a specific sequence of information sets \mathcal{F}_t . Most commonly, the latter are the natural filtration $\mathcal{F}_t = \sigma(y_t, y_{t-1}, \dots)$ (the past history of y_t) but it could be a larger information set. The only requirement is that y_t is adapted to \mathcal{F}_t , meaning that $\mathbb{E}[y_t|\mathcal{F}_t] = y_t$.

Definition 4.1.10 (Martingale Difference Sequence).

A process $\{y_t, \mathcal{F}_t\}$ is a martingale difference sequence (MDS) if y_t is adapted to \mathcal{F}_t , $\mathbb{E}[y_t] < \infty$:

$$\mathbb{E}[y_t|\mathcal{F}_{t-1}] = 0, \quad \forall t.$$

In words, a MDS y_t is unforecastable in the mean. If we apply the iterated expectations, $\mathbb{E}[y_t] = \mathbb{E}[\mathbb{E}[y_t|\mathcal{F}_{t-1}]] = 0$, thus a MDS is mean zero.

Note.

The term “martingale difference sequence” refers to the fact that the summed process $S_t = \sum_{j=1}^t y_j$, is a martingale and y_t is its first-difference.

Definition 4.1.11 (Martingale).

A process $\{S_t, \mathcal{F}_t\}$ is a martingale if S_t is adapted to \mathcal{F}_t , $\mathbb{E}[S_t] < \infty$:

$$\mathbb{E}[S_t|\mathcal{F}_{t-1}] = S_{t-1}, \quad \forall t.$$

Proposition 4.1.2 (I.I.D. and MDS).

If y_t is i.i.d. and mean zero, it is a MDS but the reverse is not the case.

Proof.

Suppose that y_t is i.i.d. and mean zero. It is then independent from $\mathcal{F}_{t-1} = \sigma(y_{t-1}, y_{t-2}, \dots)$, so $\mathbb{E}[y_t | \mathcal{F}_{t-1}] = \mathbb{E}[y_t] = 0$. Thus an i.i.d. shock is a MDS as claimed.

For the reverse, we let $u_t \sim \mathcal{N}(0, 1)$ is i.i.d. and set $y_t = u_t u_{t-1}$. By the conditioning theorem,

$$\mathbb{E}[y_t | \mathcal{F}_{t-1}] = \mathbb{E}[u_t u_{t-1} | \mathcal{F}_{t-1}] = u_{t-1} \mathbb{E}[u_t | \mathcal{F}_{t-1}] = 0.$$

So y_t is a MDS. However, y_t is not i.i.d., which can be shown by calculating the first autocovariance of y_t^2 :

$$\begin{aligned} \text{Cov}[y_t^2, y_{t-1}^2] &= \mathbb{E}[y_t^2 y_{t-1}^2] - \mathbb{E}[y_t^2] \mathbb{E}[y_{t-1}^2] \\ &= \mathbb{E}[u_t^2] \mathbb{E}[u_{t-1}^4] \mathbb{E}[u_{t-2}^2] - 1 \\ &= 2 \neq 0. \end{aligned}$$

Since the covariance is non-zero, y_t is not an independent sequence. Thus y_t is a MDS but not i.i.d. \square

Theorem 4.1.9 (Serial Uncorrelation and MDS).

If $\{y_t, \mathcal{F}_t\}$ is a MDS, and $\mathbb{E}[y_t^2] < \infty$, then y_t is serially uncorrelated.

Proof.

Take the process $y_t = u_t + u_{t-1} u_{t-2}$ with $u_t \sim \mathcal{N}(0, 1)$, i.i.d. The process is not MDS because $\mathbb{E}[y_t | \mathcal{F}_{t-1}] = u_{t-1} u_{t-2} \neq 0$, however,

$$\begin{aligned} \text{Cov}[y_t, y_{t-1}] &= \mathbb{E}[y_t y_{t-1}] \\ &= \mathbb{E}[(u_t + u_{t-1} u_{t-2})(u_{t-1} + u_{t-2} u_{t-3})] \\ &= \mathbb{E}[u_t u_{t-1} + u_t u_{t-2} u_{t-3} + u_{t-1}^2 u_{t-2} + u_{t-1} u_{t-2}^2 u_{t-3}] \\ &= \mathbb{E}[u_t] \mathbb{E}[u_{t-1}] + \mathbb{E}[u_t] \mathbb{E}[u_{t-2}] \mathbb{E}[u_{t-3}] + \mathbb{E}[u_{t-1}^2] \mathbb{E}[u_{t-2}] + \mathbb{E}[u_{t-1}] \mathbb{E}[u_{t-2}^2] \mathbb{E}[u_{t-3}] \\ &= 0 \end{aligned}$$

Similarly, $\text{Cov}[y_t, y_{t+k}] = 0$ for all $k \neq 0$. Thus the process is serially uncorrelated. \square

Definition 4.1.12 (Homoskedastic Martingale Difference Sequence).

A MDS $\{y_t, \mathcal{F}_t\}$ is said to be homoskedastic if

$$\mathbb{E}[y_t^2 | \mathcal{F}_{t-1}] = \sigma^2 < \infty, \quad \forall t.$$

where σ^2 is a constant.

Theorem 4.1.10 (MDS CLT).

If $\{y_t, \mathcal{F}_t\}$ is a martingale difference sequence, $\mathbb{E}[|y_t|^{2r}] < \infty$, and $\bar{\sigma}_T^2 = \mathbb{V} \left[\frac{1}{\sqrt{T}} \sum_{t=1}^T y_t \right] \rightarrow \bar{\sigma}^2 > 0$, then

$$\frac{1}{\sqrt{T}} \sum_{t=1}^T y_t \xrightarrow{d} \mathcal{N}(0, \bar{\sigma}^2).$$

4.2 ARMA Models

Classical regression is often insufficient for explaining all of the interesting dynamics of a time series. Instead, the introduction of correlation that may be generated through lagged linear relations leads to proposing the **autoregressive (AR)** and **autoregressive moving average (ARMA)** models that were presented in Whittle (1951) [4].

4.2.1 White Noise

In many respects the simplest kind of time series $\{X_t\}$ is one in which the random variables $X_t, t = 0, \pm 1, \pm 2, \dots$ are independently and identically distributed with zero mean and variance σ^2 . From a second order point of view i.e. ignoring all properties of the joint distributions of $\{X_t\}$ except those which can be deduced from the moments $\mathbb{E}[X_t]$ and $\mathbb{E}[X_s X_t]$, such processes are identified with the class of all stationary processes having mean zero and *autocovariance function*:

$$\gamma(h) = \text{Cov}[X_t, X_{t+h}] = \begin{cases} \sigma^2 & \text{if } h = 0, \\ 0 & \text{if } h \neq 0. \end{cases} \quad (4.1)$$

Definition 4.2.1 (White Noise).

The process u_t is called **white noise**, written as

$$u_t \sim WN(\mu, \sigma^2)$$

if and only if u_t has a mean of μ and covariance function 4.1. We call u_t strong WN if $u_t \perp u_{t-h}, \forall t, h$ (i.e. u_t and u_{t-h} are independent).

Usually, we deal with mean-zero WN: $u_t \sim WN(0, \sigma^2)$.

Note.

The distinction between (weak) WN and strong WN is only important for models where higher moments of u_t matter, e.g. conditional heteroskedasticity models.

By this definition 4.2.1, we can easily know that any mean-zero i.i.d. sequence with finite variances (MDS) is a white noise series. The converse is not true: there exist white noise series' that are not strictly stationary (MDS).

Therefore, the following types of shocks are nested: i.i.d., MDS, and white noise, with i.i.d. being the most narrow class and white noise the broadest.

If the WN process u_t is Normal, then it is strong WN, as linear independence implies independence for Normal RVs. We can simply write $u_t \stackrel{i.i.d.}{\sim} \mathcal{N}(0, \sigma^2)$. Note that the Normal WN process is SS and ergodic as it consists of i.i.d. RVs at each t .

Theorem 4.2.1 (Wold Decomposition).

Suppose that $\{y_t\}$ is *weak stationary*, with a finite variance (*we shall assume that it has mean zero, though this is not necessary*). Then y_t can be uniquely expressed as

$$y_t = \sum_{j=0}^{\infty} \phi_j u_{t-j} + V_t$$

where

1. $\{u_t\} \sim WN(0, \sigma^2)$;
2. $\phi_0 = 1$ and $\sum_{j=0}^{\infty} \phi_j < \infty$;
3. $\mathbb{E}[u_t V_s] = 0$ for all t, s ;
4. V_t is deterministic.

The Wold decomposition shows that y_t can be written as a linear function of the white noise projection errors (**a linear process** that consists of (infinitely many) WN components) plus a deterministic process. The Wold decomposition is a foundational result for linear time series analysis. Since any covariance stationary process can be written in this format this justifies linear models as approximations.

The square summability condition $\sum_{j=0}^{\infty} \phi_j^2 < \infty$ ensures that the coefficients approach zero sufficiently quickly so that the impact of any u_{t-j} on y_t diminishes as j increases. We refer to u_t as the innovation to the process y_t at time t .

Note that the Wold decomposition does not restrict the distributional family of u_t , nor does it exclude higher-order dependencies among u_t .

If u_t is i.i.d, then y_t is strictly stationary and ergodic as shown by our previous discussion of constituting ergodic process with definition 4.1.9.

4.2.2 ARMA(p, q) Models

A very wide class of stationary processes can be generated by using white noise as the forcing terms in a set of linear difference equations. This leads to the notion of an autoregressive-moving average (ARMA) process.

Definition 4.2.2 (ARMA(p, q) Process).

The process $\{y_t, t = 0, \pm 1, \pm 2, \dots\}$ is said to be an ARMA(p, q) process if y_t is stationary and for every t ,

$$y_t - \phi_1 y_{t-1} - \dots - \phi_p y_{t-p} = \alpha + \theta_1 u_{t-1} + \dots + \theta_q u_{t-q} + u_t \quad (4.2)$$

where $u_t \sim WN(0, \sigma^2)$, $\alpha = \mu(1 - \phi_1 - \dots - \phi_p)$, $\mathbb{E}[y_t] = \mu$, and $\phi_p, \theta_q \neq 0$.

The equation 4.2 can be written symbolically in the more compact form:

$$\phi(B)y_t = \alpha + \theta(B)u_t, t = 0, \pm 1, \pm 2, \dots$$

where ϕ and θ are p^{th} and q^{th} order polynomials in the **backshift operator** B :

$$\begin{aligned} \phi(B) &= 1 - \phi_1 B - \dots - \phi_p B^p, \\ \theta(B) &= 1 + \theta_1 B + \dots + \theta_q B^q \end{aligned}$$

where B is defined by:

$$B^j y_t = y_{t-j}, \quad j = 0, \pm 1, \pm 2, \dots$$

The polynomials $\phi(B)$ and $\theta(B)$ will be referred to as the **autoregressive** and **moving average** polynomials, respectively.

4.2.3 Autoregressive Models (AR(p))

Autoregressive models are based on the idea that the current value of the series, y_t , can be explained as a function of p past values: y_{t-1}, \dots, y_{t-p} , where p determines the number of steps into the past needed to forecast the current value.

If $\theta(B) \equiv 1$ in ARMA(p, q) 4.2, then we have the **autoregressive** model of order p (AR(p)):

Definition 4.2.3 (AR(p) Process).

The process $\{y_t, t = 0, \pm 1, \pm 2, \dots\}$ is said to be an AR(p) process if y_t is stationary and for every t ,

$$y_t - \phi_1 y_{t-1} - \dots - \phi_p y_{t-p} = \alpha + u_t \quad (4.3)$$

where $u_t \sim WN(0, \sigma^2)$, $\alpha = \mu(1 - \phi_1 - \dots - \phi_p)$, $\mathbb{E}[y_t] = \mu$, $\phi_p \neq 0$. This can be written symbolically in the more compact form:

$$\phi(B)y_t = \alpha + u_t, t = 0, \pm 1, \pm 2, \dots$$

where $\phi(B) = 1 - \phi_1 B - \dots - \phi_p B^p$.

Example 3 (AR(1) Process).

We initiate the investigation of AR models by considering the first-order model, AR(1), given by

$$y_t = \alpha + \phi_1 y_{t-1} + u_t, \quad u_t \sim WN(0, \sigma^2). \quad (4.4)$$

Iterating backwards k times, we get:

$$\begin{aligned} y_t &= \alpha + \phi y_{t-1} + u_t = \alpha + \phi(\alpha + \phi y_{t-2} + u_{t-1}) + u_t \\ &= \alpha + \phi\alpha + \phi^2 y_{t-2} + \phi u_{t-1} + u_t \\ &= \dots \\ &= \alpha(1 + \phi + \phi^2 + \dots + \phi^{k-1}) + \phi^k y_{t-k} + \sum_{j=0}^{k-1} \phi^j u_{t-j} \\ &= \frac{\alpha}{1 - \phi} + \phi^k y_{t-k} + \sum_{j=0}^{k-1} \phi^j u_{t-j}. \end{aligned}$$

This method suggests that, by continuing to iterate backward, and provided that $|\phi| < 1$, and $\mathbb{V}[y_t] < \infty$, we can represent an AR(1) model as a linear process given by

$$y_t = \frac{\alpha}{1 - \phi} + \sum_{j=0}^{\infty} \phi^j u_{t-j}. \quad (4.5)$$

Representation 4.5 is called the stationary solution of the model. In fact, by simple substitution,

$$\underbrace{\sum_{j=0}^{\infty} \phi^j u_{t-j}}_{y_t} = \frac{\alpha}{1 - \phi} + \phi \underbrace{\left(\sum_{k=0}^{\infty} \phi^k u_{t-k-1} \right)}_{y_{t-1}} + u_t$$

The AR(1) process defined by 4.5 is stationary with mean

$$\mathbb{E}[y_t] = \mathbb{E}\left[\frac{\alpha}{1-\phi}\right] + \sum_{j=0}^{\infty} \phi^j \mathbb{E}[u_{t-j}] = \frac{\alpha}{1-\phi}$$

and autocovariance function,

$$\begin{aligned} \gamma(h) &= \text{Cov}[y_{t+h}, y_t] = \mathbb{E}\left[\left(\sum_{j=0}^{\infty} \phi^j u_{t+h-j}\right) \left(\sum_{k=0}^h \phi^k u_{t-k}\right)\right] \\ &= \mathbb{E}\left[(u_{t+h} + \dots + \phi^h u_t + \phi^{h+1} u_{t-1} + \dots)(u_t + \phi u_{t-1} + \dots)\right] \\ &= \sigma^2 \sum_{j=0}^{\infty} \phi^{h+j} \phi^j \\ &= \sigma^2 \phi^h \sum_{j=0}^{\infty} \phi^{2j} \\ &= \frac{\sigma^2 \phi^h}{1-\phi^2}, \quad h \geq 0. \end{aligned}$$

Recall that $\gamma(h) = \gamma(-h)$, so we will only exhibit the autocovariance function for $h \geq 0$. We know it that the ACF of an AR(1) process is given by

$$\rho(h) = \frac{\gamma(h)}{\gamma(0)} = \phi^h, \quad h \geq 0.$$

and $\rho(h)$ satisfies the recursion:

$$\rho(h) = \phi \rho(h-1), \quad h \geq 1.$$

So, under $|\phi| < 1$, the AR(1) process is stable, since the effect of u_{t-h} on y_t dies out as $h \rightarrow \infty$. If instead $|\phi| > 1$, then this effect diverges to infinity and we have an exploding process. If $\phi = 1$, provided that $\alpha = 0$, y_t is simply the sum of past innovations u_t : $y_t = \sum_{h=0}^{\infty} u_{t-h}$. In that case, we call y_t a unit-root or random walk process (with drift if $\alpha \neq 0$). Under $|\phi| \geq 1$, the process is non-stationary and the ACF diverges to infinity.

Extra material Needed

Example 4 (AR(2) Process).

The AR(2) process is given by

$$y_t = \alpha + \phi_1 y_{t-1} + \phi_2 y_{t-2} + u_t, \quad u_t \sim WN(0, \sigma^2). \quad (4.6)$$

This process can also be written as

$$\phi(L)y_t = \alpha + u_t$$

where $\phi(L) = 1 - \phi_1 L - \phi_2 L^2$. Let $\frac{1}{\alpha_1}$ and $\frac{1}{\alpha_2}$ be the roots of $\phi(L) = 0$, we can rewrite the AR(2) model as:

$$\alpha + u_t = (1 - \alpha_1 L)(1 - \alpha_2 L)y_t$$

We would like to find the conditions for the stationarity of y_t . It turns out that it's convenient to transform the AR(2) model 4.6 into a VAR(1) process. Set $\tilde{y}_t = (y_t, y_{t-1})'$, which is stationary if

and only if y_t is stationary. Then equation 4.6 implies that:

$$\begin{bmatrix} y_t \\ y_{t-1} \end{bmatrix} = \begin{bmatrix} \alpha \\ 0 \end{bmatrix} + \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} y_{t-1} \\ y_{t-2} \end{bmatrix} + \begin{bmatrix} u_t \\ 0 \end{bmatrix}$$

or, in matrix form:

$$\tilde{y}_t = \mathbf{A}\tilde{y}_{t-1} + \tilde{u}_t$$

where $\mathbf{A} = \begin{bmatrix} \phi_1 & \phi_2 \\ 1 & 0 \end{bmatrix}$ and $\tilde{u}_t = \begin{bmatrix} \alpha + u_t \\ 0 \end{bmatrix}$. The VAR(1) process is strictly stationary and ergodic if the innovations satisfy $\mathbb{E}[\tilde{u}_t] < \infty$ and all the eigenvalues λ of \mathbf{A} are less than 1 in absolute value. The eigenvalues of \mathbf{A} are the roots of the characteristic polynomial:

$$\lambda^2 - \phi_1\lambda - \phi_2 = \lambda^2\phi\left(\frac{1}{\lambda}\right) = 0$$

and $\phi(z) = 1 - \phi_1z - \phi_2z^2$ is the autoregressive polynomial of the AR(2) model. By factoring this polynomial as $\phi(z) = (1 - \lambda_1z)(1 - \lambda_2z)$, we can find the roots λ_1 and λ_2 . The quadratic formula shows that these equal

$$\lambda_j = \frac{\phi_1 \pm \sqrt{\phi_1^2 + 4\phi_2}}{2}. \quad (4.7)$$

These eigenvalues are real if $\phi_1^2 + 4\phi_2 \geq 0$ and are complex conjugates otherwise. The AR(2) process is stationary if the solutions $|\lambda_j| < 1$ for $j = 1, 2$.

Theorem 4.2.2 (AR(2) Stationarity).

If $\mathbb{E}[e_t] < \infty$, and $|\lambda_j| < 1$ for λ_j defined in (4.7), then the AR(2) process is stationary, or equivalently, the following inequalities hold:

$$\phi_1 + \phi_2 < 1 \quad (4.8)$$

$$\phi_2 - \phi_1 < 1 \quad (4.9)$$

$$\phi_2 > -1 \quad (4.10)$$

then, the AR(2) process (4.6) is absolutely convergent, strictly stationary, and ergodic.

4.2.4 Moving Average Models (MA(q))

As an alternative to the autoregressive representation in which the y_t on the left-hand side of the equation are assumed to be combined linearly, the moving average model of order q , abbreviated as MA(q), assumes the white noise u_t on the right-hand side of the defining equation are combined linearly to form the observed data.

If $\phi(B) \equiv 1$ in ARMA(p, q) 4.2, then we have the **moving average** model of order q (MA(q)):

Definition 4.2.4 (MA(q) Process).

The process $\{y_t, t = 0, \pm 1, \pm 2, \dots\}$ is said to be an MA(q) process if y_t is stationary and for every t ,

$$y_t = \alpha + u_t + \theta_1 u_{t-1} + \dots + \theta_q u_{t-q} \quad (4.11)$$

where $u_t \sim WN(0, \sigma^2)$, $\alpha = \mu$, $\mathbb{E}[y_t] = \mu$, and $\theta_q \neq 0$. This can be written symbolically in the more

compact form:

$$y_t = \alpha + \theta(B)u_t, t = 0, \pm 1, \pm 2, \dots$$

where $\theta(B) = 1 + \theta_1 B + \dots + \theta_q B^q$.

It's quite easy to see that the ‘**linear process**’ we defined in 4.2.1 is $MA(\infty)$ and the white noise process is $MA(0)$. We can obtain $\mathbb{E}[y_t] = \alpha$ and:

$$\begin{aligned} \gamma(h) &= \text{Cov}[y_{t+h}, y_t] = \text{Cov}[u_{t+h} + \theta_1 u_{t+h-1} + \dots + \theta_q u_{t+h-q}, u_t + \theta_1 u_{t-1} + \dots + \theta_q u_{t-q}] \\ &= \mathbb{E}[(u_t + \theta_1 u_{t-1} + \dots)(u_{t+h} + \theta_1 u_{t+h-1} + \dots)] \\ &= \begin{cases} \sigma^2 \sum_{j=0}^{q-|h|} \theta_j \theta_{j+|h|}, & |h| < q; \\ 0, & |h| > q. \end{cases} \end{aligned}$$

Example 5 (MA(1) Process).

The MA(1) process is given by

$$y_t = \alpha + u_t + \theta u_{t-1}, \quad u_t \sim WN(0, \sigma^2). \quad (4.12)$$

The model is called a “moving average” because y_t is a weighted average of the shocks u_t and u_{t-1} . We can obtain that $\mathbb{E}[y_t] = \alpha$ and:

$$\begin{aligned} \gamma(h) &= \text{Cov}[y_{t+h}, y_t] = \text{Cov}[u_{t+h} + \theta u_{t+h-1}, u_t + \theta u_{t-1}] \\ &= \mathbb{E}[(u_t + \theta u_{t-1})(u_{t+h} + \theta u_{t+h-1})] \\ &= \begin{cases} (1 + \theta^2)\sigma^2, & \text{if } h = 0; \\ \theta\sigma^2, & \text{if } h = \pm 1; \\ 0, & \text{otherwise.} \end{cases} \end{aligned}$$

Thus y_s and y_t are uncorrelated whenever s and t are two or more times instants apart. Therefore, the MA(1) process is weak stationary regardless of the value of θ . We speak of *short range dependence* and say that the time series has *short memory*.

4.3 Inference of Univariate Time Series Models

$$\begin{aligned} y_t &\sim \text{WS} \Leftrightarrow y_t \sim I(0) \\ \Delta y_t &\sim \text{WS} \Leftrightarrow y_t \sim I(1) \\ \underbrace{\Delta y_t - \Delta y_{t-1}}_{\Delta^2 y_t} &\sim \text{WS} \Leftrightarrow y_t \sim I(2) \end{aligned}$$

4.3.1 Estimation of AR(p) Models

Example 6.

Let's consider the following AR(1) model:

$$y_t = \alpha + \beta y_{t-1} + u_t = x_t' \phi + u_t$$

$$x_t = \begin{bmatrix} 1 \\ y_{t-1} \end{bmatrix}, \phi = \begin{bmatrix} \phi_0 \\ \phi_1 \end{bmatrix}$$

where u_t is a white noise process

As u_t is the white noise, $\mathbb{E}[x_t u_t] = 0$, and that x_t is a function of u_{t-1}, u_{t-2}, \dots . If u_t is strict white noise, e.g. $u_t \sim i.i.d. \mathcal{N}(0, \sigma^2)$, then $\mathbb{E}[u_t | x_t] = \mathbb{E}[u_t] = 0$.

We consider estimation of an AR(p) model for stationary, ergodic, and non-deterministic y_t . The model is where $x_t = [1, y_{t-1}, \dots, y_{t-p}]'$, $\phi = [\phi_0, \phi_1, \dots, \phi_p]'$ and $u_t \sim WN(0, \sigma^2)$.

OLS estimator

The OLS estimator is given by:

$$\begin{aligned} \hat{\phi} &= \underset{\phi}{\operatorname{argmin}} \sum_{t=1}^T u_t^2 \\ U'U &= (Y - X\phi)'(Y - X\phi) \\ \Rightarrow \hat{\phi} &= (X'X)^{-1} X'Y = \left(\sum_{t=1}^n x_t x_t' \right)^{-1} \left(\sum_{t=1}^n x_t y_t \right). \end{aligned}$$

This notation presumes that there are $n + p$ total observations on y_t from which the first p are used as initial conditions so that $x_1 = (1, y_0, y_{-1}, \dots, y_{p+1})$ is defined.

Consistency of OLS Estimator

The OLS residuals are $\hat{u}_t = y_t - x_t' \hat{\phi}$. The error variance can be estimated by $\hat{\sigma}^2 = \frac{1}{n} \sum_{t=1}^n \hat{u}_t^2$ or $s^2 = \frac{1}{n-p-1} \sum_{t=1}^n \hat{u}_t^2$. For the true model, if y_t is strictly stationary (SS), and ergodic (E), then so are $x_t x_t'$ and $x_t y_t$. They have finite means $\mathbb{E}[y_t^2] < \infty$. Under these assumptions the Ergodic Theorem implies that

$$\begin{aligned} \frac{1}{n} \sum_{t=1}^n x_t x_t' &\xrightarrow{P} \mathbb{E}[x_t x_t'] = Q \\ \frac{1}{n} \sum_{t=1}^n x_t y_t &\xrightarrow{P} \mathbb{E}[x_t y_t] \end{aligned}$$

and we know that OLS is consistent for ϕ if x_t is uncorrelated with u_t .

$$\begin{aligned} \hat{\phi} &= (X'X)^{-1} X'Y \\ &= (X'X)^{-1} X'(X\phi + U) \\ &= \phi + (X'X)^{-1} X'U \\ &= \phi + \left(\frac{1}{T} \sum_t x_t x_t' \right)^{-1} \frac{1}{T} \sum_{t=1}^T x_t u_t \\ &\xrightarrow{P} \phi + \mathbb{E}[x_t x_t'] \mathbb{E}[x_t u_t] \\ &= \phi \end{aligned}$$

It is straightforward to show that $\hat{\sigma}^2$ is consistent as well.

Theorem 4.3.1 (AR(p) OLS Consistency).

If y_t is strictly stationary, ergodic, not purely deterministic, and $\mathbb{E}[y_t^2] < \infty$, then the OLS estimator $\hat{\phi}$ is consistent for ϕ , and $\hat{\sigma}^2 \xrightarrow{p} \sigma^2$ as $n \rightarrow \infty$.

Note.

This shows that under very mild conditions the coefficients of an AR(p) model can be consistently estimated by least squares. Once again, this does not require that the series y_t is actually an AR(p) process. It holds for any stationary process with the coefficient defined by projection.

Asymptotic Normality of OLS Estimator

The asymptotic distribution of the least squares estimator $\hat{\phi}$ depends on the stochastic assumptions. Specifically, we assume that the error u_t is a MDS (definition 4.1.10). An important implication is that since $x_t = (1, y_{t-1}, \dots, y_{t-p})'$ is part of the information set \mathcal{F}_{t-1} , by the conditioning theorem,

$$\mathbb{E}[x_t u_t | \mathcal{F}_{t-1}] = x_t \mathbb{E}[u_t | \mathcal{F}_{t-1}] = 0.$$

Thus $x_t u_t$ is a MDS. It has a finite variance if u_t has a finite fourth moment. We know that y_t can be written in MA(∞) form:

$$y_t = \mu + \sum_{j=0}^{\infty} \phi_j u_{t-j}$$

with $\sum_{j=0}^{\infty} |\phi_j| < \infty$. Using Minkowski's Inequality,

$$(\mathbb{E}|y_t|^4)^{\frac{1}{4}} \leq \sum_{j=0}^{\infty} |\phi_j| (\mathbb{E}|u_{t-j}|^4)^{\frac{1}{4}} < \infty.$$

Thus $\mathbb{E}[y_t^4] < \infty$. The Cauchy-Schwarz inequality then shows that $\mathbb{E}\|x_t u_t\|^2 < \infty$. We can then apply the MDS CLT 4.1.10 to show that:

$$\frac{1}{\sqrt{n}} \sum_{t=1}^n x_t u_t \xrightarrow{d} \mathcal{N}(0, \Sigma)$$

where $\Sigma = \mathbb{E}[x_t x_t' u_t^2]$.

Theorem 4.3.2 (MDS AR(p) Asymptotic Normality).

If y_t follows the AR(p) model, all roots of $\phi(z)$ lie outside the unit circle, (all eigenvalues have absolute value smaller than 1), $\mathbb{E}[u_t | \mathcal{F}_{t-1}] = 0$, $\mathbb{E}[u_t^4] < \infty$, and $\mathbb{E}[u_t^2] > 0$, then as $n \rightarrow \infty$, the OLS estimator $\hat{\phi}$ is asymptotically normal:

$$\sqrt{n} (\hat{\phi} - \phi) \xrightarrow{d} \mathcal{N}(0, V)$$

where $V = Q^{-1} \Sigma Q^{-1}$. ($Q = \mathbb{E}[x_t x_t']$)

This is identical in form to the asymptotic distribution of least squares in cross-section regression. The implication is that asymptotic inference is the same. In particular, the asymptotic covariance matrix is estimated just as in the cross-section case.

Theorem 4.3.3 (MDS AR(p) Asymptotic Normality(continue)).

If we further assume that $\mathbb{E}[u_t^2 | \mathcal{F}_{t-1}] = \sigma^2$, then as $n \rightarrow \infty$,

$$\sqrt{n}(\hat{\phi} - \phi) \xrightarrow{d} \mathcal{N}(0, V)$$

where $V = \sigma^2 Q^{-1}$.

These results show that under correct specification (a MDS error) the format of the asymptotic distribution of the least squares estimator exactly parallels the cross-section case.

If the AR(p) model holds WN errors, as we originally assumed ($u_t \sim WN(0, \sigma^2)$), the MDS CLT doesn't hold³. Instead, if y_t is strong mixing, we can use the central limit theorem for mixing processes.

Theorem 4.3.4 (Mixing AR Asymptotic Normality).

Assume y_t is strictly stationary, ergodic, and for some $r > 4$, $\mathbb{E}|y_t|^r < \infty$ and the mixing coefficients $\sum_{l=1}^{\infty} \phi(l)^{1-\frac{r}{4}} < \infty$. Let ϕ be defined as the best linear projection coefficients, $\phi = (\mathbb{E}[x_t x_t'])^{-1} \mathbb{E}[x_t y_t]$ from an AR(p) model, with projection errors u_t . Let $\hat{\phi}$ be the OLS estimator of ϕ .

$$\Omega = \sum_{l=-\infty}^{\infty} \mathbb{E}[x_{t-l} x_t' u_t u_{t-l}]$$

is convergent and

$$\sqrt{n}(\hat{\phi} - \phi) \xrightarrow{d} \mathcal{N}(0, V)$$

where $V = Q^{-1} \Omega Q^{-1}$.

To simplify our case, we can assume that our model is with strict WN $u_t \sim \mathcal{N}(0, \sigma^2)$. We can thus have the following result using theorem 4.3.4:

$$\begin{aligned} \sqrt{T}(\hat{\phi} - \phi) &= \sqrt{T} \left(\frac{1}{T} \sum_{t=1}^T x_t u_t \right)' \left(\frac{1}{T} \sum_{t=1}^T x_t x_t' \right)^{-1} \\ &\xrightarrow{d} \mathbb{E}[x_t x_t']^{-1} \mathcal{N} \left(0, \mathbb{V} \left[\frac{1}{\sqrt{T}} \sum_{t=1}^T x_t u_t \right] \right) \end{aligned}$$

where

$$\begin{aligned} \frac{1}{T} \mathbb{V} \left[\sum_{t=1}^T x_t u_t \right] &= \frac{1}{T} \sum_t \mathbb{V}[x_t u_t] + \frac{1}{T} \sum_t \sum_{\tau \neq t} \text{Cov}(x_t u_t, x_\tau u_\tau) \\ &= \mathbb{E}[u_t^2 x_t x_t'] + \frac{1}{T} \sum_t \sum_{\tau \neq t} \mathbb{E}[u_t u_\tau] x_t x_\tau' \end{aligned}$$

$$\begin{aligned} \text{Assume that } u_t \sim \mathcal{N}(0, \sigma^2) &= \mathbb{E}[u_t^2] \mathbb{E}[x_t x_t'] + \frac{1}{T} \sum_t \sum_{\tau} \mathbb{E}[u_t] \mathbb{E}[u_\tau] \mathbb{E}[x_t x_\tau'] \\ &= \sigma^2 \mathbb{E}[x_t x_t'] \end{aligned}$$

$$\Rightarrow \sqrt{T}(\hat{\phi} - \phi) \xrightarrow{d} \mathcal{N}(0, Q^{-1} \sigma^2 Q Q^{-1}) = \mathcal{N} \left(0, \underbrace{\sigma^2 \mathbb{E}[x_t x_t']^{-1}}_V \right)$$

³As we have previously shown under the definition 4.2.1, there exist white noises that are not strictly stationary, hence not MDS.

For the estimators:

$$\hat{V} = \hat{\sigma}^2 \left(\frac{1}{T} \sum_{t=1}^T x_t x_t' \right)^{-1}$$

$$\hat{\sigma}^2 = \frac{1}{T} \sum_{t=1}^T \hat{u}_t^2$$

like with cross sectional data.

Unlike the cross-sectional case, x_t consists of lagged values of y_t , we can also use the moments implied by our supposed model to estimate $\mathbb{E}[x_t x_t']$. In addition, if we consider $AR(1)$ model, where we have $x_t = [1, y_{t-1}]'$:

$$\mathbb{E}[x_t x_t'] = \mathbb{E} \left[\begin{bmatrix} 1 \\ y_{t-1} \end{bmatrix} \begin{bmatrix} 1 & y_{t-1} \end{bmatrix} \right] = \mathbb{E} \begin{bmatrix} 1 & y_{t-1} \\ y_{t-1} & y_{t-1}^2 \end{bmatrix} = \begin{bmatrix} 1 & \frac{\phi_0}{1-\phi_1} \\ \frac{\phi_0}{1-\phi_1} & \mathbb{E}[y_t^2] \end{bmatrix}$$

where $\mathbb{E}[y_t] = \frac{\phi_0}{1-\phi_1}$, $\mathbb{V}[y_t] = \frac{\sigma^2}{1-\phi_1^2}$, $\mathbb{E}[y_t^2] = \mathbb{V}[y_t] + \mathbb{E}[y_t]^2 = \frac{\sigma^2}{1-\phi_1^2} + \frac{\phi_0^2}{(1-\phi_1)^2}$

For the estimator,

$$\hat{\mathbb{E}}[x_t x_t'] = \begin{bmatrix} 1 & \frac{\hat{\phi}_0^2}{1-\hat{\phi}_1} \\ \frac{\hat{\phi}_0^2}{1-\hat{\phi}_1} & \frac{\hat{\sigma}^2}{1-\hat{\phi}_1^2} + \frac{\hat{\phi}_0^2}{(1-\hat{\phi}_1)^2} \end{bmatrix}$$

MLE estimator

$$\hat{\phi}_{ML} = \underset{\phi}{\operatorname{argmax}} p(y_T | y_1, \dots, y_{T-1}, \phi) p(y_{T-1}, y_{T-2}, \dots, y_1 | \phi, y_0)$$

$$= \prod_{t=1}^T p(y_t | y_{1:t-1}, \phi)$$

Under $AR(1)$, only the first lag is important for the distribution of y_t .

$$p(y_t | y_{1:t-1}, \phi) = p(y_t | \phi, y_{t-1})$$

We further assume $u_t \sim \mathcal{N}(0, \sigma^2)$, then we get: $y_t | y_{t-1} \sim \mathcal{N}(\phi y_{t-1}, \sigma^2)$, thus

$$\begin{aligned} p(y_{1:T} | \phi, y_0) &= \prod_{t=1}^T (2\pi\sigma^2)^{-1/2} \exp \left\{ -\frac{1}{2\sigma^2} (y_t - \phi y_{t-1})^2 \right\} \\ &= (2\pi\sigma^2)^{-T/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{t=1}^T (y_t - \phi y_{t-1})^2 \right\} \\ &= (2\pi\sigma^2)^{-T/2} \exp \left\{ -\frac{1}{2\sigma^2} (Y - X\phi)'(Y - X\phi) \right\} \end{aligned}$$

Unit-Root

Consider the $AR(1)$ process under the presence of a unit-root:

$$y_t = \phi_0 + \phi_1 y_{t-1} + u_t, \quad \phi_1 = 1$$

In this case, the process is not WS, thus not Strictly stationary and Ergodic.

If $\phi_0 = 0$, then

$$\sqrt{T}(\hat{\phi}_1 - \phi_1) \xrightarrow{d} \mathcal{N}\left(0, \frac{\sigma^2}{\mathbb{V}[y_t]}\right)$$

If we further assume that $y_0 = 0$, the process y_t can be represented as

$$y_t = \sum_{\tau=1}^t u_\tau$$

The central limit theorem for i.i.d. random variables implies that:

$$\frac{y_T}{\sqrt{T}} = \frac{1}{\sqrt{T}} \sum_{\tau=1}^T u_\tau \xrightarrow{d} \mathcal{N}(0, \sigma^2)$$

which suggests that

$$\frac{1}{T} \sum_{t=1}^T y_t = \frac{1}{\sqrt{T}} \sum_{t=1}^T \left[\sqrt{\frac{t}{T}} \frac{1}{\sqrt{t}} \sum_{\tau=1}^t u_\tau \right]$$

will not converge to a constant in probability but instead to a random variable.

Thus, we have a special distribution derived by Dickey and Fuller:

$$T(\hat{\phi}_1 - \phi_1) \xrightarrow{d} \text{Dickey-Fuller distribution}$$

4.3.2 Estimating Regressions with Autocorrelated Errors

Suppose we want to estimate $y_t = x_t' \beta + u_t$, where x_t contains k variables measured over time (and potentially a constant). In such time series, the error u_t is likely autocorrelated, $\mathbb{E}[u_t u_\tau] \neq 0$ for $t \neq \tau$. This leads to different standard errors for the OLS estimator $\hat{\beta}$ under the i.i.d. (cross-sectional) case.

Assuming that $\mathbb{E}[x_t u_t] = 0$, and that both u_t and x_t are strictly stationary and ergodic, we can apply the LLN for such processes.

$$\begin{aligned} \hat{\beta} &= (X'X)^{-1} X'Y \\ &= (X'X)^{-1} X'(X\beta + U) \\ &= \beta + (X'X)^{-1} X'U \\ &= \beta + \left(\frac{1}{T} \sum_{t=1}^T x_t x_t' \right)^{-1} \frac{1}{T} \sum_{t=1}^T x_t u_t \\ &\rightarrow \beta + \mathbb{E}[x_t x_t']^{-1} \mathbb{E}[x_t u_t] \\ &= \beta \text{ for SS+E, TS} \end{aligned}$$

and that

$$\begin{aligned} \sqrt{T}(\hat{\beta} - \beta) &= \sqrt{T} \left(\frac{1}{T} \sum_{t=1}^T x_t x_t' \right)^{-1} \left(\frac{1}{T} \sum_{t=1}^T x_t u_t \right) \\ &\xrightarrow{d} \mathcal{N}(0, V) \end{aligned}$$

where

$$V = \mathbb{E}[x_t x_t']^{-1} \mathbb{V} \left[\frac{1}{\sqrt{T}} \sum_{t=1}^T x_t u_t \right] \mathbb{E}[x_t x_t']^{-1}.$$

The term in the middle can be written as:

$$\begin{aligned}
 \mathbb{V} \left[\frac{1}{\sqrt{T}} \sum_{t=1}^T x_t u_t \right] &= \frac{1}{T} \sum_{t=1}^T \mathbb{V} [x_t u_t] + \frac{1}{T} \sum_{t=1}^T \sum_{\tau \neq t}^T \text{Cov}[x_t u_t, x_\tau u_\tau] \\
 &= \frac{1}{T} \sum_{t=1}^T \mathbb{E}[u_t^2 x_t x_t'] + \frac{1}{T} \sum_t \sum_{\tau \neq t} \mathbb{E}[x_t x_\tau' u_t u_\tau] \\
 &= \mathbb{E}[x_t x_t' u_t^2] + \sum_{h=1}^{T-1} \mathbb{E} [u_t u_{t-h} (x_t x_{t-h}' + x_{t-h} x_t')] .
 \end{aligned}$$

1. $x_t u_t$ is not autocorrelated, then the second term is zero.

- (a) if $\mathbb{E}[u_t^2] = \sigma^2$ (homoskedasticity), then $\mathbb{E}[u_t^2 x_t x_t'] = \sigma^2 \mathbb{E}[x_t x_t']$, the same expression as obtained under a homoskedastic linear regression model for cross-sectional data.
- (b) if $\mathbb{E}[u_t^2] = f(x_t)$ (heteroskedasticity), the first term does not simplify, and we construct the heteroskedasticity-robust variance $\mathbb{E}[x_t x_t']^{-1} \mathbb{E}[x_t x_t' u_t^2] \mathbb{E}[x_t x_t']^{-1}$. We just estimate it by replacing the expectation operators with time averages and \hat{u}_t : $\frac{1}{T} \sum_t \hat{u}_t^2 x_t x_t'$.

2. $x_t u_t$ is autocorrelated, then we can write V as:

$$\begin{aligned}
 &\mathbb{E}[u_t^2 x_t x_t'] + \frac{1}{T} \sum_t \sum_{h=1}^{T-1} \mathbb{E}[u_t u_{t-h} x_t x_{t-h}'] \\
 &= \frac{1}{T} \sum_t \hat{u}_t^2 x_t x_t' + \frac{1}{T} \sum_t \sum_{h=1}^{T-1} 2 \hat{u}_t \hat{u}_{t-h} x_t x_{t-h}'
 \end{aligned}$$

Remark.

For the second case, we recall the original version of variance:

$$\mathbb{V} \left[\frac{1}{T} \sum_{t=1}^T x_t u_t \right] = \frac{1}{T} \sum_{t=1}^T \sum_{\tau=1}^T \text{Cov}[u_t, u_\tau] x_t x_\tau'$$

We assume that $|\text{Cov}[u_t, u_\tau]| \leq |t - \tau|^\kappa$ where $\kappa > 1$. Since $u_t u_\tau$ can be treated as a “preestimator” of $\text{Cov}[u_t, u_\tau]$, we can replace the covariance with $u_t u_\tau$ and weight it with $\lambda(t - \tau)$ to yield the Newey-West/HAC estimator:

$$\hat{\sigma}^2 = \frac{1}{T} \sum_{t=1}^T \sum_{\tau=1}^T \lambda_m(t - \tau) u_t u_\tau x_t x_\tau'.$$

Multivariate Time Series

Many time series arising in practice are best considered as components of some vector-valued (multivariate) time series $y_t = (y_{1t}, y_{2t}, \dots, y_{nt})'$, which is an $n \times 1$ vector process observed in sequence series y_t , whose specification includes two parts:

- (1) The serial dependence of each component series $\{y_{it}\}$;
- (2) The interdependence between different component series $\{y_{it}\}$ and $\{y_{jt}\}$.

The most common multivariate time series models used by economists are vector autoregressions (VARs).

5.1 Multivariate Processes

Stationarity and ergodicity of a multivariate process are defined analogously to the univariate case. If all the finite dimensional joint distributions of the random variables $\{y_{it}\}$ were multivariate normal, then the distributional properties of $\{y_{it}\}$ would be completely determined by the means,

$$\mu_{it} = \mathbb{E}[y_{it}]$$

and covariances,

$$\gamma_{ij}(t, t-h) = \text{Cov}[y_{i,t}, y_{j,t-h}] = \mathbb{E}[(y_{i,t} - \mu_{i,t})(y_{j,t-h} - \mu_{j,t-h})].$$

It is more convenient when dealing with n interrelated series to use vector notation. The second-order properties of the multivariate process are specified by the mean vectors:

$$\mu_t = \mathbb{E}[y_t]$$

and covariance matrices:

$$\Gamma_t(h) = \text{Cov}[y_t, y_{t-h}] = \mathbb{E}[(y_t - \mu_t)(y_{t-h} - \mu_{t-h})'].$$

Definition 5.1.1 (Stationary Multivariate Time Series).

A multivariate time series $\{y_t\}$ is said to be weakly stationary if:

- (1) $\mu_t = \mu$ for all t ;
- (2) $\Gamma_t(h) = \Gamma(h) = \mathbb{E}[(y_t - \mu)(y_{t-h} - \mu)']$ for all t and h .

We shall refer to μ as the mean of the series and $\Gamma(h)$ as the covariance function of the series at lag h .

Example 7.

Take the bivariate case ($n = 2$) as an example, we have:

$$\Gamma(h) = \begin{bmatrix} \text{Cov}[y_{1,t}, y_{1,t-h}] & \text{Cov}[y_{1,t}, y_{2,t-h}] \\ \text{Cov}[y_{1,t-h}, y_{2,t}] & \text{Cov}[y_{2,t}, y_{2,t-h}] \end{bmatrix} = \begin{bmatrix} \gamma_{11}(h) & \gamma_{12}(h) \\ \gamma_{21}(h) & \gamma_{22}(h) \end{bmatrix}$$

The diagonal elements are simply the univariate autocovariances of y_{1t} and y_{2t} , while the off-diagonal elements are the cross-covariances between y_{1t} and y_{2t} , with the former (latter) leading the latter (former) by a displacement of h lags. We should notice that $\gamma_{12}(h) \neq \gamma_{21}(h)$, as the correlation of y_{1t} with the past movement in y_{2t} is not necessarily the same as the correlation of y_{2t} with the past movement in y_{1t} .

As a result, in contrast to the univariate case, in multivariate case, $\Gamma(h) \neq \Gamma(-h)$. Instead, we have $\gamma_{12}(h) = \gamma_{21}(-h)$, because:

$$\Gamma_t(-h)' = \mathbb{E}[(y_t - \mu_t)(y_{t+h} - \mu_{t+h})']' = \mathbb{E}[(y_{t+h} - \mu_{t+h})(y_t - \mu_t)'] = \Gamma_{t+h}(h)$$

and under weakly stationarity, time doesn't matter: $\Gamma_{t+h}(h) = \Gamma_t(h) = \Gamma(h)$, and therefore, $\Gamma(h) = \Gamma(-h)'$

A multivariate White Noise (WN) process and the multivariate General Linear Process (GLP) are also defined analogously to the univariate case.

Definition 5.1.2 (Multivariate White Noise).

The n -variate series u_t is said to be a WN process with mean zero and covariance Σ , written $u_t \sim WN(0, \Sigma)$, if and only if u_t is weakly stationary with mean vector $\mathbf{0}$ and covariance matrix function:

$$\Gamma(h) = \text{Cov}[u_t, u_{t-h}] = \mathbb{E}[u_t u_{t-h}'] \begin{cases} \Sigma, & h = 0 \\ 0, & h \neq 0 \end{cases}$$

Note that Σ is not necessarily diagonal, i.e. the individual components of the vector valued process u_t may be contemporaneously correlated. However, the defining feature of a WN process is that each of the components is uncorrelated with its own past (and future) movements as well as those of any of the other components.

The multivariate GLP is defined as:

$$y_t = B(L)u_t = \sum_{l=0}^{\infty} B_l u_{t-l}, \quad u_t \sim WN(0, \Sigma)$$

where $B_0 = I$ and $\sum_{l=0}^{\infty} \|B_l\|^2 < \infty$.

5.2 Vector Autoregressions (VARs)

A VAR is a common way to approximate the multivariate GLP. As in the previous lecture, we have built a VAR(1) model for solving the AR(2) process.

$$y_t = \Phi_0 + \Phi_1 y_{t-1} + u_t, \quad u_t \sim WN(0, \Sigma)$$

$$\begin{bmatrix} y_{1t} \\ y_{2t} \end{bmatrix} = \begin{bmatrix} \Phi_{0,1} \\ \Phi_{0,2} \end{bmatrix} + \begin{bmatrix} \Phi_{1,11} & \Phi_{1,12} \\ \Phi_{1,21} & \Phi_{1,22} \end{bmatrix} \begin{bmatrix} y_{1,t-1} \\ y_{2,t-1} \end{bmatrix} + \begin{bmatrix} u_{1t} \\ u_{2t} \end{bmatrix}, \quad u_t \sim WN(0, \Sigma)$$

where $u_t \sim WN\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} \sigma_{11} & \sigma_{12} \\ \sigma_{21} & \sigma_{22} \end{bmatrix}\right)$. There are two sources of interaction between y_{1t} and y_{2t} :

- (1) The value of each series today depends on its own, but also on the other variable's value at time $t - 1$;
- (2) The innovations of the two series are potentially correlated (contemporaneously).

If the equations include p lags of each variable we obtain the p -th order vector autoregressive (VAR) model:

$$y_t = \Phi_0 + \Phi_1 y_{t-1} + \cdots + \Phi_p y_{t-p} + u_t \quad (5.1)$$

Our notation allows $y_t = [y_{1t}, y_{2t}, \dots, y_{nt}]'$ to be a vector of dimension n in which case the matrices Φ_i are $n \times n$ and error u_t is $n \times 1$. We will denote the elements of Φ_i using the notation:

$$\Phi_i = \begin{bmatrix} \Phi_{i,11} & \Phi_{i,12} & \cdots & \Phi_{i,1n} \\ \Phi_{i,21} & \Phi_{i,22} & \cdots & \Phi_{i,2n} \\ \vdots & \vdots & \ddots & \vdots \\ \Phi_{i,n1} & \Phi_{i,n2} & \cdots & \Phi_{i,nn} \end{bmatrix}.$$

The error $u_t = [u_{1t}, u_{2t}, \dots, u_{nt}]'$ is the component of y_t which is unforecastable at time $t - 1$. However, the components of y_t are contemporaneously correlated, therefore the contemporaneous covariance matrix $\Sigma = \mathbb{E}[uu']$ is non-diagonal.

Definition 5.2.1 (VAR(p)).

A VAR(p) model is a system of n equations of the form:

$$y_t = \Phi_0 + \sum_{i=1}^p \Phi_i y_{t-i} + u_t, \quad u_t \sim WN(0, \Sigma) \quad (5.2)$$

where y_t is an n -dimensional vector, Φ_0 is an n -dimensional vector of constants, Φ_i are $n \times n$ matrices of coefficients, and u_t is an n -dimensional vector of white noise errors.

5.2.1 Stationarity of VAR(p)

We are interested in conditions under which y_t is a stationary process. Let $\lambda_i(\Phi)$ denote the i -th eigenvalue of the matrix Φ .

Theorem 5.2.1 (SS and Ergodic of VAR(1)).

If u_t is strictly stationary and ergodic, with $\mathbb{E}[u_t] < \infty$, and $|\lambda_i(\Phi_1)| < 1$, then the VAR(1) process y_t is also strictly stationary and ergodic.

Example 8.

Consider the VAR(2) model:

$$y_t = \Phi_0 + \Phi_1 y_{t-1} + \Phi_2 y_{t-2} + u_t$$

or we write in companion form:

$$\begin{bmatrix} y_t \\ y_{t-1} \end{bmatrix} = \underbrace{\begin{bmatrix} \Phi_0 \\ 0 \end{bmatrix}}_{F_0} + \underbrace{\begin{bmatrix} \Phi_1 & \Phi_2 \\ I & 0 \end{bmatrix}}_{F_1 (np \times np)} \begin{bmatrix} y_{t-1} \\ y_{t-2} \end{bmatrix} + \begin{bmatrix} u_t \\ 0 \end{bmatrix}$$

we can simply write:

$$Y_t = F_0 + F_1 Y_{t-1} + v_t$$

We take mean on both sides,

$$\mathbb{E}[y_t] = \Phi_0 + \Phi_1 \mathbb{E}[y_{t-1}] + \Phi_2 \mathbb{E}[y_{t-2}] + 0$$

as $\mathbb{E}[y_t] = \mathbb{E}[y_{t-1}] = \mathbb{E}[y_{t-2}] = \mu$, we know that $\mu = (1 - \Phi_1 - \Phi_2)^{-1} \Phi_0$.

For the $VAR(p)$ model 5.2, we can write using the lag operator notation:

$$\Phi(L)y_t = \Phi_0 + u_t$$

where

$$\Phi(z) = I_n - \Phi_1 z - \Phi_2 z^2 - \dots - \Phi_p z^p.$$

The condition for stationarity of the system can be expressed as a restriction on the roots of the determinantal equation of the autoregressive polynomial. Recall that a root r of $\det(\Phi(z))$ is a solution to $\det(\Phi(r)) = 0$.

Theorem 5.2.2 (SS and Ergodicity of $VAR(p)$).

If all roots r of the polynomial $\det(\Phi(z))$ satisfy $|r| > 1$, then the $VAR(p)$ process y_t is strictly stationary and ergodic.

Or, equivalently, we write the $VAR(p)$ model in companion form:

$$\begin{bmatrix} y_t \\ y_{t-1} \\ \vdots \\ y_{t-p+1} \end{bmatrix} = \begin{bmatrix} \Phi_0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} + \underbrace{\begin{bmatrix} \Phi_1 & \Phi_2 & \dots & \Phi_p \\ I & 0 & \dots & 0 \\ 0 & I & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & \dots & I & 0 \end{bmatrix}}_{F_{np \times np}} \begin{bmatrix} y_{t-1} \\ y_{t-2} \\ \vdots \\ y_{t-p} \end{bmatrix} + \begin{bmatrix} u_t \\ 0 \\ \vdots \\ 0 \end{bmatrix} \quad (5.3)$$

We can simply write $Y_t = F_0 + F Y_{t-1} + v_t$, hence Theorem 5.2.1 can be reorganize to:

Theorem 5.2.3 (SS and Ergodicity of $VAR(p)$).

If all eigenvalues $\lambda_i(F)$ of the matrix F , given by:

$$|I_n \lambda^p - \Phi_1 \lambda^{p-1} - \Phi_2 \lambda^{p-2} - \dots - \Phi_p| = 0$$

satisfy $|\lambda_i(F)| < 1$, then the $VAR(p)$ process y_t is strictly stationary and ergodic.

5.2.2 Moments & GLP-Representation

Once we know that a $VAR(p)$ is Weakly Stationary, it is quite trivial to calculate its mean.

Consider first for a VAR(1). Repeatedly inserting for lags y_{t-1} , we get:

$$\begin{aligned}
 y_t &= \Phi_0 + \Phi_1 y_{t-1} + u_t \\
 &= \Phi_0 + \Phi_1(\Phi_0 + \Phi_1 y_{t-2} + u_{t-1}) + u_t \\
 &= \Phi_0 + \Phi_1 \Phi_0 + \Phi_1^2 y_{t-2} + \Phi_1 u_{t-1} + u_t \\
 &= \dots \\
 &= \sum_{l=0}^{\infty} \Phi_1^l \Phi_0 + \sum_{l=0}^{\infty} \Phi_1^l u_{t-l} + \lim_{k \rightarrow \infty} \Phi_1^k y_{t-k} \\
 &= \frac{\Phi_0}{1 - \Phi_1} + \sum_{l=0}^{\infty} \Phi_1^l u_{t-l}
 \end{aligned}$$

if Φ_1 has all eigenvalues in the unit circle, then $\lim_{k \rightarrow \infty} \Phi_1^k y_{t-k} = 0$. Based on this expresion, we can get the ACF:

$$\begin{aligned}
 \Gamma(h) &= \text{Cov}[y_t, y_{t-h}] = \mathbb{E} \left[\left(\sum_{l=0}^{\infty} \Phi_1^l u_{t-l} \right) \left(\sum_{k=0}^{\infty} \Phi_1^k u_{t-h-k} \right)' \right] \\
 &= \sum_{k=0}^{\infty} \mathbb{E} \left[(\Phi_1^{k+h} u_{t-h-k}) (\Phi_1^k u_{t-h-k})' \right] \\
 &= \sum_{k=0}^{\infty} \Phi_1^{h+k} \mathbb{E} [u_{t-h-k} u_{t-h-k}'] \Phi_1^{k'} \\
 &= \sum_{k=0}^{\infty} \Phi_1^{h+k} \Sigma \Phi_1^{k'}
 \end{aligned}$$

The previous discussion can be extended easily to VAR(p) processes with $p > 1$ because any VAR(p) process can be written in VAR(1) form. More precisely, if y_t is a VAR(p) process as in definition 5.2.1, a corresponding np -dimensional VAR(1) is shown as equation 5.3: $Y_t = F_0 + FY_{t-1} + v_t$. Recall that

$$Y_t = \begin{bmatrix} y_t \\ y_{t-1} \\ \vdots \\ y_{t-p+1} \end{bmatrix}, \quad F = \begin{bmatrix} \Phi_1 & \Phi_2 & \dots & \Phi_p \\ I_n & 0 & \dots & 0 \\ 0 & I_n & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & \dots & I_n & 0 \end{bmatrix}, \quad v_t = \begin{bmatrix} u_t \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

Following previous discussion, we know that Y_t is strictly stationary and ergodic, if $|\lambda_i(F)| < 1$ (Theorem 5.2.3). It's mean vector is:

$$\mu = \mathbb{E}[Y_t] = (I_{np} - F)^{-1} F_0$$

and the autocovariances are

$$\Gamma_Y(h) = \sum_{k=0}^{\infty} F^{h+k} \Sigma^v F^{k'}$$

where $\Sigma^v = \mathbb{E}[v_t v_t'] = \mathbb{V}[u_t] = \begin{bmatrix} \Sigma & 0 \\ 0 & 0 \end{bmatrix}$. Using the $(n \times np)$ matrix

$$M = \begin{bmatrix} I_n & 0 & 0 & 0 \end{bmatrix},$$

the process y_t is obtained as $y_t = MY_t$. Because Y_t is strictly stationary and ergodic, y_t is also strictly

stationary and ergodic, meaning that $\mathbb{E}[y_t] = M\mu$ is constant for all t and the autocovariances $\Gamma_y(h) = M\Gamma_Y(h)M'$ are also time invariant. And we can have the ACF for the original y_t :

$$\begin{aligned}\Gamma_y(h) &= \text{Cov}[y_t, y_{t-h}] = \text{Cov}[MY_t, MY_{t-h}] \\ &= M \text{Cov}[Y_t, Y_{t-h}] M' \\ &= \sum_{k=0}^{\infty} M F^{h+k} \Sigma^v F^{k'} M' \\ &= \sum_{k=0}^{\infty} (F^{k+h})_{11} \Sigma (F^k)'_{11}\end{aligned}$$

where $(F^k)_{11}$ denotes the upper-left $(n \times n)$ block of the $(np \times np)$ matrix F^k .

Computation of Autocovariances and Autocorrelations of VAR Processes

In order to illustrate the computation of the autocovariances when the process coefficients are given, suppose that y_t is strictly stationary and ergodic VAR(1) process:

$$y_t = \Phi_0 + \Phi_1 y_{t-1} + u_t, \quad u_t \sim WN(0, \Sigma)$$

We can rewrite the equation in mean-adjusted form:

$$y_t - \mu = \Phi_1 (y_{t-1} - \mu) + u_t$$

where $\mu = \mathbb{E}[y_t] = \frac{\Phi_0}{1-\Phi_1}$. Postmultiplying by $(y_{t-h} - \mu)'$ and taking the expectation, we have:

$$\text{Cov}[y_t, y_{t-h}] = \mathbb{E}[(y_t - \mu)(y_{t-h} - \mu)'] = \Phi_1 \mathbb{E}[(y_{t-1} - \mu)(y_{t-h} - \mu)'] + \mathbb{E}[u_t(y_{t-h} - \mu)'].$$

Thus, for $h = 0$, we have:

$$\Gamma_y(0) = \text{Cov}[y_t, y_t] = \Phi_1 \text{Cov}[y_{t-1}, y_t] + \mathbb{E}[u_t u_t'] = \Phi_1 \Gamma_y(-1) + \Sigma = \Phi_1 \Gamma_y(1)' + \Sigma \quad (5.4)$$

and for $h > 0$,

$$\Gamma_y(h) = \Phi_1 \Gamma_y(h-1) + \mathbb{E}\left[u_t \left(\sum_{j=0}^{\infty} \Phi_1^j u_{t-1-j}\right)'\right] = \Phi_1 \Gamma_y(h-1) \quad (5.5)$$

as u_t is strict white noise, indicating that $\mathbb{E}[u_t u_s] = 0$ for $t \neq s$.

These equations are usually referred to as the *Yule-Walker equations* (??). If Φ_1 and the covariance matrix $\Gamma_y(0) = \Sigma_y$ of y_t are unknown, the $\Gamma_y(h)$ can be computed recursively using the equation above.

If Φ_1 and Σ are given, $\Gamma_y(0)$ can be determined as follows. For $h = 1$, we get from equation 5.4 that $\Gamma_y(1) = \Phi_1 \Gamma_y(0)$, and substitute back into equation 5.4 gives:

$$\Gamma_y(0) = \Phi_1 \Gamma_y(0) \Phi_1' + \Sigma^1$$

This is a matrix equation in $\Gamma_y(0)$, which can be solved using the Kronecker product and vector operator

¹Or using the stationarity of y_t , we know that $\mathbb{V}[y_t] = \mathbb{V}[y_{t-1}]$, so

$$\Gamma_y(0) = \mathbb{V}[y_t] = \mathbb{V}[\Phi_1 y_{t-1} + u_t] = \Phi_1 \Gamma_y(0) \Phi_1' + \Sigma$$

(See appendix ??).

$$\begin{aligned} \text{vec } \Gamma_y(0) &= \text{vec}(\Phi_1 \Gamma_y(0) \Phi_1') + \text{vec } \Sigma \\ &= (I_n \otimes \Phi_1) \text{vec } \Gamma_y(0) + \text{vec } \Sigma \end{aligned}$$

Hence,

$$\text{vec } \Gamma_y(0) = (I_{n^2} - \Phi_1 \otimes \Phi_1)^{-1} \text{vec } \Sigma$$

Note that the invertibility of $I_{n^2} - \Phi_1 \otimes \Phi_1$ follows from the strict stationarity of y_t because the eigenvalues of $\Phi_1 \otimes \Phi_1$ are the products of the eigenvalues of Φ_1 . Hence, the eigenvalues of $\Phi_1 \otimes \Phi_1$ are all less than 1 in absolute value.

5.2.3 Reduced-Form vs. Structural Representation

The above expressions for the VAR model are called reduced-form representation and the u_t that appear in it are called reduced-form errors. They are the forecasting errors obtained when predicting y_t one step ahead:

$$\begin{aligned} y_t &= \Phi_0 + \Phi_1 y_{t-1} + u_t \\ \Rightarrow \mathbb{E}_{t-1}[y_t] &= \Phi_0 + \Phi_1 y_{t-1} + \mathbb{E}_{t-1}[u_t] \\ \Rightarrow u_t &= y_t - \mathbb{E}_{t-1}[y_t] \end{aligned}$$

This is why u_t is also referred to as the innovation to the process y_t at time t ; it contains everything that affects y_t and is not known to the researcher at time $t - 1$.

Generally, in impulse response analysis the emphasis has shifted from specifying the relations between the observable variables directly to interpreting the unexpected part of their changes or the shocks. To make causal statements in VAR, we need to decompose the reduced-form errors into underlying, independent driving forces of y_t , referred to as shocks (structural innovations). One way to do so is to think of the forecast errors as linear functions of the structural innovations. In that case, we have the relations

$$u_t = \Phi_\varepsilon \varepsilon_t, \quad \varepsilon_t \sim WN(0, I_n)$$

where ε_t is the vector of shocks. Hence $\Sigma = \Phi_\varepsilon \Sigma_\varepsilon \Phi_\varepsilon'$, as we normalize the variances of the structural innovations to one, we have $\Sigma = \Phi_\varepsilon \Phi_\varepsilon'$.

Note.

Due to the symmetry of the covariance matrix, these relations specify only $\frac{n(n+1)}{2}$ different equations and we need again $\frac{n(n-1)}{2}$ further relations to identify all n^2 elements of Φ_ε . We shall discuss this in a following section.

Let's review the two forms of VAR representation:

$$\begin{aligned} \text{Reduced-form: } y_t &= \Phi_0 + \Phi_1 y_{t-1} + u_t, \quad u_t \sim WN(0, \Sigma) \\ \text{Structural: } y_t &= \Phi_0 + \Phi_1 y_{t-1} + \Phi_\varepsilon \varepsilon_t, \quad \varepsilon_t \sim WN(0, I_n) \end{aligned}$$

For a general VAR(p) model

$$y_t = \Phi_0 + \sum_{i=1}^{\infty} \Phi_i y_{t-i} + u_t, \quad u_t \sim WN(0, \Sigma)$$

Sometimes, we write the structural representation as:

$$Ay_t = B_0 + B_1y_{t-1} + \cdots + B_py_{t-p} + \varepsilon_t, \quad \varepsilon_t \sim WN(0, I_n)$$

where $A = \Phi_\varepsilon^{-1}$, $B_0 = \Phi_\varepsilon^{-1}\Phi_0$, $B_i = \Phi_\varepsilon^{-1}\Phi_i$ and $\varepsilon_t = \Phi_\varepsilon^{-1}u_t \sim (0, \Sigma_\varepsilon = A\Sigma A')$. For the structural representation of VAR(1), we can write:

$$y_t = (I - \Phi_1)^{-1}\Phi_0 + \sum_{l=0}^{\infty} \Phi_1^l \Phi_\varepsilon \varepsilon_{t-l}$$

$$\Gamma_y(0) = \sum_{l=0}^{\infty} \Phi_1^l \Sigma_v \Phi_1^{l'}, \quad \Sigma_v = \mathbb{E}[\Phi_\varepsilon \varepsilon_t \varepsilon_t' \Phi_\varepsilon'] = \Phi_\varepsilon \Phi_\varepsilon'$$

Suppose that we know all parameters in the VAR, including the matrix Φ_ε that maps reduced-form errors into structural shocks. This allows us to compute impulse-response functions (IRFs), variance decompositions (VDs) and historical decompositions (HDs).

An IRF illustrates the dynamic effects of a shock on a series in y_t . Having identified the structural shocks ε_t , our interest usually is not in the shocks themselves, however, but in the responses of each element of y_t to a one-time impulse of ε_t .

Definition 5.2.2 (Impulse Response Function (IRF)).

The effect of shock j , ε_{jt} on series i , y_{it} , h periods into the future is given by:

$$\Theta_{ij} = \frac{\partial y_{i,t+h}}{\partial \varepsilon_{j,t}} = \left[\frac{\partial y_{t+h}}{\partial \varepsilon_t} \right]_{ij} = [(F^h)_{11} \Phi_\varepsilon]_{ij} = [\Phi_1^h \Phi_\varepsilon]_{ij} \quad (5.6)$$

The sequence Θ_{ij} is referred to as the IRF of variable i to shock j , which tells us the dynamic effect of ε_{jt} on the series y_{it} . As by definition, ε_{jt} is i.i.d., this effect has a causal interpretation.

A variance decomposition of y_t determines the contribution of each shock to the variance of y_t .

$$\Gamma_0 = \mathbb{V}[y_t] = \sum_{l=0}^{\infty} \Phi_1^l \Phi_\varepsilon I \Phi_\varepsilon' \Phi_1^{l'}$$

Let $I^{(j)}$ be the identity matrix with all but the j -th diagonal element equal to zero. Then

$$\Gamma_0^{(j)} = \sum_{l=0}^{\infty} \Phi_1^l \Phi_\varepsilon I^{(j)} \Phi_\varepsilon' \Phi_1^{l'}$$

We have that $\Gamma_0 = \sum_{j=1}^n \Gamma_0^{(j)}$, as a result, the fraction of the variance of a particular series y_{it} , $\mathbb{V}[y_{it}] = [\Gamma_0]_{ij}$ is explained by ε_{jt} given by:

$$[\Gamma_0^{(j)}]_{ij} / [\Gamma_0]_{ij}.$$

5.3 Estimation of Reduced-Form VARs

Consider the general VAR(p) model:

$$y_t = \Phi_0 + \Phi_1 y_{t-1} + \cdots + \Phi_p y_{t-p} + u_t, \quad u_t \sim WN(0, \Sigma) \quad (5.7)$$

and we write it into a simple regression form:

$$y'_t = x'_t \Phi + u'_t, \quad x_t = \begin{bmatrix} y_{t-1} \\ y_{t-2} \\ \vdots \\ 1 \end{bmatrix}, \quad \Phi = \begin{bmatrix} \Phi'_1 \\ \Phi'_2 \\ \vdots \\ \Phi'_0 \end{bmatrix}$$

in matrix form:

$$Y = X\Phi + U, \quad Y = \begin{bmatrix} y'_1 \\ y'_2 \\ \vdots \\ y'_T \end{bmatrix}.$$

To simplify, we only analyze for the OLS estimator:

$$\begin{aligned} \min_{\Phi} \sum_t \sum_i u_{it}^2 &= \min_{\Phi} \sum_t u'_t u_t \\ &= \min_{\Phi} \text{tr}[(Y - X\Phi)'(Y - X\Phi)] \\ \Rightarrow \hat{\Phi} &= (X'X)^{-1} X'Y \\ \Sigma = \mathbb{V}[u_t] &= \begin{bmatrix} \mathbb{V}[u_{1t}] & \text{Cov}[u_{1t}, u_{2t}] \\ \cdot & \mathbb{V}[u_{2t}] \end{bmatrix} \leftarrow \begin{bmatrix} \frac{1}{T} \sum_t \hat{u}_{1t}^2 & \frac{1}{T} \sum_t \hat{u}_{1t} \hat{u}_{2t} \\ \cdot & \frac{1}{T} \sum_t \hat{u}_{2t}^2 \end{bmatrix} = \hat{\Sigma} \\ \hat{\Sigma} &= \frac{1}{T} \sum_{t=1}^T \hat{u}_t \hat{u}'_t = \frac{1}{T} \hat{U}' \hat{U} = \frac{1}{T} (Y - X\hat{\Phi})' (Y - X\hat{\Phi}) \end{aligned}$$

This holds whether or not x_t is actually a VAR(p) process. By the properties of projection errors, $\mathbb{E}[x_t u'_t] = 0$. The projection coefficient matrix $\hat{\Phi}$ is identified if $X'X = \mathbb{E}[x_t x'_t]$ is invertible.

Theorem 5.3.1 (Identification of SVAR).

If y_t is strictly stationary, and $\Sigma = \mathbb{E}[u_t u'_t] < \infty$, then $Q = \mathbb{E}[x_t x'_t] > 0$ and the coefficient vector $\hat{\Phi}$ is identified.

5.4 Estimation of Structural VARs

We consider the structural VAR(1) model in the form:

$$y_t = \Phi_0 + \Phi_1 y_{t-1} + \Phi_\varepsilon \varepsilon_t, \quad \varepsilon_t \sim WN(0, I)$$

As we have previously discussed in a note, there are $\frac{n(n+1)}{2}$ unique elements in Σ , but n^2 elements in Φ_ε . So, *point identify* Φ_ε , we need $\frac{n(n-1)}{2}$ additional restrictions. Even with fewer restrictions, we can decrease the identified set, and we can possibly point-identify the effects of some of the n shocks in ε_t (by point-identifying the corresponding columns of Φ_ε) or the responses of some of the n series in y_t by point-identifying the corresponding rows of Φ_ε .

We can further analyze the structural VAR model by decomposing Φ_ε into a part that we can point-identify and a part that we cannot. Let Ω be an orthogonal matrix such that $\Omega\Omega' = I$. Then we can decompose Φ_ε as:

$$\Phi_\varepsilon = \Sigma_{\text{tr}} \Omega$$

where Σ_{tr} is the Cholesky decomposition of Σ , satisfying $\Sigma_{\text{tr}} \Sigma'_{\text{tr}} = \Sigma$, and Ω is any orthogonal matrix

with $\Omega\Omega' = I$.

Because we can point-identify Σ , and because Cholesky decomposition is unique up to the ordering of the diagonal elements, we can point-identify the diagonal elements of Σ_{tr} from the data. And that any Ω is consistent with the data (gives the same likelihood), as for any Ω , $\Phi_\varepsilon\Phi'_\varepsilon = \Sigma_{\text{tr}}\Omega\Omega'\Sigma_{\text{tr}} = \Sigma$.

From a frequentist point of view, this means that we cannot find a (unique) point estimator for Ω from the data.

Example 9.

For $n = 2$, we can obtain a nice representation of the space of orthogonal matrices:

$$\Omega(\theta, \xi) = \begin{bmatrix} \cos \theta & -\xi \sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}, \quad \text{where } -\pi < \theta < \pi, \quad \xi = \{-1, 1\}$$

We notice that $\Omega(\frac{\pi}{2}) = \Omega(-\frac{\pi}{2})$, which means that only the signs of the impulse responses change but not the shape.

5.5 Point Restrictions

A point-restriction fixes some of the elements of Ω . For example, one might impose a point restriction on short-term dynamics, such as real activity reacting to monetary policy shocks only with a lag. If y_{it} is GDP and ε_{jt} is the MP shock, this implies that:

$$\frac{\partial y_{it}}{\partial \varepsilon_{jt}} = ((F^0)_{11}\Phi_\varepsilon)_{ij} = (\Sigma_{\text{tr}}\Omega) = 0$$

and restrict the dot product of the i -th row of Σ_{tr} and the j -th column of Ω to be zero.

Example 10 (Short-run Restrictions).

Consider the following example: the central bank does not react contemporaneously to technology shocks because data on aggregate output only become available with a one-quarter lag. Suppose that

$$y_t = \begin{bmatrix} \text{Fed Funds} \\ \text{Output} \end{bmatrix}, \quad \varepsilon_t = \begin{bmatrix} \varepsilon_{R,t} \\ \varepsilon_{z,t} \end{bmatrix} = \begin{bmatrix} \text{Technology shock} \\ \text{Monetary Policy shock} \end{bmatrix}$$

Formalization in our framework:

- (i) $\theta = 0, \xi = 1$;
- (ii) $\theta = 0, \xi = -1$;
- (iii) $\theta = \pi, \xi = 1$;
- (iv) $\theta = \pi, \xi = -1$;

Under (i), we have:

$$u_t = \begin{bmatrix} \Sigma_{11}^{\text{tr}} & 0 \\ \Sigma_{21}^{\text{tr}} & \Sigma_{22}^{\text{tr}} \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} \varepsilon_{R,t} \\ \varepsilon_{z,t} \end{bmatrix}$$

Choosing between (i)-(iv) normalizes the direction of responses.

Example 11 (Long-run Restrictions).

Now, let's assume that monetary policy shocks do not raise output in the long-run. Suppose that

$$y_t = \begin{bmatrix} \text{Inflation} \\ \text{Output Growth} \end{bmatrix}, \quad \varepsilon_t = \begin{bmatrix} \varepsilon_{R,t} \\ \varepsilon_{z,t} \end{bmatrix} = \begin{bmatrix} \text{Technology shock} \\ \text{Monetary Policy shock} \end{bmatrix}$$

Moreover, we have (as in the definition 4.2.4):

$$y_t = \left(\sum_{j=0}^{\infty} \phi_j L^j \right) u_t = \phi(L) u_t$$

Notice that ξ only affects the response to the second shock. Hence we set it to 1.

Let's examine the moving average representation of y_t in terms of the structural shocks:

$$\begin{aligned} y_t &= \begin{bmatrix} \phi_{11}(L) & \phi_{12}(L) \\ \phi_{21}(L) & \phi_{22}(L) \end{bmatrix} \begin{bmatrix} \Sigma_{11}^{\text{tr}} & 0 \\ \Sigma_{21}^{\text{tr}} & \Sigma_{22}^{\text{tr}} \end{bmatrix} \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} \varepsilon_{R,t} \\ \varepsilon_{z,t} \end{bmatrix} \\ &= \begin{bmatrix} \Sigma_{11}^{\text{tr}} \cos \theta \phi_{21}(L) + (\Sigma_{21}^{\text{tr}} \cos \theta + \Sigma_{22}^{\text{tr}} \sin \theta) \phi_{22}(L) & \cdot \\ d_{11}(L) & d_{12}(L) \end{bmatrix} \begin{bmatrix} \varepsilon_{R,t} \\ \varepsilon_{z,t} \end{bmatrix} \\ &= \begin{bmatrix} d_{11}(L) & d_{12}(L) \\ d_{21}(L) & d_{22}(L) \end{bmatrix} \begin{bmatrix} \varepsilon_{R,t} \\ \varepsilon_{z,t} \end{bmatrix} \end{aligned}$$

Suppose that in period $t = 0$ log output and log prices are equal to zero. Then the log-level of output and prices in period $t = T > 0$ is given by:

$$y_T = \sum_{t=1}^T y_t = \sum_{t=1}^T \sum_{j=0}^{\infty} D_j \varepsilon_{t-j}$$

Now consider the derivative

$$\frac{\partial y_T}{\partial \varepsilon'_1} = \sum_{j=0}^{T-1} D_j$$

Letting $T \rightarrow \infty$ gives us the long-run response of the level of prices and output to the shock ε_1 :

$$\frac{\partial y_{\infty}}{\partial \varepsilon'_1} = \sum_{j=0}^{\infty} D_j = D(1)$$

Here, we want to restrict the long-run effect of monetary policy shocks on output: $d_{21}(1) = 0$. This leads us to the equation

$$[\Sigma_{11}^{\text{tr}} \phi_{21}(1) + \Sigma_{21}^{\text{tr}} \phi_{22}(1)] \cos \theta + \Sigma_{22}^{\text{tr}} \phi_{22}(1) \sin \theta = 0.$$

Notice that the equation has two solutions for $\theta \in (-\pi, \pi]$. Under one solution a positive monetary policy shock is contractionary, under the other solution it is expansionary. The shape of the responses is, of course, the same.

Lecture 6.

Appendix

6.1 Yule-Walker

6.2 Kronecker Products & Vector Operator

Recommended Resources

Books

- [1] Peng Ding. *A First Course in Causal Inference*. 2023. arXiv: [2305.18793](https://arxiv.org/abs/2305.18793) [stat.ME]. URL: <https://arxiv.org/abs/2305.18793> (p. 8)
- [2] Bruce E. Hansen. *Econometrics*. Princeton, New Jersey: Princeton University Press, 2022 (pp. 11, 19, 23, 25)
- [3] Jeffrey M. Wooldridge. *Econometric Analysis of Cross Section and Panel Data*. 2nd ed. Cambridge, Massachusetts: The MIT Press, 2010 (p. 14)
- [4] Peter Whittle. *Hypothesis Testing in Time Series Analysis*. Uppsala: Almqvist & Wiksells, 1951 (p. 36)
- [5] Helmut Lütkepohl. *New Introduction to Multiple Time Series Analysis*. New York: Springer, 2005
- [6] James H. Stock and Mark W. Watson. *Introduction to Econometrics*. 4th ed. New York: Pearson, 2003
- [7] Jeffrey M. Wooldridge. *Introductory Econometrics: A Modern Approach*. 7th ed. Cengage Learning, 2020
- [8] Fumio Hayashi. *Econometrics*. Princeton, New Jersey: Princeton University Press, 2000
- [9] Joshua Chan et al. *Bayesian Econometric Methods*. 2nd ed. Cambridge, United Kingdom: Cambridge University Press, 2019
- [10] Badi H. Baltagi. *Econometric Analysis of Panel Data*. 6th ed. Cham, Switzerland: Springer, 2021
- [11] James D. Hamilton. *Time Series Analysis*. Princeton, New Jersey: Princeton University Press, 1994. ISBN: 9780691042893
- [12] Takeshi Amemiya. *Advanced Econometrics*. Cambridge, MA: Harvard University Press, 1985
- [13] David Walters. *An Introduction to Stochastic Processes and Their Applications*. New York: Dover Publications, 1982
- [14] Colorado P. J. Brockwell Fort Collins. *Time Series: Theory and Methods*. New York: Springer, 1991

Others

- [15] Donald B. Rubin. “Bayesian Inference for Causality: The Importance of Randomization”. In: *The Annals of Statistics* 3.1 (1975), pp. 121–131. DOI: [10.1214/aos/1176343238](https://doi.org/10.1214/aos/1176343238) (p. 8)
- [16] Paul W. Holland. “Statistics and Causal Inference(with discussion)”. In: *Journal of the American Statistical Association* 81.396 (1986), pp. 945–960. DOI: [10.1080/01621459.1986.10478373](https://doi.org/10.1080/01621459.1986.10478373) (p. 8)
- [17] Jerzy Neyman. “On the Application of Probability Theory to Agricultural Experiments. Essay on Principles. Section 9”. In: *Statistical Science* 5.4 (1923), pp. 465–472 (p. 8)
- [18] Donald B. Rubin. “Estimating Causal Effects of Treatments in Randomized and Nonrandomized Studies”. In: *Journal of Educational Psychology* 66.5 (1974), pp. 688–701. DOI: [10.1037/h0037350](https://doi.org/10.1037/h0037350) (p. 8)

- [19] Donald B. Rubin. “Comment on "Randomization Analysis of Experimental Data: The Fisher Randomization Test" by D. Basu”. In: *Journal of the American Statistical Association* 75.371 (1980), pp. 591–593. DOI: [10.1080/01621459.1980.10477410](https://doi.org/10.1080/01621459.1980.10477410) (pp. 8, 9)
- [20] Donald B. Rubin. “Causal Inference Using Potential Outcomes: Design, Modeling, Decisions”. In: *Journal of the American Statistical Association* 100.469 (2005), pp. 322–331. DOI: [10.1198/016214504000001880](https://doi.org/10.1198/016214504000001880) (p. 9)
- [21] Roger Bowden. “The Theory of Parametric Identification”. In: *Econometrica* 41.6 (1973), pp. 1069–1074. DOI: [10.2307/1914036](https://doi.org/10.2307/1914036)
- [22] Robert I. Jennrich. “Asymptotic Properties of Non-linear Least Squares Estimators”. In: *The Annals of Mathematical Statistics* 40.2 (1969), pp. 633–643. DOI: [10.1214/aoms/1177697731](https://doi.org/10.1214/aoms/1177697731)
- [23] Michael P. Keane. “A Note on Identification in the Multinomial Probit Model”. In: *Journal of Business & Economic Statistics* 10.2 (1992), pp. 193–200. DOI: [10.1080/07350015.1992.10509906](https://doi.org/10.1080/07350015.1992.10509906)
- [24] Thomas J. Rothenberg. “Identification in Parametric Models”. In: *Econometrica* 39.3 (1971), pp. 577–591. DOI: [10.2307/1913267](https://doi.org/10.2307/1913267)
- [25] George Tauchen. “Diagnostic Testing and Evaluation of Maximum Likelihood Models”. In: *Journal of Econometrics* 30 (1985), pp. 415–443. DOI: [10.1016/0304-4076\(85\)90149-6](https://doi.org/10.1016/0304-4076(85)90149-6)
- [26] Abraham Wald. “Note on the Consistency of the Maximum Likelihood Estimate”. In: *The Annals of Mathematical Statistics* 20.4 (1949), pp. 595–601. DOI: [10.1214/aoms/1177729952](https://doi.org/10.1214/aoms/1177729952)
- [27] Halbert White. “Maximum Likelihood Estimation of Misspecified Models”. In: *Econometrica* 50.1 (1982), pp. 1–25. DOI: [10.2307/1912526](https://doi.org/10.2307/1912526)